

GETTING STARTED WITH RHRV

VERSION 2.0



CONSTANTINO A. GARCÍA*, ABRAHAM OTERO, XOSÉ VILA, ARTURO MÉNDEZ,
LEANDRO RODRÍGUEZ-LIÑARES AND MARÍA JOSÉ LADO

*CONSTANTINOANTONIO.GARCIA@USC.ES

JANUARY 31, 2014

Acknowledgements

We would like to thank Matias Garcia-Constantino for his kind contributions and suggestions.

Contents

Acknowledgements	iv
List of Tables	x
List of Figures	xi
1 Overview	4
1.1 Aim	7
1.2 Structure of the document	7
2 Heart Rate Variability	9
2.1 Obtaining HRV time series	13
2.1.1 QRS detection	13
2.1.2 Constructing HRV time series	16
2.1.3 Preprocessing HRV time series	17
2.2 HRV analysis techniques	18
2.2.1 Time domain methods	18
2.2.2 Frequency domain methods	21
2.2.3 Nonlinear Methods	25

CONTENTS

2.2.3.1	Phase space reconstruction	25
2.2.3.2	Correlation dimension	27
2.2.3.3	Generalized correlation dimension	28
2.2.3.4	Information dimension	29
2.2.3.5	Sample entropy	29
2.2.3.6	Maximum Lyapunov exponent	30
2.2.3.7	Detrended Fluctuation Analysis (DFA)	31
2.2.3.8	Recurrence Quantification Analysis (RQA)	32
2.2.3.9	Poincaré plot	32
2.3	HRV alterations related to specific pathologies	34
3	Installation	38
3.1	Installation	38
3.2	WFDB applications	39
3.3	Troubleshooting	40
3.3.1	tkrplot dependency	40
4	A 15-minutes guide to RHRV	41
4.1	Preprocessing the Heart Rate series	43
4.1.1	Load heart beat positions	43
4.1.2	Calculating HR and filtering	45
4.1.3	Interpolating	46
4.1.4	Plotting	47
4.2	Analyzing the Heart Rate series	47

CONTENTS

4.2.1	Accessing “raw data”	47
4.2.2	Time-domain analysis techniques	49
4.2.3	Frequency-domain analysis techniques	53
4.2.3.1	Fourier	54
4.2.3.2	Wavelets	56
4.2.3.3	Creating several analyses	59
4.2.3.4	Plotting	61
4.2.3.5	A brief comparison: wavelets Vs. Fourier	62
5	Some more advanced features of RHRV	65
5.1	Completing our first tour	65
5.1.1	Creating the structure	66
5.1.2	Reading heart beats	67
5.1.3	Constructing the time series	68
5.1.4	Filtering the time series	69
5.1.5	Interpolation	72
5.1.6	Time analysis	72
5.1.7	Frequency analysis	73
5.2	Reading several file formats	77
5.2.1	Reading <i>RR</i> files	77
5.2.2	Reading files in WFDB format	77
5.2.3	Other formats	79
5.2.4	A general function	79
5.3	Performing analysis in different intervals of a recording	81

5.3.1	AddEpisodes	81
5.3.2	Plotting episodic information	83
5.3.3	LoadEpisodesAscii	85
5.3.4	LoadApneaWFDB	88
5.3.5	Analyzing HRV inside and outside the episodes	90
5.4	Storing and reading HRVData	95
6	HRV nonlinear analysis	97
6.1	An introduction to nonlinear analysis techniques	97
6.1.1	Nonlinearity Test	99
6.1.2	Phase space reconstruction	103
6.1.2.1	Time lag estimation	103
6.1.2.2	Embedding dimension estimation	107
6.1.3	Computing nonlinear statistics	109
6.1.3.1	The classic correlation dimension	109
6.1.3.2	Sample entropy	115
6.1.3.3	Maximum Lyapunov exponent	119
6.1.3.4	Detrended Fluctuation Analysis	123
6.1.3.5	Recurrence Quantification Analysis	127
6.2	Advanced nonlinear analysis techniques	131
6.2.1	Nonlinear noise reduction	134
6.2.2	Generalized correlation dimensions	135
6.2.3	Information dimension	136
6.2.4	Poincaré Plot	141

CONTENTS

Bibliography	146
---------------------	------------

List of Tables

2.1	Summary of the most broadly used nonlinear statistics.	26
2.2	Most important RQA statistics.	33
2.3	Clinical value of HRV analysis in cardiological diseases	35
5.1	<i>LoadBeat</i> operation depending on the <i>fileType</i> parameter.	80

List of Figures

2.1	Modulation of the heart rate by the ANS	10
2.2	Heart rate variation	11
2.3	Influence of the ANS system over the different HRV frequency bands.	14
2.4	Normal electrocardiogram.	15
2.5	High and low frequencies illustrated with sines.	22
2.6	Two wavelets	24
4.1	Non interpolated Heart Rate time plot example.	48
4.2	Interpolated Heart Rate time plot example.	49
4.3	The most important fields stored in the <i>HRVData</i> structure.	50
4.4	Plot obtained with the <i>PlotPowerBand</i> for the Fourier-based analysis.	62
4.5	Plot obtained with the <i>PlotPowerBand</i> for the wavelet-based analysis.	63
5.1	All the fields stored in the <i>HRVData</i> structure.	66
5.2	Modifying default values in the <i>FilterNIHR</i> function	70
5.3	Manually removal of artifacts with <i>EditNIHR</i>	71
5.4	Plot obtained with the <i>PlotSpectrogram</i> function.	75

LIST OF FIGURES

5.5	<i>PlotSpectrogram</i> using the <i>freqRange</i> parameter	76
5.6	Episodic information in the Non interpolated Heart Rate time series.	84
5.7	Episodic information in the interpolated Heart Rate time series.	85
5.8	Episodic information in all the power bands.	86
5.9	Loading Apnea episodes using the <i>LoadApneaWFDB</i> function.	91
6.1	Surrogate data testing.	102
6.2	Calculation of the optimum time lag using the autocorrelation	106
6.3	Estimation of the embedding dimension using the Cao's algorithm.	108
6.4	Correlation sums calculation.	112
6.5	114
6.6	Sample entropy computations.	117
6.7	Sample entropy estimation.	119
6.8	Divergence computations for the maximum Lyapunov exponent.	122
6.9	Maximum Lyapunov exponent estimation.	123
6.10	DFA.	126
6.11	Recurrence plot.	129
6.12	Recurrence plot.	130
6.13	Autocorrelation function for the RR time series.	132
6.14	Automatic estimation of the embedding dimension.	133
6.15	Generalized correlation dimension computations.	137
6.16	Generalized correlation dimension estimation.	138
6.17	Information dimension computations in <i>RHRV</i>	140
6.18	Information dimension estimation.	141

LIST OF FIGURES

6.19 “Classic” Poincaré Plot	143
6.20 Poincaré Plot using the confidence region estimation.	145

Acronyms

ANS Autonomic Nervous System.

bpm Beats Per Minute.

DFT Discrete Fourier Transform.

ECG electrocardiogram.

FFT Fast Fourier Transform.

HF High Frequency.

HR Heart Rate.

HRV Heart Rate Variability.

IRRR length of the interval determined by the first and the third quantile of the ΔRR time series.

LF Low Frequency.

Acronyms

MADRR Median of the Absolute values of the successive Differences between the RR intervals.

MODWPT Maximal Overlap Discrete Wavelet Packet Transform.

niHR Non Interpolated Heart Rate.

pNN50 proportion of successive RR intervals greater than 50 ms.

PSD Power Spectrum Density.

RMSSD Root Mean Square of Successive Differences.

RSA Respiratory Sinus Arrhythmia.

SA sinoatrial node.

SDANN Standard Deviation of the Average NN/(RR) intervals calculated over short periods.

SDNN Standard Deviation of the NN interval.

SDNN index the mean of the standard deviation calculated over the windowed RR intervals.

SDSD Standard Deviation of Successive Differences.

STFT Short Time Fourier Transform.

TINN Triangular Interpolation of NN (RR) interval histogram.

Acronyms

ULF Ultra Low Frequency.

VLF Very Low Frequency.

Chapter 1

Overview

It has been recognized in the past two decades that there is a significant relationship between the Autonomic Nervous System (ANS) and cardiovascular mortality, including sudden cardiac death. Experimental evidence for a connection between a propensity for cardiac failure and either increased sympathetic or reduced parasympathetic activity has encouraged the search of quantitative markers of autonomic activity.

One of the most promising non-invasive markers is Heart Rate Variability (HRV). HRV refers to the variation over time of both the intervals between consecutive heart beats and the instantaneous Heart Rate (HR). As the heart rhythm is modulated by the ANS, HRV is thought to reflect the activity of the sympathetic and parasympathetic branches of the ANS. The continuous modulation of the ANS results in continuous variations in heart rate. HRV has been recognized to be a

useful non-invasive tool as a predictor of several pathologies such as myocardial infarction, diabetic neuropathy, sudden cardiac death and ischemia, among others [17].

The existence of several software tools (Kubios HRV [32], the HRV toolkit for MatLab [24] or aHRV [27], just to mention a few) have helped to popularize its use. Some of these software packages are commercial and require the purchase of expensive licenses (e.g., aHRV). Even although others are free, they require the purchase of expensive commercial software on which they depend (e.g., the HRV toolkit for MatLab). Kubios is free (though not open source), but it is based on a graphical user interface, which makes it extremely tedious to perform systematic analyses of a large database of recordings, as the user must manually load and analyze through the user interface each recording. In this context, we have developed RHRV, an open-source package for the statistical environment R [12], [13], [30], [34]. To the best of our knowledge, RHRV is the only completely free and open source software package for performing HRV analysis and that is based on scripting commands; thus it enables the easy automation of analyses of a large number of recordings.

RHRV provides a complete set of tools for HRV analysis which can be used for developing new HRV analysis algorithms or for performing clinical experiments. Although this software is mainly designed for the analysis of the HRV in humans, it may also be used by animal researchers. Among the main characteristics of RHRV, we may highlight:

-
- RHRV can read heart rate data in multiple formats such as ASCII, Polar, Suunto and WFDB.
 - RHRV can compute the HRV time series from the beat positions as well as preprocessing and filtering the HRV time series to eliminate outliers or spurious points.
 - RHRV includes functionality for the visualization and manipulation of the HRV time series.
 - RHRV includes the most commonly HRV analysis techniques, with facilities for tuning the most important analysis parameters. It is possible to:
 - Perform time-domain analysis.
 - Perform frequency-domain analysis; they provide information on the renin -angiotensin system (Very Low Frequency component), both sympathetic and parasympathetic systems (Low Frequency component) and the parasympathetic system (High Frequency component). The components can be calculated using both Fourier analysis and wavelet analysis.
 - Perform nonlinear analysis techniques; they can extract some valuable information from the HRV since it responds to a complex control system.
 - RHRV can split HRV series into different segments that may correspond with different pathological states (i.e.: HRV inside and outside apnea episodes). This simplifies the statistical comparison of the heart rate inside and outside episode events.

1.1. AIM

- RHRV provides flexibility for accessing directly the internal data structures that it uses in its calculations.

The RHRV package can be freely downloaded from the R-CRAN repository [2].

1.1 Aim

The aim of this tutorial is to help the user to get started with the RHRV package for the R environment. This document supposes that the user has some basic knowledge about both the R environment and HRV. However, a short introduction to HRV will be given, and further references are provided.

1.2 Structure of the document

The remainder of this document is structured as follows. First, a brief review of several HRV topics is given in Chapter 2. This chapter contains a short discussion on the physiological origins of heart rate variability, as well as a review of the frequency components of HRV. Section 2.1 continues discussing the extraction of heart beat periods. The derivation and the preprocessing of HRV time series are also described. In Section 2.2, the most common HRV analysis methods are summarized (although they will be covered in more depth when they are introduced in the document). The descriptions of the methods are divided into time-domain, frequency-domain, and nonlinear. A discussion on the important issue of stationarity is included. The rest of the chapter (Section 2.3) is focused on the use of HRV as a predictor of different

1.2. STRUCTURE OF THE DOCUMENT

pathologies and its clinical applications.

Chapter 3 explains how to get RHRV installed in your computer. This guide assumes that you have already installed R in your computer.

Chapter 4 presents a “15-minutes guide to RHRV”. This chapter presents the essential functions needed to perform basic time and frequency domain analysis with RHRV. Chapter 5 completes the functionality introduced in Chapter 4 and presents more advanced features available in RHRV focusing on reading RR intervals stored in different formats and episodic information analysis.

Chapter 6 introduces the functionality needed to perform HRV nonlinear analysis with RHRV.

Chapter 2

Heart Rate Variability

Heart Rate Variability (HRV) describes variations over time of both instantaneous HR and the intervals between consecutive heart beats. The rhythm of the heart is modulated by the sinoatrial node (SA), which is largely influenced by both the sympathetic and parasympathetic branches of the ANS (see Figure 2.1). Sympathetic activity increases the heart rate and its response is slow (a few seconds). On the other hand the parasympathetic activity decreases the heart rate and its response is faster (0.2-0.6 seconds). Parasympathetic influence on heart rate is mediated by the action of the vagus nerve. There are also some feedback mechanisms modulating the heart rates, that try to maintain cardiovascular homeostasis by responding to the perturbations sensed by baroreceptors and chemoreceptors.

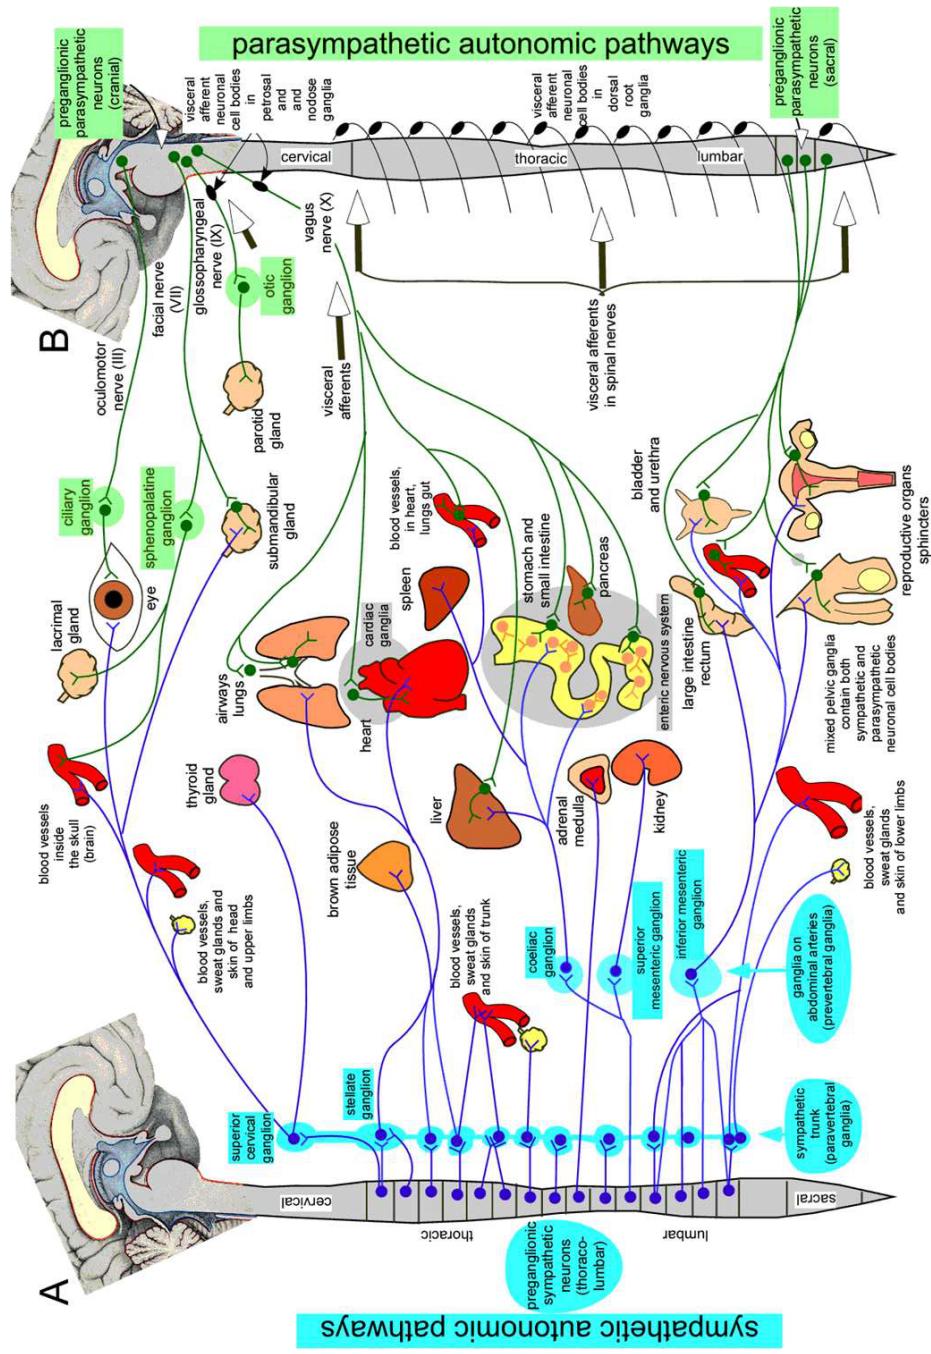


Figure 2.1: Modulation of the heart by the sympathetic and parasympathetic systems.
Figure taken from [1].

Under resting conditions, vagal tone prevails. However, parasympathetic and sympathetic activity constantly interact. The continuous modulation of the ANS results in continuous variations in heart rate as shown in Figure 2.2. The beat to beat interval variations are the result of the interaction of the beat-to-beat control mechanisms.

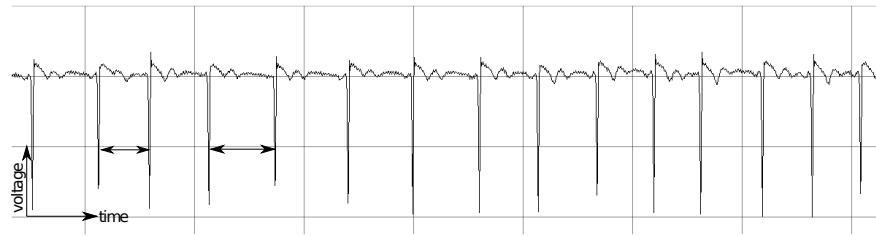


Figure 2.2: *Heart rate variation as a consequence of the modulation of the ANS.*

Due to the different speed of response of both branches of the ANS, it is possible to use the frequency analysis to discriminate between the sympathetic and parasympathetic contributions to the HRV. Akselrod et al. [4] described three components in the HRV power spectrum with physiological relevance: the Very Low Frequency (VLF) component (frequencies below 0.03 Hz), the Low Frequency (LF) component (0.03-0.15 Hz) and the High Frequency (HF) component (0.15-0.4 Hz). However, at present there is no absolute consensus on the precise limits of their boundaries.

Among all the HF mechanisms involved in the heart rate modulation we find the so called Respiratory Sinus Arrhythmia (RSA): the heartbeat synchronization with the respiratory rhythm [7]. In addition to the breathing frequencies, the HF component is believed to be of parasympathetic origin. It should be noted that, although it is

common to set the upper limit of the HF band to 0.4-0.5 Hz, it may extend up to 1 Hz for children or adults during exercise.

The LF component is a subject of controversy. Some consider that the LF phenomena is of both sympathetic and parasympathetic origin [4], [5], although some authors have suggested that the sympathetic system predominates [16], [23]. This discrepancy is due to the fact that, in conditions of sympathetic excitation, a decrease in the absolute power of the LF band is observed. This band also includes the component referred to as the 10-second rhythm or the Mayer wave, caused by oscillations in baroreceptor and chemoreceptor reflex control systems.

Spectral analysis of 24-hour recordings shows that in healthy individuals both LF and HF bands exhibit a circadian pattern and reciprocal fluctuations, with higher values of the LF in the daytime and of HF at night [10], [23].

LF and HF power can increase under different conditions. An increase of LF is observed during mental stress, standing and moderate exercise in healthy subjects, and during hypotension, physical activity and occlusion of a coronary artery or common carotid arteries in conscious dogs. On the other hand, an increase of the HF activity is observed during cold stimulation of the face, rotational stimuli and controlled respiration [9].

2.1. OBTAINING HRV TIME SERIES

The LF/HF ratio is often used by some investigators [9] as a quantitative mirror of the sympatho/vagal balance. However, other researchers disagree about the usefulness of the LF/HF index [7].

Finally, the rhythms associated with VLF have not been studied as deeply as the higher frequencies. Indeed, some authors doubt that there is a specific physiological process attributable to these heart period changes. Furthermore, the VLF band is affected by algorithms of baseline removal [9]. Despite all these objections, some authors have related the Very Low Frequency with the renin-angiotensin system. Finally, it is possible to split this band into another two: the Very Low Frequency Band (VLF, 0.003-0.03 Hz) and the Ultra Low Frequency (ULF) Band(0-0.003 Hz). Unless explicitly mentioned, the VLF band will be used to refer the (0 - 0.03 Hz) band.

Figure 2.3 summarizes the influence of the ANS system over the different HRV frequency bands.

2.1 Obtaining HRV time series

2.1.1 QRS detection

The aim of HRV analysis is to analyze the sinus rhythms while it is modulated by the ANS. Thus, the starting point for HRV analysis should be the extraction of

2.1. OBTAINING HRV TIME SERIES

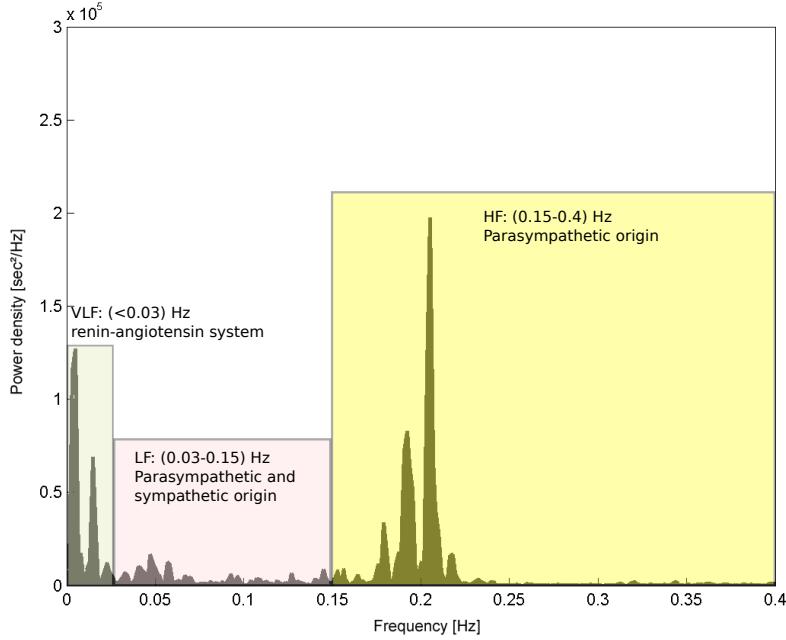


Figure 2.3: Influence of the ANS system over the different HRV frequency bands.

the SA-node action potentials from the electrocardiogram (ECG). A typical ECG showing a heartbeat consists of a P wave, a QRS complex and a T wave (see Figure 2.4). The P wave represents the wave of depolarization that spreads from the SA-node throughout the atria. The QRS complex reflects the rapid depolarization of the right and left ventricles. Since the ventricles are the largest part of the heart, in terms of mass, the QRS complex usually has a much larger amplitude than the P-wave. The T wave represents the ventricular repolarization of the ventricles. On rare occasions, a U wave can be seen following the T wave. The U wave is believed to be related to the last remnants of ventricular repolarization.

2.1. OBTAINING HRV TIME SERIES

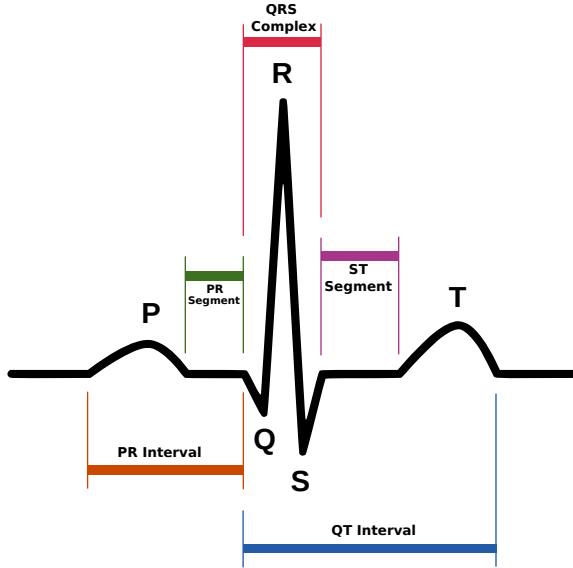


Figure 2.4: Normal electrocardiogram.

The observable that is closest related to the action of the SA-node is the P wave and, thus, the heartbeat period is defined as the time difference between two different P waves. However, the signal to noise ratio (SNR) of the P wave is smaller than the QRS complex SNR. Therefore, the QRS complexes are more easily detected than the P waves and, for convenience, the heart beat period is computed as the time difference between two successive QRS complexes. For the sake of simplicity, we will not discuss the QRS detectors in this tutorial. Further information about QRS detection may be found in [20].

2.1.2 Constructing HRV time series

After the QRS complex occurrences have been detected, the HRV time series (sometimes called the RR time series) may be calculated. The intervals between consecutive heart beats needed to construct the time series are called RR intervals, inter-beat intervals or interval function. In some context, normal-to-normal intervals (NN) may also be used when referring to these intervals.

RR intervals are computed as the difference between successive R-wave occurrence times t_n . That is, the n-th RR interval RR_n will be computed as

$$RR_n = \alpha \cdot (t_n - t_{n-1}), \quad (2.1)$$

where α is a conversion parameter that may vary depending of the units in which the RR time series will be expressed. Usually, the RR intervals are expressed in ms and thus, if the occurrence times are expressed in seconds, α is setted as $\alpha = 1000$. It must be noticed that, in some studies, the HRV is constructed as the sequence of the instantaneous heart rates. That is

$$HR_n = \frac{\beta}{t_n - t_{n-1}}. \quad (2.2)$$

Again, β is used as a conversion parameter. Since the HR is usually expressed in Beats Per Minute (bpm), $\beta = 60$ if the occurrence times are expressed in seconds. In this section, for the sake of simplicity, the RR_n construction will be used.

2.1. OBTAINING HRV TIME SERIES

The resulting RR series will consist of a set of pairs (t_n, RR_n) . It should be noted that this time series is not equidistantly sampled (that is why the time value, t_n , must be specified). This must be taken into account before frequency-domain analysis, since it requires an uniformly sampled time series. There are several approaches to overcome this issue [9]. RHRV uses interpolation for transforming the non-uniformly sampled RR series into an equidistantly sampled one. After interpolation, regular frequency analysis may be applied. A second approach, maybe the simplest one, assumes equidistant sampling and constructs a signal, called tachogram, using RR intervals as a function of a beat number. However, when using this approach, the spectrum is not a function of the frequency, rather of cycles per beat. A third approach receives the name of the spectrum of the counts, that is, it uses a series of impulses (delta functions) positioned at beat occurrence times. This approach relies on the commonly accepted Integral Pulse Frequency Modulator (IPFM) model [6], [15], that simulates the modulation of the sinoatrial node.

2.1.3 Preprocessing HRV time series

Before performing the analysis of any RR time series, a filtering operation must be carried out in order to eliminate outliers or spurious points present in the signal with unacceptable physiological values. Outliers present in the series originate from the detection of an artifact as a heartbeat (RR interval too short), or from the loss of a heartbeat in the detection procedure (RR interval too large). The RR time series may also contain some physiological artifacts. Physiological artifacts include ectopic

beats (an ectopic beat occurs when the heart beat is not triggered by the SA-node, causing an “extra” beat) and arrhythmic events. If detection of the heartbeat has been revised and corrected manually by a physician, this step can be skipped.

2.2 HRV analysis techniques

The purpose of analysis techniques usually is to extract useful physiological information that may help researchers to create new disease markers or predictors. There are several tools to perform HRV analysis, however these are usually classified into three categories: time domain methods, frequency domain methods and non-linear methods. A brief review of the main techniques of time domain, frequency domain and nonlinear methods is presented. Further information may be found at [9].

2.2.1 Time domain methods

The simplest HRV analysis techniques are the time domain measures. Since there exist a wide variety of time domain techniques, we will focus on those included in the RHRV software.

The best known time analysis statistic may be the standard deviation of the RR interval: Standard Deviation of the NN interval (SDNN).

$$SDNN = \sqrt{\frac{1}{N-1} \sum_{j=1}^N (RR_j - \bar{RR})^2}$$

2.2. HRV ANALYSIS TECHNIQUES

Since the variance is mathematically equal to the total power of spectral analysis, SDNN reflects the power of the components responsible for variability. The SDNN reflects both short-term and long-term variations within the RR series. However, it should be noted that total variance of HRV increases with the length of the analyzed recording [31]. Thus, on arbitrarily ECGs, SDNN may not be an appropriate HRV analysis variable because of its dependence with the recording's length. To avoid this issue, statistical variables calculated from segments of the total monitoring period may be used. Among this type of variables are the SDANN, the standard deviation of the average NN (RR) intervals calculated over short periods (usually 5 minutes); and the SDNN index, the mean of the standard deviation calculated over the windowed RR intervals, usually 5 minutes.

Other measures use the time series constructed as successive RR interval differences, defined as

$$\Delta RR_j = RR_{j+1} - RR_j.$$

The Standard Deviation of Successive Differences (SDSD) is given by

$$SDSD = \sqrt{\frac{1}{N-1} \sum_{j=1}^N (\Delta RR_j - \overline{\Delta RR})^2}.$$

The Root Mean Square of Successive Differences (RMSSD) is given by

$$RMSSD = \sqrt{\frac{1}{N-1} \sum_{j=1}^N (\Delta RR_j)^2}.$$

2.2. HRV ANALYSIS TECHNIQUES

Other measures using the successive RR interval differences include the length of the interval determined by the first and the third quantile of the ΔRR time series (IRRR); and the median of the absolute values of the ΔRR time series (MADRR, Median of the Absolute Differences of the RR intervals).

Other commonly used measures derived from interval differences include NN50, the number of interval differences of successive RR intervals greater than 50 ms, and pNN50, the proportion derived by dividing NN50 by the total number of RR intervals.

All these measures derived from interval differences estimate the HF variation in heart rhythm and thus, they are highly correlated.

Finally, in addition to these statistical parameters, there are some geometric measures that can be calculated from the RR interval histogram. The HRV triangular index measurement is the integral of the density distribution (that is, the number of all RR intervals) divided by the maximum of the density distribution. The density distribution may be estimated by using a histogram, thus the size of the bins should be specified. Another geometrical measure is the triangular interpolation of NN (RR) interval histogram (TINN), which is calculated as the baseline width of the distribution measured as the base of a triangle (a triangular interpolation of the histogram may be used). The TINN measure is usually expressed in milliseconds.

The major advantage of geometric methods lies in their relative insensitivity to the analytical quality of the RR series. Their major disadvantage is that they need a large number of RR intervals for performing correctly.

2.2.2 Frequency domain methods

The basic frequency domain analysis technique is the Power Spectrum Density (PSD). It provides basic information on how power distributes as a function of frequency in the RR time series. Since the sympathetic and parasympathetic branches of the ANS are associated with different frequency bands, the PSD may be a useful tool to discriminate its different contributions to the HR. The most common approach to spectral analysis of HRV is based on the Fourier transform. The Fourier transform is a tool that is able to extract the frequencies of a signal. For those unfamiliar with the “frequency” language, we will say that a signal with fast and sharp changes has “high frequencies”, whereas a signal with slow transitions is referred to as a signal with “low frequencies” (see Figure 2.5). Of course, a signal can contain both low and high frequencies. In this sense, the Fourier transform acts as a prism, separating the high frequency contributions from the low frequency contributions. The discrete implementation is referred to as the Discrete Fourier Transform (DFT) and its efficient implementation is called the Fast Fourier Transform (FFT).

The Fourier transform is one of the most powerful tools for signal processing. However, it may not be the most suitable tool for studying transient phenomena: the Fourier transform might be able to determine all the frequencies present in a

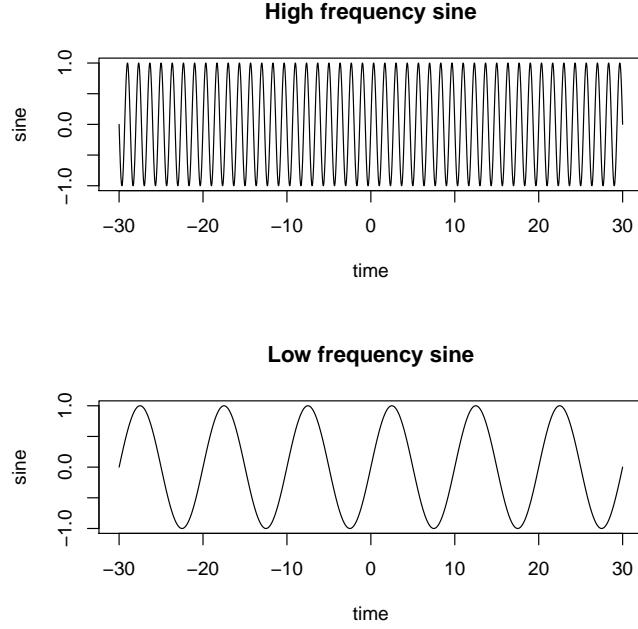


Figure 2.5: High and low frequencies illustrated with sines.

signal, but not when they are present. To address this issue, several techniques able to represent a signal in both time and frequency domain have been developed.

Following Gabor [11], the idea behind these time-frequency joint representations is to define elementary time-frequency atoms as waveforms with minimum spread in the time-frequency plane. To measure time-frequency information content, Gabor proposed decomposing signals over these elementary atoms. Selecting the time-frequency atoms is not a trivial problem because of the existence of a time-frequency uncertainty principle. This uncertainty principle states that the energy spread of a function and its Fourier transform cannot simultaneously be arbitrarily small.

2.2. HRV ANALYSIS TECHNIQUES

The simplest transform that uses this idea is the windowed Fourier transform, that is constructed by using a symmetric window that selects the portion of the signal that is going to be analyzed. The remaining portions of signal can be selected by translating the window in time. When this transform is applied to discrete signals, it is referred to as the Short Time Fourier Transform (STFT).

Another widely used transform that uses time-frequency atoms is the wavelet transform. A wavelet is a “small wave” with zero mean that grows and decays in a limited time period. Since any of these small waves results in different wavelets, there are several wavelet families. Figure 2.6 shows two such wavelets. The reference wavelet fulfilling the above conditions is called “mother wavelet”. The mother wavelet can be translated and dilated in time, yielding a set of wavelet functions with different sizes and centered in different time positions. This set of functions is used to extract time-frequency information by correlating them with the signal being analyzed.

Although the idea of the wavelet transform is similar to that used in the STFT, the wavelet transform often provides a better compromise between time and frequency resolution. This is due to the fact that the STFT uses just one window for “exploring” all the frequency bands. However, the ideal approximation would be using short windows at high frequencies and long windows at low frequencies. Thus, the “global” performance of the STFT will depend on the choice of the length of the window and the displacement time used for moving it. The wavelet transform, in contrast to the STFT, follows the ideal approximation, leading to a multiresolution

2.2. HRV ANALYSIS TECHNIQUES

analysis. RHRV has support for both approaches, and they both have a similar computational efficiency.

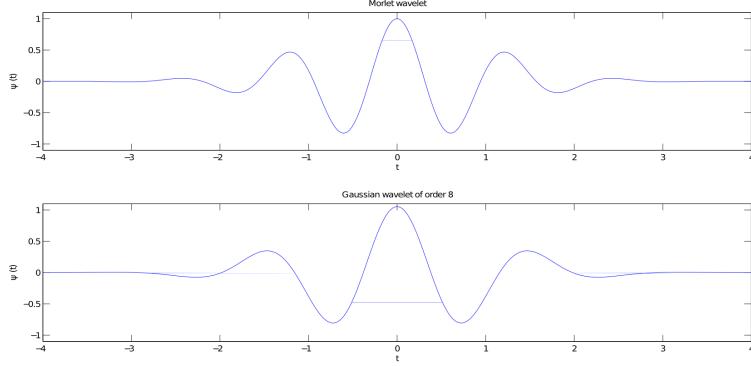


Figure 2.6: Two wavelets. The top of the figure shows the Morlet wavelet. The bottom of the figure shows a Gaussian wavelet.

When working with frequency methods, researchers are especially interested in the VLF, LF and HF frequency bands. Some authors also include the ULF band. When selecting the frequency bands, the researchers should take into account whether they are working with short (2-5 min) or long term recordings (up to 24-hours). Three main spectral components are distinguished in a spectrum calculated from short-term recordings: VLF, LF and HF components. However, VLF assessed from short-term recordings is a dubious parameter and, therefore, it should be avoided when interpreting the PSD in this type of recordings [9]. Spectral analysis resulting from long-term recordings include VLF, LF and HF bands. In the long recordings, the VLF band may be split into the ULF and the VLF components.

2.2.3 Nonlinear Methods

There is a profound connection between nonlinear phenomena and HRV. HRV is determined by complex interactions of electrophysiological and humoral variables, as well as by autonomic and central nervous regulations. Considering these complex control systems modulating the heart rhythm, it has been speculated that methods of the nonlinear dynamics might extract some valuable information from the HRV series.

A wide variety of nonlinear statistics have already been used in the HRV literature, including largest Lyapunov exponent, generalized correlation dimension, SD1/SD2 of Poincaré plots, detrended fluctuation analysis, sample entropy and recurrence quantification analysis. Table 2.1 summarizes the most important nonlinear statistics that have been included in *RHRV*. In the next sections, we shall present a quick theoretical review of all these methods. More details will be given later in the tutorial when we show how to use these methods in RHRV.

2.2.3.1 Phase space reconstruction

A large amount of nonlinear algorithms is based on the concept of phase space. For a deterministic system, the phase space is the collection of all possible system states. That is, each point of the phase space represents all the information needed to determine the evolution of the system. Of course, the problem now is: How can

2.2. HRV ANALYSIS TECHNIQUES

Statistic	Interpretation
Maximum Lyapunov Exponent	Quantifies the rate of divergence of close trajectories
Generalized Correlation Dimension	Quantifies the dimensionality of the reconstructed phase space
Sample Entropy	Measures the complexity of the time series being studied
Recurrence Quantification Analysis (RQA)	Quantifies the number and duration of recurrences of a time series in its phase space
Detrended Fluctuation Analysis (DFA)	Quantifies the presence of fractal correlation properties in non-stationary data
Poincaré Plot	Characterizes the system dynamics by using a two dimensional embedding

Table 2.1: Summary of the most broadly used nonlinear statistics.

we transform an univariate time series (the RR time series) in a multivariate phase space?

The Takens embedding theorem answers this question. Takens proved that phase space reconstruction from a single time series $x(n)$ could be achieved by using the vectors:

$$\mathbf{x}_i = [x(i), x(i + \tau), \dots, x(i + (m - 1) \cdot \tau)], \quad (2.3)$$

Sections 6.1.2.1 and 6.1.2.2 deals with the problem of selecting both m (the so-called embedding dimension) and τ (the time lag parameter) using *RHRV*.

2.2.3.2 Correlation dimension

The correlation dimension is the most common measure of the fractal dimensionality of a geometrical object embedded in a phase space. The estimation of the correlation dimension requires the computation of the so-called correlation sum $C(r)$. The correlation sum is defined over the N points from the phase space as follows:

$$C(r) = \frac{\#\{(\mathbf{x}_i, \mathbf{x}_j) : \text{distance}(\mathbf{x}_i, \mathbf{x}_j) < r\}}{N^2},$$

where $\#$ represents the cardinality of the set and r a radius in the embedding dimension. However, this estimator is biased when the pairs in the sum are not statistically independent. For example, Taken's vectors that are close in time, are usually close in the phase space due to the non-zero autocorrelation of the original time series. This is solved by using the so-called Theiler window: two Takens' vectors must be separated by, at least, the time steps specified by this window in order to be considered neighbours. By using a Theiler window, we exclude temporally correlated vectors from our estimations.

Chaotic attractors are expected to fulfill

$$C(r) \propto r^D,$$

being D the correlation dimension that we want to estimate. Thus, the correlation dimension may be estimated using the slope obtained by performing a linear regression of $\log_{10}(C(r))$ Vs. $\log_{10}(r)$. Since this dimension is supposed to be an invariant

of the system, it should not depend on the dimension of the Taken's vectors used to estimate it (provided that we are using an embedding dimension that is large enough to reconstruct the phase space). Thus, the user should plot $\log(C(r))$ Vs. $\log(r)$ for several embedding dimensions when looking for the correlation dimension and, if for some range $\log(C(r))$ shows a similar linear behaviour in different embedding dimensions (i.e. parallel slopes), these slopes are an estimate of the correlation dimension. This is very important! If the slope depends on the embedding dimension (provided that we are embedding the time series in a phase space with sufficient dimensions) we cannot talk about a correlation dimension. Furthermore, the time series may not be chaotic. More details about this requirement shall be given when presenting the *RHRV* functionality for computing the correlation dimension.

2.2.3.3 Generalized correlation dimension

Note that the correlation sum $C(r)$ may be interpreted as: $C(r) = \langle p(r) \rangle$, that is: the mean probability of finding a neighbour in a ball of radius r surrounding a point in the phase space. Thus, it is possible to define a generalization of the correlation dimension by writing:

$$C_q(r) = \langle p(r)^{(q-1)} \rangle .$$

With this notation, the “classic” correlation sum is $C(r) = C_2(r)$. It is possible to determine generalized dimensions D_q using the slope obtained by performing a linear regression of $\log(C_q(r))$ Vs. $(q-1)\log(r)$. The case $q = 1$ leads to the information

dimension, that is treated separately in this package. see Sections 2.2.3.4 and 6.2.3). The considerations discussed for the correlation dimension estimate are also valid for these generalized dimensions.

2.2.3.4 Information dimension

The information dimension is a particular case of the generalized correlation dimension when setting the order $q = 1$. It is possible to demonstrate that it can be defined as:

$$D_1 = \lim_{r \rightarrow 0} \langle \log p(r) \rangle / \log(r), \quad (2.4)$$

being D_1 the information dimension, $p(r)$ is the probability of finding a neighbour in a neighbourhood of size r and $\langle \cdot \rangle$ is the mean value. Thus, the information dimension specifies how the average Shannon information scales with the radius r .

In order to estimate D_1 in practical applications, a variation of equation 2.4 is used. This algorithm looks for the scaling behaviour of the average radius that contains a given portion (a "fixed-mass" p) of the total points in the phase space. By performing a linear regression of $\log(p)$ Vs. $\log(\langle r \rangle)$, an estimate of D_1 is obtained.

2.2.3.5 Sample entropy

The sample entropy measures the complexity of a time series. Large values of the Sample Entropy indicate high complexity whereas that smaller values characterize more regular signals. The sample entropy of order q is computed in the RHRV

package by using the correlation sums calculated with the *CalculateCorrDim*. We first define the function:

$$h_q(m, r) = \log \left(\frac{C_q(m, r)}{C_q(m + 1, r)} \right),$$

where m is the embedding dimension and q the order of the correlation sum. The Sample entropy (or Renyi entropy) of order q H_q fulfills

$$H_q = \lim_{\substack{r \rightarrow 0 \\ m \rightarrow \infty}} h_q(m, r). \quad (2.5)$$

2.2.3.6 Maximum Lyapunov exponent

Close trajectories diverge exponentially fast in a chaotic system. The averaged exponent that determines the divergence rate is called the Lyapunov exponent (usually denoted with λ). If $\delta(0)$ is the distance between two Takens' vectors in a m -dimensional space, we expect that the distance after a time t between the two trajectories arising from this two vectors fulfills:

$$\delta(t) \propto \delta(0) \cdot \exp(\lambda t).$$

Thus, the Lyapunov exponent is estimated using the slope obtained by performing a linear regression of $S(t) = \lambda \cdot t \approx \log(\delta(t)/\delta(0))$ on t . In practical applications, we should check the existence of a linear region when plotting $S(t)$ Vs. t . If for some temporal range this plot shows a linear behaviour, its slope is an estimate of the maximal Lyapunov exponent per unit of time. If such a region does not exist, the

estimation should be discarded.

Also, just as for the correlation dimension computations, the maximal Lyapunov exponent should be computed for several embedding dimensions in order to check that it does not depend on the embedding dimension.

2.2.3.7 Detrended Fluctuation Analysis (DFA)

The Detrended Fluctuation Analysis (DFA) is a widely used technique for detecting correlations in time series. These functions are able to estimate several scaling exponents from the RR time series being analyzed. These scaling exponents characterize short or long-term fluctuations. The DFA procedure may be summarized as follows:

1. Integrate the time series to be analyzed. The time series resulting from the integration will be referred to as the *profile*.
2. Divide the profile into N non-overlapping segments.
3. Calculate the local trend for each of the segments using least-square regression.
Compute the total error for each of the segments.
4. Compute the average of the total error over all segments and take its root square. By repeating the previous steps for several segment sizes (let's denote it by t : number of beats), we obtain the so-called *fluctuation function* $F(t)$.
5. If the data presents long-range power law correlations: $F(t) \propto t^\alpha$, we can estimate the exponent using regression.

6. Usually, when plotting $\log(F(t))$ Vs $\log(t)$ we may distinguish two linear regions. By performing a regression on each of them separately, we obtain two scaling exponents, α_1 (the exponent for small values of t , characterizing short-term fluctuations) and α_2 (the exponent for large values of t , characterizing long-term fluctuations).

2.2.3.8 Recurrence Quantification Analysis (RQA)

The Recurrence Quantification Analysis (RQA) is an advanced technique for the nonlinear analysis that allows to quantify the number and duration of the recurrences in the phase space. A recurrence is a time instant in which the trajectory returns to a phase space region it has visited before. Thus, it is a representation of those instants of time in which $\mathbf{x}_i \approx \mathbf{x}_j$ for every i and j . The recurrence plot is the graphical representation of the recurrence matrix of the RR time series. The *RQA* function allows to compute several statistics derived from the RQA analysis of the RR time series. Table 2.2 summarize the most important RQA statistics and its meaning.

2.2.3.9 Poincaré plot

The Poincaré plot is a graphical representation of the dependance between successive RR intervals obtained by plotting the $RR_{j+\tau}$ as a function of RR_j . This dependance is often quantified by fitting an ellipse to the plot. In this way, two parameters are obtained characterizing the ellipse: SD_1 and SD_2 . When $\tau = 1$, SD_1 is usually calculated as the standard deviation of the points perpendicular to the line of identity

2.2. HRV ANALYSIS TECHNIQUES

Statistic	RHRV name	Interpretation
Recurrence	<i>REC</i>	Percentage of recurrence points in a Recurrence Plot
Determinism	<i>DET</i>	Percentage of recurrence points that form diagonal lines
Laminarity	<i>LAM</i>	Percentage of recurrent points that form vertical lines
Ratio	<i>RATIO</i>	Ratio between DET and RR
Longest diagonal line	<i>Lmax</i>	Length of the longest diagonal line
Averaged diagonal line length	<i>Lmean</i>	Mean length of the diagonal lines. The main diagonal is not taken into account
Divergence	<i>DIV</i>	Inverse of Lmax
Longest vertical line	<i>Vmax</i>	Longest vertical line
Trapping time	<i>Vmean</i>	Average length of the vertical lines.
Entropy	<i>ENTR</i>	Shannon entropy of the diagonal line lengths distribution
Trend	<i>TREND</i>	Trend of the number of recurrent points depending on the distance to the main diagonal
Recurrence Rate	<i>recurrenceRate</i>	Number of recurrent points depending on the distance to the main diagonal

Table 2.2: Most important RQA statistics.

2.3. HRV ALTERATIONS RELATED TO SPECIFIC PATHOLOGIES

and SD_2 is calculated as the standard deviation along the line of identity. In terms of time-domain parameters:

$$SD_1^2 = \frac{1}{2} SDSD^2$$
$$SD_2^2 = 2 \cdot SDNN^2 - \frac{1}{2} SDSD^2$$

In this way SD_1 characterizes short-term variability whereas that SD_2 characterizes long-term variability. However, sometimes the ellipse that is fitted using this approach is too small. *RHRV* also allows the user to fit a ellipse by estimating a confidence region. If $\tau > 1$, the confidence region approach is always used. More details shall be given in 6.2.4.

2.3 HRV alterations related to specific pathologies

In the course of the last two decades numerous studies have shown HRV to be a useful tool as a predictor of several pathologies such as myocardial infarction, sudden cardiac death, heart failure, hypertension, and ischemia, among others [22]. However, it should be noted that the practical use of HRV has reached general consensus only in two clinical applications: as a predictor of risk after myocardial infarction and as an early warning of diabetic neuropathy. [9], [18].

Table 2.3 resumes some HRV applications to other diseases.

2.3. HRV ALTERATIONS RELATED TO SPECIFIC PATHOLOGIES

Table 2.3: Summary of the clinical value of HRV analysis in cardiological diseases. Inspired by [9].

Disease state	Clinical finding	Potential value
Myocardial infarction (MI)	↓ HRV after myocardial infarction (MI). In the severe phase of MI, there is a ↓ standard deviation of the HRV signal	Depressed HRV is a powerful predictor of mortality and of arrhythmic complications in patients following acute MI
		HRV analysis is useful for risk stratification of patients following MI
Diabetic neuropathy	↓ time-domain parameters of HRV preceded the clinical detection of autonomic neuropathy. ↓ LF and HF bands in diabetic patients with no signs of autonomic neuropathy	HRV analysis may be used as predictor of diabetic autonomic neuropathy occurrence
Hypertension	↑ LF found in hypertensives with circadian patterns	Hypertension is characterized by depressed circadian rhythmicity of LF
	Reduced parasympathetic activity in hypertensive patients	
Congestive heart failure (CHF)	↓ spectral power in all frequencies, especially > 0.04 Hz	In CHF, there is ↓ vagal, but relatively preserved sympathetic modulation of HR
	Low HRV	Reduced vagal activity in CHF patients
	↓ HF power in CHF. ↑ LF/HF	Low parasympathetic tone in CHF. CHF produces imbalance of autonomic tone with ↓ parasympathetic and predominance of sympathetic tone
Continued on next page		

2.3. HRV ALTERATIONS RELATED TO SPECIFIC PATHOLOGIES

Table 2.3 – continued from previous page

Disease state	Clinical findings	Potential value
	Alterations in HRV not tightly linked to severity of CHF. ↓ HRV was related to sympathetic excitation	
	↑ HRV during ACE (angiotensin-converting-enzyme) inhibitor treatment	Increase of the sympathetic tone associated with ACE inhibitor therapy
Heart Transplantation	HRV from 0.02 to 1 Hz is 90% reduced	Patients with rejection show less variability
Chronic mitral regurgitation	HR techniques correlated with ventricular performance and predicted clinical events	Prognostic indicator of atrial fibrillation, mortality and progression to valve surgery
Mitral Valve prolapse (MVP)	↓ HF power	MVP patients had low vagal tone
Cardiomyopathies	Global and specific vagal tone measurements of HRV were ↓ in symptomatic patients	
Sudden death (SD) or cardiac arrest (CA)	LF power and standard deviation of HRV signals were related to 1 year mortality	HRV is useful to risk stratify CA survivors for 1 year mortality
	↓ HF power in CA survivors	
	Both time and frequency domain indexes separated controls from SD patients. ↓ HF power was the best separator between heart disease patients with and without SD	HF power may be useful predictor of SD
	SDNN index was lower in SD patients	Time domain indexes may identify increased risk of SD
Continued on next page		

2.3. HRV ALTERATIONS RELATED TO SPECIFIC PATHOLOGIES

Table 2.3 – continued from previous page

Disease state	Clinical findings	Potential value
Ventricular arrhythmias	HRV indexes do not change consistently before ventricular fibrillation (VF). All power spectra of HRV were significantly ↓ before the onset of sustained ventricular tachycardia (VT) than before non sustained VT	A temporal relation exists between the decrease of HRV and the onset of sustained VT

Chapter 3

Installation

3.1 Installation

This guide assumes that the user has some basic knowledge of the R environment. If this is not your case, you can find a nice introduction to R in the R project homepage [3]. The R project homepage also provides an “R Installation and Administration” guide. Once you have download and installed R, you can install RHRV by typing:

```
> install.packages("RHRV")
```

You can also install it by downloading it from the CRAN [2]. Once the download has finished, open R, move to the directory where you have download it (by using the R command *setwd*) and type:

```
> install.packages("RHRV_XXX", repos=NULL)
```

Here, XXX is the version number of the library. To start using the library, you should load it by using the *library* command:

> *library(RHRV)*

3.2 WFDB applications

Some functions of the RHRV package (such as the *LoadApneaWFDB*) require the installation of the WFDB functions [25]. If the user is not going to work with WFDB formatted files, the installation of these libraries is not required for the proper functioning of RHRV. The WFDB functions is a large collection of specialized software for processing and manipulating the PhysioNet's databases [14]. On Windows and Mac OSX operating systems is necessary to define a .Renvironment file in the user workspace indicating the directory of the WFDB commands. Examples for both OS are given below:

```
## .Renvironment on Windows  
PATH = "c:\\cygwin\\bin"  
DYLD_LIBRARY_PATH = "c:\\cygwin\\lib"  
  
## .Renvironment on Macosx  
PATH = "/opt/local/bin"  
DYLD_LIBRARY_PATH = "/opt/local/bin"
```

3.3 Troubleshooting

3.3.1 When installing the RHRV package in linux, sometimes the installation fails when installing the tkplot dependency.

...

```
tcltkimg.c:2:16: fatal error: tk.h: No such file or directory  
compilation terminated.
```

```
ERROR: compilation failed for package 'tkplot'
```

...

```
ERROR: dependency 'tkplot' is not available for package 'RHRV'
```

This is usually because there are some missing libraries in your system. Generally, the problem will be fixed by installing the *tclX.X*, *tkX.X*, *tclX.X-dev* and *tkX.X-dev* libraries (X.X stands for the version of the libraries).

Chapter 4

A 15-minutes guide to RHRV

In this chapter, a brief description of the RHRV package is presented [30]. Due to the large collection of features that RHRV offers, in this chapter we shall refer only to the most important functionality for performing a basic HRV analysis. In the next chapter we will present more advanced functionality of the package, or functionality geared to certain particular types of analysis. RHRV can be freely downloaded from the R-CRAN repository [2].

We propose the following basic program flow to perform HRV analysis using the RHRV package:

1. Load heart beat positions. For the sake of simplicity, in this section we will focus in ASCII files.
2. Build the instantaneous HR series and filter it to eliminate spurious points.
3. Plot the instantaneous HR series.

-
4. Interpolate the instantaneous HR series to obtain a HR series with equally spaced values.
 5. Plot the interpolated HR series.
 6. Perform the desired analysis. The user can perform time-domain analysis, frequency-domain analysis and/or nonlinear analysis.
 7. Plot the results of the analysis that has been performed and access the “raw” data.

In section 4.1 we will address points 1-5, whereas in section 4.2 we will deal with points 6 and 7. All the examples of this chapter will use the example beats file “example.beats” that may be downloaded from <http://rhrv.r-forge.r-project.org/>. Additionally, the data from this file has been included in RHRV. The user can access this data executing:

```
> # HRVData structure containing the heart beats  
> data("HRVData")  
> # HRVData structure storing the results of processing the  
> # heart beats: the beats have been filtered, interpolated, ...  
> data("HRVPprocessedData")
```

The example file is an ASCII file that contains the beat positions obtained from a 2 hours ECG (one beat position per row). The subject of the ECG is a patient suffering from paraplegia and hypertension (systolic blood pressure above 200 mmHg). During the recording, he is supplied with prostaglandin E1 (a vasodilator

that is rarely employed) and systolic blood pressure fell to 100 mmHg for over an hour. Then, the blood pressure increased slowly up to approximately 150 mmHg.

The console output shall be shown for every example.

4.1 Preprocessing the Heart Rate series

4.1.1 Load heart beat positions

RHRV uses a custom data structure called *HRVData* to store all HRV information related to the signal being analyzed. *HRVData* is implemented as a list object in R language. This list contains all the information corresponding to the imported signal to be analyzed, some parameters generated by the pre-processing functions and the HRV analysis results. A new *HRVData* structure is created using the *CreateHRVData* function. In order to obtain detailed information about the operations performed by the program, we can activate a verbose mode using the *SetVerbose* function.

```
> hrv.data = CreateHRVData()  
> hrv.data = SetVerbose(hrv.data, TRUE )
```

After creating the empty *HRVData* structure the next step should be loading the signal that we want to analyze into this structure. RHRV imports data files containing heart beat positions. Supported formats include ASCII (*LoadBeatAscii* function), EDF (*LoadBeatEDFPlus*), Polar (*LoadBeatPolar*), Suunto (*LoadBeatSuunto*) and WFDB data files (*LoadBeatWFDB*) [26]. For the sake of simplicity, we will focus

4.1. PREPROCESSING THE HEART RATE SERIES

in ASCII files containing one heart beat occurrence time per line. We also assume that the beat occurrence time is specified in seconds (further details will be given in chapter 5). For example, let's try to load the “example.beats” file, whose first lines are shown below. Each line denotes the occurrence time of each heartbeat.

```
0  
0.3280001  
0.7159996  
1.124  
1.5  
1.88
```

In order to load this file, we may write:

```
> hrv.data = LoadBeatAscii(hrv.data, "example.beats",  
+           RecordPath = "beatsFolder")  
  
** Loading beats positions for record: example.beats **  
Path: beatsFolder  
Scale: 1  
Date: 01/01/1900  
Time: 00:00:00  
Number of beats: 17360
```

The console information is only displayed if the verbose mode is on. The *Scale* parameter is related to the time units of the file. 1 denotes seconds, 0.1 tenth

of seconds and so on. The *Date* and *Time* parameters specify when the file was recorded. More details about these parameters will be given in section 5.1.2. The *RecordPath* can be omitted if the *RecordName* is in the working directory.

4.1.2 Calculating HR and filtering

To compute the HRV time series the *BuildNIHR* function can be used (*Build Non Interpolated Heart Rate*). This function constructs both the RR (Equation 2.1) and instantaneous heart rate (HR) series (Equation 2.2) described in Section 2.1. We will refer to the instantaneous heart rate (HR) as the Non Interpolated Heart Rate (niHR) series. Both series are stored in the *HRVData* structure.

```
> hrv.data = BuildNIHR(hrv.data)

** Calculating non-interpolated heart rate **

Number of beats: 17360
```

A filtering operation must be carried out in order to eliminate outliers or spurious points present in the niHR time series with unacceptable physiological values. Outliers present in the series originate both from detecting an artifact as a heartbeat (RR interval too short) or not detecting a heartbeat (RR interval too large). The outliers removal may be both manual or automatic. In this quick introduction, we will use the automatic removal. The automatic removal of spurious points can be performed by the *FilterNIHR* function. The *FilterNIHR* function also eliminates

points with unacceptable physiological values.

```
> hrv.data = FilterNIHR(hrv.data)

** Filtering non-interpolated Heart Rate **

Number of original beats: 17360

Number of accepted beats: 17259
```

4.1.3 Interpolating

In order to be able to perform spectral analysis in the frequency domain, a uniformly sampled HR series is required. It may be constructed from the niHR series by using the *InterpolateNIHR* function, which uses linear (default) or spline interpolation (further details on chapter 5). The frequency of interpolation may be specified. 4 Hz (the default value) is enough for most applications.

```
> # Note that it is not necessary to specify freqhr since it matches with
> # the default value: 4 Hz

> hrv.data = InterpolateNIHR (hrv.data, freqhr = 4)

** Interpolating instantaneous heart rate **

Frequency: 4Hz

Number of beats: 17259

Number of points: 29594
```

4.1.4 Plotting

Before applying the different analysis techniques that RHRV provides, it is usually interesting to plot the time series with which we are working. The *PlotNIHR* function permits the graphical representation of the niHR series whereas the *PlotHR* function permits to graphically represent the interpolated HR time series.

```
> PlotNIHR(hrv.data)
```

```
> PlotHR(hrv.data)
```

The plots obtained with *PlotNIHR* and *PlotHR* are shown in Figures 4.1 and 4.2, respectively.

As seen in the Figures 4.1 and 4.2, the patient initially had a heart rate of approximately 160 beats per minute. Approximately half an hour into record the prostaglandina E1 was provided, resulting in a drop in heart rate to about 130 beats per minute during about 40 minutes, followed by a slight increase in heart rate.

4.2 Analyzing the Heart Rate series

4.2.1 Accessing “raw data”

In the previous sections, we have used the *HRVData* structure to store all HRV information related to the signal being analyzed with no knowledge about its internal

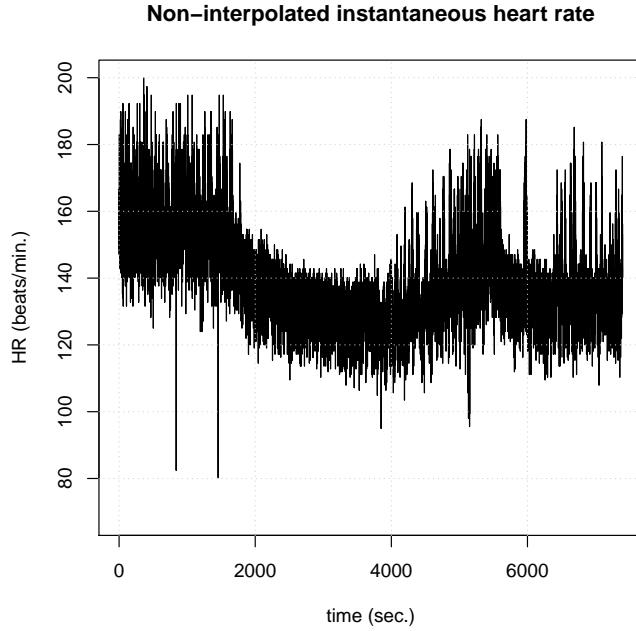


Figure 4.1: Non interpolated Heart Rate time plot example.

structure. However, sometimes, in order to make some particular analysis of the data, it may be interesting to access them directly. Figure 4.3 summarizes the most important fields in the *HRVData* structure. Since all the data in this structure is stored as an R list, each of its fields can be accessed using the \$ operator of the R language. For example, if we want to access the RR time series of the *hrv.data*, we would use:

```
> RR = hrv.data$Beat$RR
```

Although it is an advantage to be familiarized with the *HRVData* structure, there is no need to memorize it since we can use the useful *name* R function. Thus, if we want to know which fields are stored into the *hrv.data\$Beat* subfield, we could use:

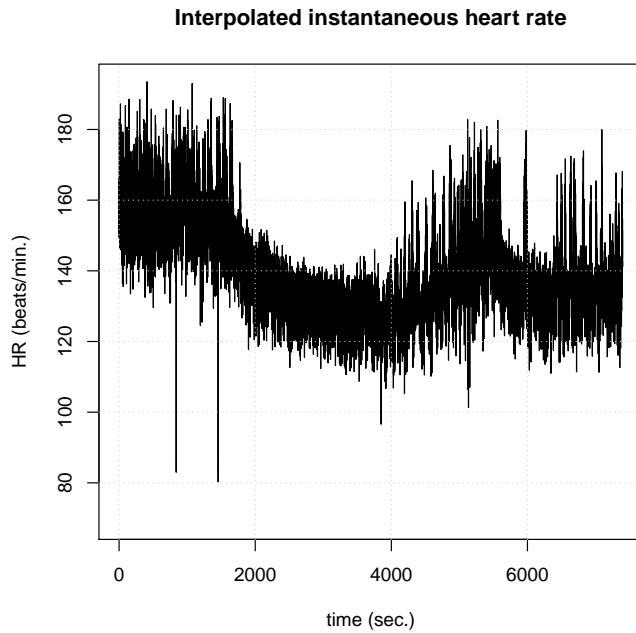


Figure 4.2: Interpolated Heart Rate time plot example.

```
> names(hrv.data$Beat)
```

```
[1] "Time" "niHR" "RR"
```

As we can see, *hrv.data\$Beat* stores the occurrence time of each beat (“*Time*”), the niHR time series (“*niHR*”) and the RR time series (“*RR*”).

4.2.2 Time-domain analysis techniques

The simplest way of performing a HRV analysis in RHRV is using the time analysis techniques provided by the *CreateTimeAnalysis* function. This function computes the time-domain parameters presented in Section 2.2.1 and stores them in the *HRV-Data* structure. The most interesting parameter that the user may specify is the

4.2. ANALYZING THE HEART RATE SERIES

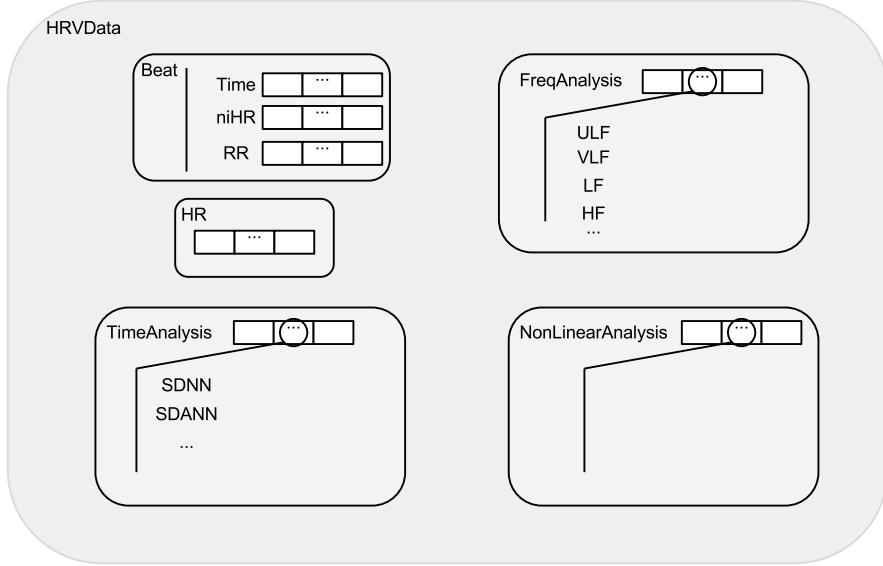


Figure 4.3: The most important fields stored in the `HRVData` structure.

width of the window that will be used to analyze short segments from the RR time series (`size` parameter, in seconds). Specifically, several statistics will be computed for each window. By studying how these statistics evolve through the recording, a set of time parameters will be computed (For example, the `SDANN` and `SDNNIDX` parameters). Other important argument that can be tuned is the interval width of the bins that will be used to compute the histogram (`interval` parameter). As an alternative to the `interval` parameter, the user may use the `numofbins` parameter to specify the number of bins in the histogram. A typical value for the `size` parameter is 300 seconds (which is also the default value), whereas that a typical value for the `interval` is about 7.8 milliseconds (also default value).

4.2. ANALYZING THE HEART RATE SERIES

```
> hrv.data = CreateTimeAnalysis(hrv.data, size = 300,
+                                interval = 7.8125)

** Creating time analysis

Size of window: 300 seconds

Width of bins in histogram: 7.8125 milliseconds

Number of windows: 24

Data has now 1 time analyses

SDNN: 39.81504 msec.

SDANN: 31.11223 msec.

SDNNIDX: 25.05384 msec.

pNN50: 9.39854 %

SDSD: 31.07026 msec.

r-MSSD: 31.06936 msec.

IRRR: 32 msec.

MADRR: 16 msec.

TINN: 86.10213 msec.

HRV index: 11.02107
```

If the verbose mode is on, the program will display the results of the calculations on the screen. Otherwise, the user must access the “raw” data as explained before to obtain the results.

Finally, we show a complete example for performing a basic time-domain analysis. The console output is also shown. It should be noted that it is not necessary to

4.2. ANALYZING THE HEART RATE SERIES

perform the interpolation process before applying the time-domain techniques since these parameters are calculated directly from the RR-time series.

```
> hrv.data = CreateHRVData()  
  
> hrv.data = SetVerbose(hrv.data, FALSE)  
  
> hrv.data = LoadBeatAscii(hrv.data, "example.beats", "beatsFolder")  
  
> hrv.data = BuildNIHR(hrv.data)  
  
> hrv.data = FilterNIHR(hrv.data)  
  
> PlotNIHR(hrv.data)  
  
> hrv.data = SetVerbose(hrv.data, TRUE)  
  
> hrv.data = CreateTimeAnalysis(hrv.data, size=300, interval = 7.8125)  
  
** Creating time analysis  
  
Size of window: 300 seconds  
  
Width of bins in histogram: 7.8125 milliseconds  
  
Number of windows: 24  
  
Data has now 1 time analyses  
  
SDNN: 39.81504 msec.  
  
SDANN: 31.11223 msec.  
  
SDNNIDX: 25.05384 msec.  
  
pNN50: 9.39854 %  
  
SDSD: 31.07026 msec.  
  
r-MSSD: 31.06936 msec.  
  
IRRR: 32 msec.
```

```
MADRR: 16 msec.
```

```
TINN: 86.10213 msec.
```

```
HRV index: 11.02107
```

```
> # We can access "raw" data... let's print separately, the SDNN  
> # parameter  
> cat("The SDNN has a value of ",hrv.data$TimeAnalysis[[1]]$SDNN," msec.\n")
```

```
The SDNN has a value of 39.81504 msec.
```

4.2.3 Frequency-domain analysis techniques

A major part of the functionality of the RHRV package is dedicated to the spectral analysis of HR signals. Before performing the frequency analysis, a data analysis structure must be created. Such structure shall store the information extracted from a variability analysis of the HR signal as a member of the *FreqAnalysis* list, under the *HRVData* structure. Each analysis structure created is identified by a unique number (in order of creation). To create such an analysis structure, the *CreateFreqAnalysis* function is used.

```
> hrv.data = CreateFreqAnalysis(hrv.data)  
  
** Creating frequency analysis  
  
Data has now 1 frequency analysis
```

Notice that, if verbose mode is on, the *CreateFreqAnalysis* function informs us about the number of frequency analysis structures that have been created. In order to

4.2. ANALYZING THE HEART RATE SERIES

select a particular spectral analysis, we will use the *indexFreqAnalysis* parameter in the frequency analysis functions.

The most important function to perform spectral HRV analysis is the *CalculatePowerBand* function. The *CalculatePowerBand* function computes the spectrogram of the HR series in the ULF, VLF, LF and HF frequency bands using STFT or wavelets. Boundaries of the bands may be chosen by the user. If boundaries are not specified, default values are used: ULF, [0, 0.03] Hz; VLF, [0.03, 0.05] Hz; LF, [0.05, 0.15] Hz; HF, [0.15, 0.4] Hz. The type of analysis can be selected by the user by specifying the *type* parameter of the *CalculatePowerBand* function. The possible options are either “fourier” or “wavelet”. Because of the backwards compatibility, the default value for this parameter is “fourier”.

4.2.3.1 Fourier

When using the STFT to compute the spectrogram employing the *CalculatePowerBand* function, the user may specify the following parameters related with the STFT:

- *Size*: the size of window for calculating the spectrogram measured in seconds.

The RHRV package employs a Hamming window to perform the STFT.

- *Shift*: the displacement of window for calculating the spectrogram measured in seconds.

4.2. ANALYZING THE HEART RATE SERIES

- *Sizesp*: the number of points for calculating each window of the STFT. Thus, it is highly recommended to select *sizesp* so that $sizesp = 2^N$. If the user does not specify it, the program selects a proper length for the calculations.

When using *CalculatePowerBand*, the *indexFreqAnalysis* parameter (in order to indicate which spectral analysis we are working with) and the boundaries of the frequency bands may also be specified.

As an example, let's perform a frequency analysis in the typical HRV spectral bands based on the STFT . We may select 300 s (5 minutes) and 30 s as window size and displacement values because these are typical values when performing HRV spectral analysis. The value of the zero-padding should be chosen to be greater than the number of samples of the window size. Assuming that the sampling frequency is 4 Hz, the zero-padding value must fulfill $sizesp \geq size \cdot f_s$. In this occasion, we select the smallest power of 2 that meets the previous condition: $sizesp = 2048 = 2^{11} > 1200 = 300 \cdot 4$. Thus, we may write:

```
> hrv.data = CreateHRVData( )  
> hrv.data = SetVerbose(hrv.data, FALSE)  
> hrv.data = LoadBeatAscii(hrv.data, "example.beats", "beatsFolder")  
> hrv.data = BuildNIHR(hrv.data)  
> hrv.data = FilterNIHR(hrv.data)  
> hrv.data = InterpolateNIHR (hrv.data, freqhr = 4)  
> hrv.data = CreateFreqAnalysis(hrv.data)  
> hrv.data = SetVerbose(hrv.data, TRUE)
```

4.2. ANALYZING THE HEART RATE SERIES

```
> # Note that it is not necessary to write the boundaries  
> # for the frequency bands, since they match  
> # the default values  
  
> hrv.data = CalculatePowerBand( hrv.data , indexFreqAnalysis= 1,  
+ size = 300, shift = 30, sizesp = 2048, type = "fourier",  
+ ULFmin = 0, ULFmax = 0.03, VLFmin = 0.03, VLFmax = 0.05,  
+ LFmin = 0.05, LFmax = 0.15, HFmin = 0.15, HFmax = 0.4 )  
  
** Calculating power per band **  
** Using Fourier analysis **  
  
Widowing signal... 237 windows  
  
Power per band calculated
```

Alternatively, we could not specify the *sizesp* parameter and let the program decide for us. In fact, the program would use the same criteria that we used in the previous example. Thus, we could have used the following sentence to obtain exactly the same results:

```
> hrv.data = CalculatePowerBand( hrv.data , indexFreqAnalysis= 1,  
+ size = 300, shift = 30 )
```

4.2.3.2 Wavelets

When using wavelet analysis with the *CalculatePowerBand* function, the user may specify:

4.2. ANALYZING THE HEART RATE SERIES

- *Wavelet*: mother wavelet used to calculate the spectrogram. Some of the most widely used wavelets are available: Haar (“haar”), extremal phase (“d4”, “d6”, “d8” and “d16”) and the least asymmetric Daubechies (“la8”, “la16” and “la20”) and the best localized Daubechies (“bl14” and “bl20”) wavelets among others. The default value is “d4”. The name of the wavelet specifies the “family” (the family determines the shape of the wavelet and its properties) and the length of the wavelet. For example, “la8” belongs to the Least Asymmetric family and has a length of 8 samples. We may give a simple advice for wavelet selection based on the wavelet’s length: shorter wavelets usually have better temporal resolution, but worse frequency resolution. On the other hand, longer wavelets usually have worse temporal resolution, but they provide better frequency resolution. Better temporal resolution means that we can study shorter time intervals. On the other hand, a better frequency resolution means better “frequency discrimination”. That is, shorter wavelets will tend to fail when discriminating close frequencies.
- *Bandtolerance*: maximum error allowed when the wavelet-based analysis is performed [12], [13]. It can be specified as either an absolute or a relative error depending on the “*relative*” parameter value. Default value is 0.01.
- *Relative*: logic value specifying which type of band tolerance shall be used: relative (in percentage) or absolute (default value).

Let $[f_l, f_u]$ be any frequency band specified by the user and let $[f_1, f_2]$ be a frequency interval associated with some node in the

4.2. ANALYZING THE HEART RATE SERIES

Maximal Overlap Discrete Wavelet Packet Transform (MODWPT) tree [29].

The relative error ϵ_r of f_l over the $[f_1, f_2]$ interval is computed as

$$\epsilon_r = \left| \frac{f_l - f_1}{f_u - f_l} \right| \cdot 100\%.$$

Similarly, we may define the error ϵ_r of the upper frequency f_u as

$$\epsilon_r = \left| \frac{f_u - f_2}{f_u - f_l} \right| \cdot 100\%.$$

The relative error can be used to avoid introducing large errors at small frequency bands (usually both ULF and VLF bands).

The absolute value ϵ is defined as usual: $\epsilon = |f_2 - f_u|$ for the upper frequency and $\epsilon = |f_1 - f_l|$ for the lower frequency.

Let's analyze the same frequency bands as before but using the wavelet-algorithm. For the sake of simplicity, we will use an absolute tolerance of 0.01 Hz . We may select the least asymmetric Daubechies of width 8 ("la8") as wavelet, since it provides a good compromise between frequency and time resolution. Thus, we may write:

```
> hrv.data = CreateHRVData( )
> hrv.data = SetVerbose(hrv.data, FALSE)
> hrv.data = LoadBeatAscii(hrv.data, "example.beats", "beatsFolder")
> hrv.data = BuildNIHR(hrv.data)
> hrv.data = FilterNIHR(hrv.data)
```

4.2. ANALYZING THE HEART RATE SERIES

```
> hrv.data = InterpolateNIHR (hrv.data, freqhr = 4)

> hrv.data = CreateFreqAnalysis(hrv.data)

> hrv.data = SetVerbose(hrv.data,TRUE)

> # Note that it is not necessary to write the boundaries

> # for the frequency bands, since they match the default values

> hrv.data = CalculatePowerBand( hrv.data , indexFreqAnalysis= 1,

+ type = "wavelet", wavelet = "la8", bandtolerance = 0.01, relative = FALSE,

+ ULFmin = 0, ULFmax = 0.03, VLFmin = 0.03, VLFmax = 0.05,

+ LFmin = 0.05, LFmax = 0.15, HFmin = 0.15, HFmax = 0.4 )



** Calculating power per band **

** Using Wavelet analysis **

Power per band calculated
```

4.2.3.3 Creating several analyses

In the previous examples we have used just one frequency analysis to illustrate the basic use of *CalculatePowerBand*. However, it is possible to create and use the same *HRVData* for performing several spectral analysis. When we do this, we use the parameter “indexFreqAnalysis” to indicate which spectral analysis we are working with. For example, we could perform both Fourier and wavelet based analysis:

```
> # ...

> # create structure, load beats, filter and interpolate

> hrv.data = CreateFreqAnalysis(hrv.data)

> hrv.data = SetVerbose(hrv.data,TRUE)
```

4.2. ANALYZING THE HEART RATE SERIES

```
> # use freqAnalysis number 1 for performing
> # Fourier analysis. This time, we do not
> # write the band's boundaries
> hrv.data = CalculatePowerBand( hrv.data , indexFreqAnalysis= 1,
+ size = 300, shift = 30, sizesp = 2048, type = "fourier")

** Calculating power per band **

** Using Fourier analysis **

Windowing signal... 237 windows

Power per band calculated

> # use freqAnalysis number 2 for performing
> # wavelet analysis. Note the indexFreqAnalysis = 2!!!
> hrv.data = CreateFreqAnalysis(hrv.data)

** Creating frequency analysis

Data has now 2 frequency analysis

> hrv.data = CalculatePowerBand( hrv.data , indexFreqAnalysis= 2,
+ type = "wavelet", wavelet = "la8", bandtolerance = 0.01, relative = FALSE)

** Calculating power per band **

** Using Wavelet analysis **

Power per band calculated
```

4.2.3.4 Plotting

RHRV also includes plotting utilities for representing the spectrogram of each frequency band: the *PlotPowerBand* function. The *PlotPowerBand* receives as inputs the *HRVData* structure and the index of the frequency analysis that the user wants to plot (*indexFreqAnalysis* argument). Optionally, the user can specify additional parameters for modifying the plots (whether to use or not normalized plots, specify the y-axis, etc.). For the sake of simplicity we will only use the *ymax* parameter (for specifying the maximum y-axis of the power bands plot) and the *ymaxratio* parameter (for specifying the maximum y-axis in the *LF/HF* plot).

If we want to plot the power bands computed in the previous example, we may use:

```
> # Plotting Fourier analysis  
> PlotPowerBand(hrv.data, indexFreqAnalysis = 1, ymax = 200, ymaxratio = 1.7)  
  
> # Plotting wavelet analysis  
> PlotPowerBand(hrv.data, indexFreqAnalysis = 2, ymax = 700, ymaxratio = 50)
```

The plots obtained with *PlotPowerBand* are shown in Figures 4.4 and 4.5, respectively.

4.2. ANALYZING THE HEART RATE SERIES

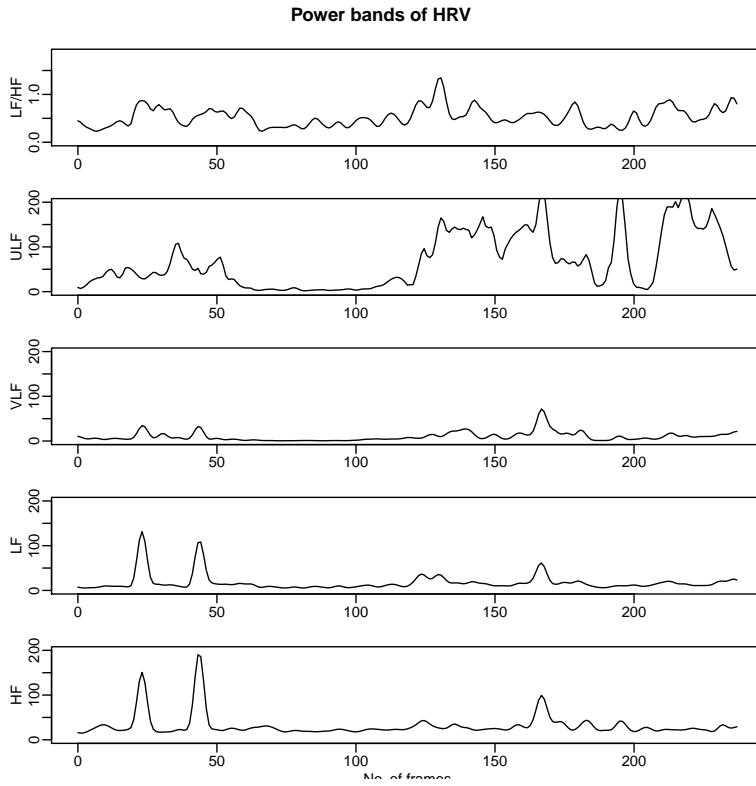


Figure 4.4: Plot obtained with the *PlotPowerBand* for the Fourier-based analysis.

4.2.3.5 A brief comparison: wavelets Vs. Fourier

Figures 4.4 and 4.5 illustrate some of the most important differences between Fourier and wavelet-based analysis. The most important differences may be summarized as follows:

- The power range is not the same when using Fourier than when using wavelets due to the windowing used in both techniques. Thus, we should avoid direct comparisons between the numerical results obtained with Fourier with those obtained using wavelets.

4.2. ANALYZING THE HEART RATE SERIES

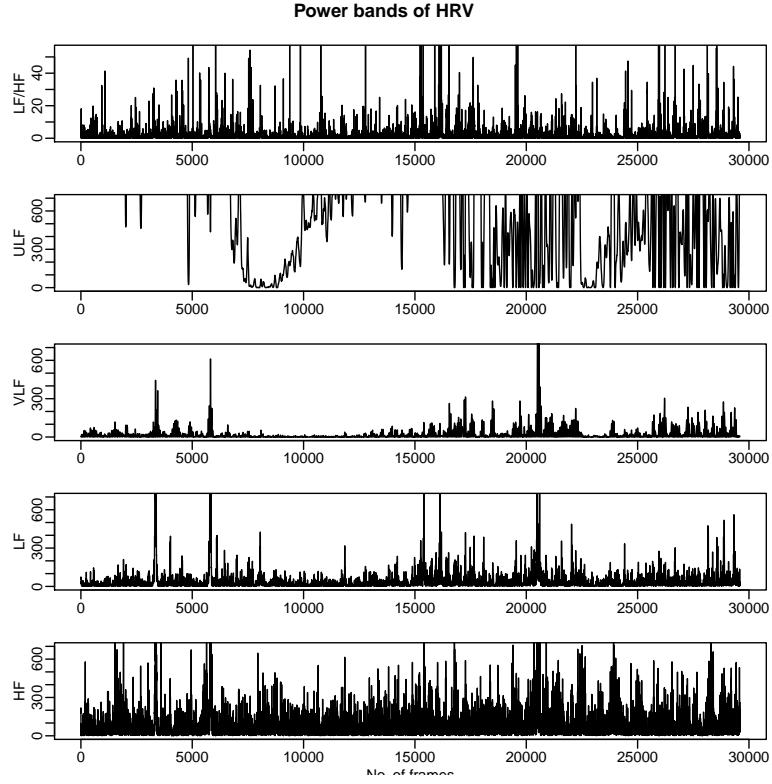


Figure 4.5: Plot obtained with the *PlotPowerBand* for the wavelet-based analysis.

- The Fourier's power spectrum is smoother than the wavelet's power spectrum. This is a consequence of the higher temporal resolution that the wavelet-based analysis provides. We could try to increase Fourier's frequency resolution by decreasing the window's size used in the analysis. The shorter window we use, the sharper spectrum we get. Similarly, we can increase/decrease temporal resolution using shorter/larger wavelets when performing wavelet-based analysis.
- The power spectrum obtained from the Fourier-based analysis has a smaller number of samples than the original signal as a consequence of the use of

4.2. ANALYZING THE HEART RATE SERIES

windows. Conversely, the power spectrum obtained from the wavelet-based analysis has the same number of samples as the original *RR* time series.

Chapter 5

Some more advanced features of RHRV

5.1 Completing our first tour

In chapter 4 we have presented a brief description of the RHRV package. In this section, we introduce some more advanced functionality of RHRV, or functionality that has narrower applications than the one presented in the previous chapter. Thus, we will introduce some new functions (*EditNIHR*, *CalculateSpectrogram* and *PlotSpectrogram*) and we will finish the description of all the function parameters introduced in the previous chapter. Also, further information about the *HRVData* structure will be given. Figure 5.1 shows a detailed view of the internal organization of the *HRVData* structure. This figure should be used as a roadmap through the explanations concerning the *HRVData* structure.

5.1. COMPLETING OUR FIRST TOUR

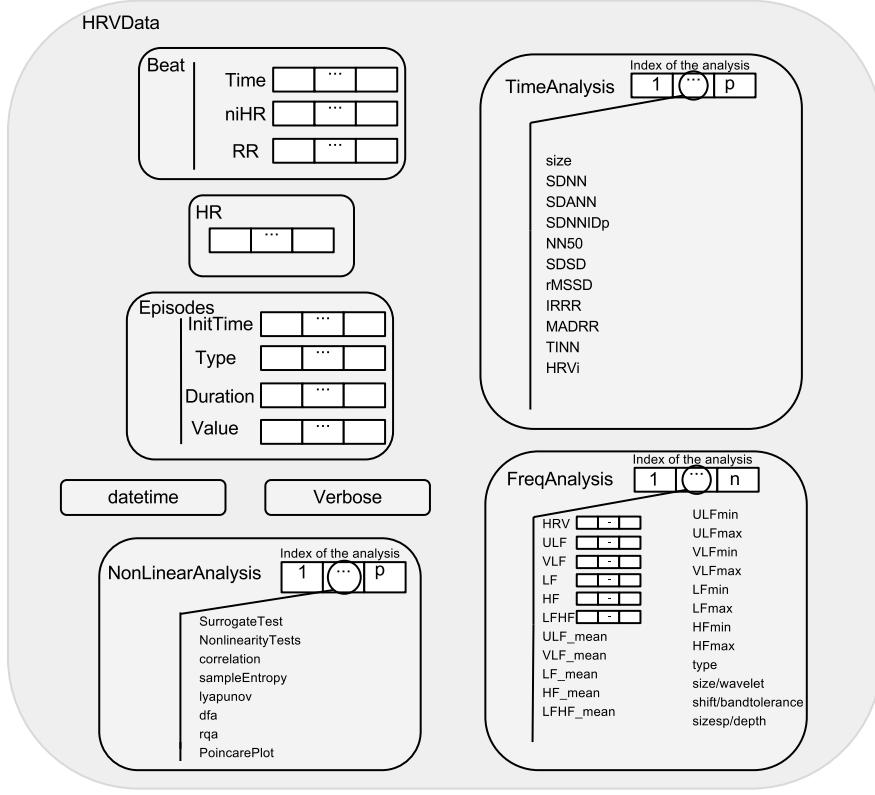


Figure 5.1: All the fields stored in the `HRVData` structure.

5.1.1 Creating the structure

When a new `HRVData` is created using the `CreateHRVData` function, it contains the following fields (among others that are not useful for the final user):

- `Ext`: A string that will be used as file extension by the loading/writing functions included in the package. The default value is “`hrv`”.
- `TimeAnalysis`: This field stores the information generated using time-domain analysis techniques. It is implemented as a list in the R language.

- *FreqAnalysis*: This field stores the results of one or more frequency analysis. Frequency analysis can be based on Fourier or on wavelets. It is implemented as a list in the R language.
- *NonLinearAnalysis*: This field stores the results of one or more nonlinear analysis. It is implemented as a list in the R language.
- *Verbose*: Boolean flag that specifies if all the functions should return additional information (mode verbose on). The *SetVerbose* function sets verbose mode on or off.

5.1.2 Reading heart beats

After creating the empty *HRVData* structure we will usually read the corresponding data file containing the heart beat positions. For the sake of simplicity, we keep on focussing on the ASCII files. The reader may have been wondering: “what happens if my ASCII file is specified in milliseconds?”. We shall use the *scale* parameter to overcome this issue. By setting this parameter to 1 (default value) we are indicating that the beat positions are specified in seconds; by setting it to 0.1, we are indicating that the beats are in deciseconds and so on. The function transforms the heart beat positions to seconds so all the other functions can be used as before. Other interesting parameter that can be specified by the user is the date-time when the file was recorded (*datetime* parameter). This is particularly useful for following a patient’s evolution over a set of recordings. The string format for the *datetime* parameter is “DD/MM/YYYY HH:MM:SS”. Thus, let’s read the “example.beats”

5.1. COMPLETING OUR FIRST TOUR

file (as in chapter 4) specifying that it was recorded on “30/04/2012 12:00:00” in seconds:

```
> hrv.data = LoadBeatAscii(hrv.data, "example.beats",
+     RecordPath = "beatsFolder", scale = 1,
+     datetime = "30/04/2012 12:00:00")

** Loading beats positions for record: example.beats **

Path: beatsFolder

Scale: 1

Date: 30/04/2012

Time: 12:00:00

Number of beats: 17360
```

When importing the data into the *HRVData* structure, two new fields are created (see Figure 5.1):

- *Datetime*: Date and time associated with the record.
- *Beat*: A *dataframe* object which stores the positions of the beats in the sub-field *Time* .

5.1.3 Constructing the time series

To compute the HRV time series the *BuildNIHR* function is used. Since we have already used all the parameters of this function, we will focus on the *HRVData structure*. As we know, this function constructs both the RR (Equation 2.1) and

instantaneous heart rate series (Equation 2.2). Both series are stored in the *Beat* dataframe of the *HRVData* structure: the RR series is stored in the sub-field called *RR* whereas the instantaneous HR is stored in the *niHR* sub-field (see Figure 5.1).

5.1.4 Filtering the time series

The automatic removal of outliers is performed with the *FilterNIHR* function. This function implements an algorithm that uses adaptive thresholds for rejecting or accepting beats [33]. The rule for beat acceptance or rejection is to compare the present beat with the previous one, the following one and with an updated mean of the RR interval. The different adaptive thresholds establish an upper limit for the relative errors of each of these comparisons. The *long* parameter allows the user to specify the number of beats that shall be used to calculate the updated mean (default value is 50 heartbeats). Also, the *last* parameter permits the user to specify the initial threshold value in % (default value is 13%). Finally, the algorithm also applies a comparison with acceptable physiological values. The user can specify the range of acceptable physiological values by using the *minbpm* and *maxbpm* (minimum beats per minute and maximum beats per minute, respectively). Default values are designed for human beings(*minbpm*=25, *maxbpm*=200), but they can be specified in such a way that it may also be used by animal researchers. As an illustrative example we could modify the *last* parameter in such a way that it does not allow quick fluctuations (by decreasing *last* to 1%) in our example file. Also, we could decrease

5.1. COMPLETING OUR FIRST TOUR

the *maxbpm* parameter to 180 bpm. The results are shown in Figure 5.2 (compare it with Figure 4.1).

```
> hrv.data = FilterNIHR(hrv.data, long=50, last=1, minbpm=25, maxbpm=180)
> PlotNIHR(hrv.data)
```

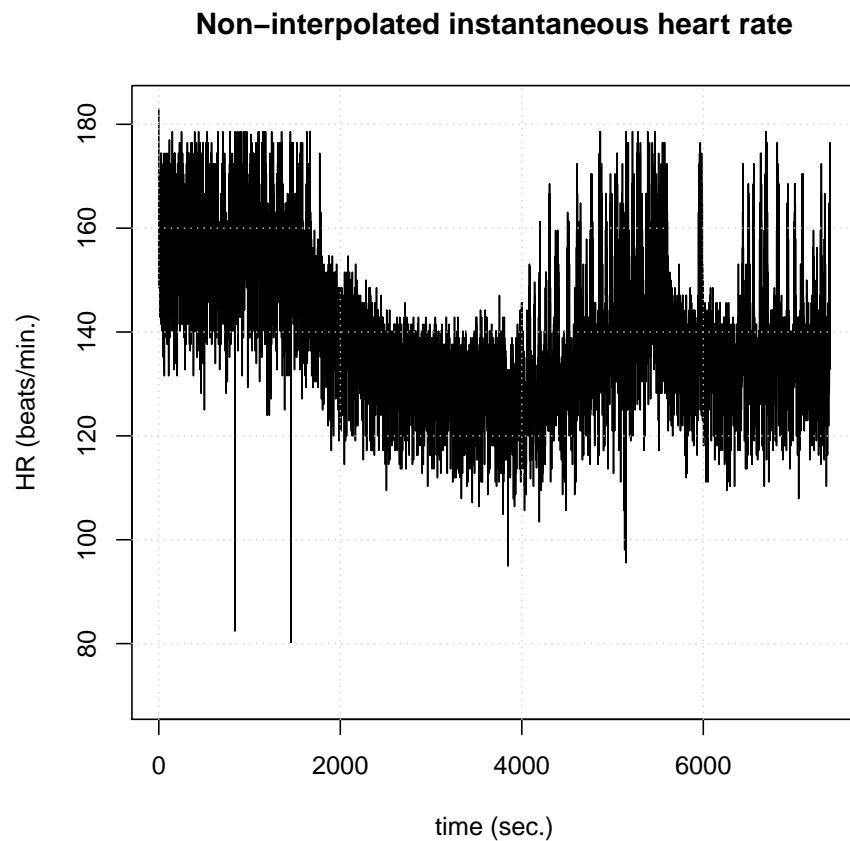


Figure 5.2: Effects of the modification of the default values in the *FilterNIHR* function.

RHRV also provides functionality for manually removing the spurious heartbeats. In order to delete outliers manually, a graphical editor can be used. The graphical editor is launched by executing the *EditNIHR* function.

5.1. COMPLETING OUR FIRST TOUR

```
> hrv.data = EditNIHR(hrv.data)
```

This interactive editor allows the user to select a rectangular area defined by two points that are the top left corner and bottom right corner, respectively, of a rectangle. (see Figure 5.3). The points included in this rectangle can then be removed by pressing the “remove outliers” button. If we make a mistake in the outliers selection, we can reset the window by pressing “clear”. The outliers removal ends when the user presses “End”.

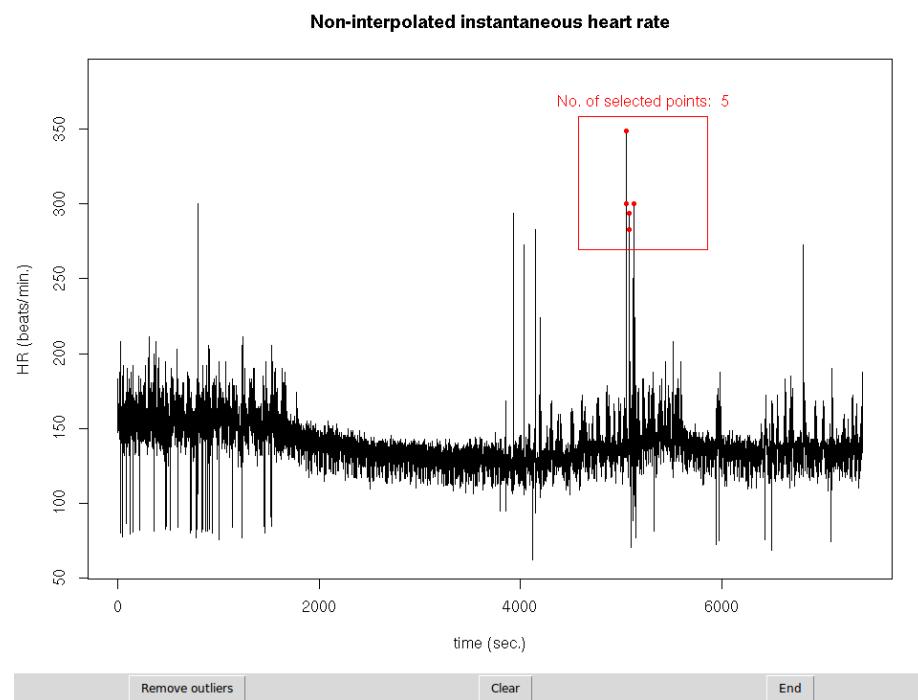


Figure 5.3: Manually removal of artifacts with *EditNIHR*.

5.1.5 Interpolation

The uniformly sampled HR series is obtained using the *InterpolateNIHR* function, which by default uses linear interpolation. However, it is possible to select a cubic spline interpolation by setting *method* = “*spline*” (default value is *method* = “*linear*”). Thus, as an illustrative example, let’s interpolate the RR data using splines and a sampling frequency of 8 Hz (This is just an illustrative example, for most of the situations 4 Hz will be enough. By setting an unnecessarily high sampling frequency, we are overloading the computer):

```
> hrv.data = InterpolateNIHR (hrv.data, freqhr = 8, method = "spline")
```

This function creates two new fields in the *HRVData* structure:

- *Freq_HR*: Sampling frequency used in the interpolation. The default sampling frequency value is 4 Hz.
- *HR*: Heart Rate signal with equally spaced values at a certain sampling frequency obtained from the niHR series (Figure 5.1).

5.1.6 Time analysis

The *CreateTimeAnalysis* function has been kept simple, in such a manner that the user only has to specify the window that will be used to compute successive differences of intervals (*size* parameter, in seconds) and the interval width of the bins that will be used to compute the histogram (*interval* parameter, in milliseconds), as shown in chapter 4.

This function fills (one position of) the *TimeAnalysis* in the *HRVData* structure by computing the following parameters: SDNN, SDANN, SDNNIDX, pNN50, rMSSD, IRRR, MADRR, TINN and HRVi (see Figure 5.1 and Section 2.2.1). The size of the window involved in the computations is also stored in the *size* field.

5.1.7 Frequency analysis

All the main *parameters* of the *CalculatePowerBand* function have already been used in chapter 4. As shown in Figure 5.1, the *CalculatePowerBand* function fills the corresponding *FreqAnalysis* data structure with the following fields:

- *Type*: a string identifying the type of analysis that has been used. The possible values are either “fourier” or “wavelet”.
- *ULFmin*, *ULFmax*, *VLFmin*, *VLFmax*, *LFmin*, *LFmax*, *HFmin* and *HFmax*: These fields store the boundaries of each frequency band.
- *ULF*, *VLF*, *LF* and *HF*: These fields store the spectrogram of the HR signal in the ULF, VLF, LF and HF bands, respectively.
- *HRV*: Stores the total energy of the signal as a function of time. This Energy time series is estimated from the spectrogram signal.
- *LFHF*: Stores the LF/HF ratio (Section 2.2) by dividing the LF time series by the HF time series.

Some additional parameters are incorporated to the *FreqAnalysis* structure depending on the type of analysis used. When using the STFT, the *size*, *shift* and *sizesp*

5.1. COMPLETING OUR FIRST TOUR

fields store information about the window, the window shift, and the number of points per DFT that have been used. When using the wavelet transform, the *wavelet*, *bandtolerance* and *depth* fields store information about the mother wavelet and the tolerance used, as well as the number of levels that the algorithm has descended in the MODWPT tree [12], [13].

RHRV provides another function for computing the spectrogram without being restricted to the four bands: ULF, VLF, LF and HF. This function, called *CalculateSpectrogram*, uses the STFT approach. Thus, the user has to specify the size of window (*size* parameter), the displacement of window (*shift*) and the zero-padding (*sizesp*), as in the *CalculatePowerBand* function. The spectrogram is returned in a real matrix in a way that, as the number of the row increases, the time increases and, as the column's number increases, the frequency increases. This matrix is not stored in the *HRVData* structure since it can be very expensive in terms of memory. As an example, let's compute the spectrogram of the example file with the same parameters used with the *CalculatePowerBand* function:

```
> # Plotting Fourier analysis  
> spectrogram = CalculateSpectrogram( hrv.data ,size = 300,  
+ shift = 30, sizesp = 2048)
```

The user can obtain a graphical representation of the spectrogram by using the *image* R function. Alternatively, the user can use the *PlotSpectrogram* function, that also returns the spectrogram matrix (see Figure 5.4). By using the *scale* function, the

5.1. COMPLETING OUR FIRST TOUR

user may choose a linear axis (“linear”) or a logarithmic axis (“logarithmic”). The user must also specify the *size*, *shift* and *sizesp* parameter.

```
> # Plotting wavelet analysis  
> spectrogram = PlotSpectrogram(HRVData=hrv.data, size=300, shift=60,  
+ sizesp=2048)
```

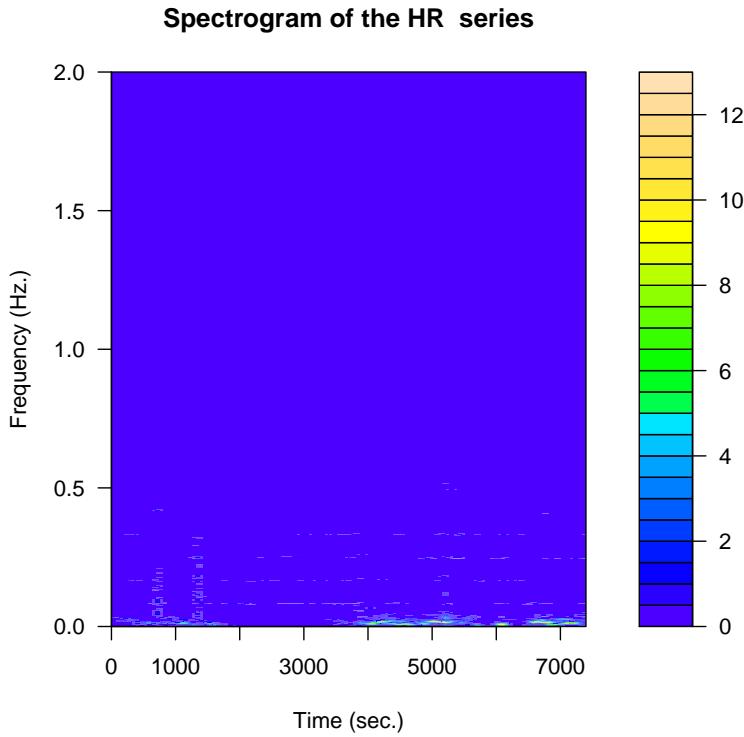


Figure 5.4: Plot obtained with the *PlotSpectrogram* function.

Note that most of the energy shown in Figure 5.4 is concentrated in the low frequencies. To obtain a more detailed graphic in this zone of the spectrum, we can use the

5.1. COMPLETING OUR FIRST TOUR

freqRange parameter. For example, if we wish to plot the spectrum in the $[0, 0.2]$ Hz band, we write:

```
> # Plotting wavelet analysis  
> PlotSpectrogram(HRVData=hrv.data, size=300, shift=60,  
+ sizesp=2048, freqRange = c(0,0.2))
```

The result of this script is shown in Figure 5.5.

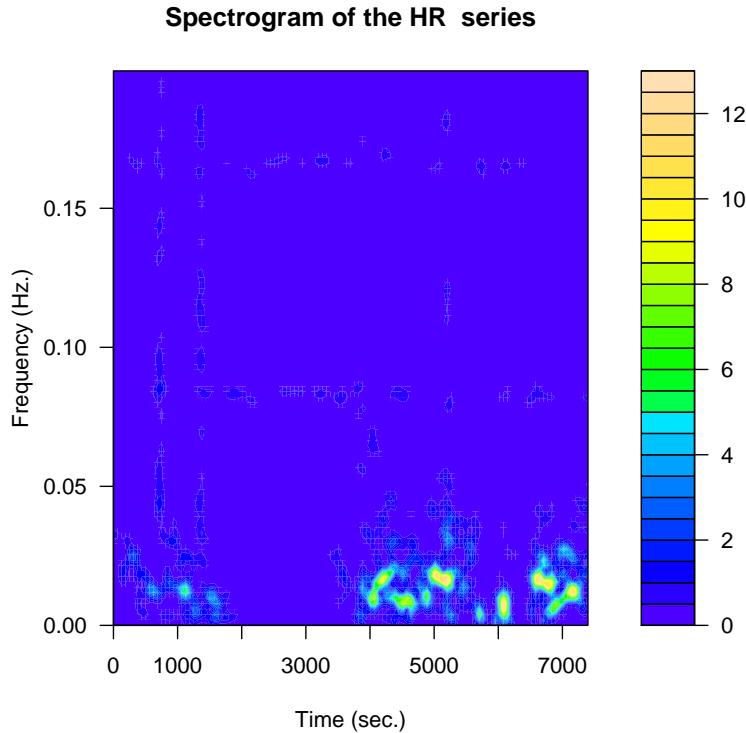


Figure 5.5: Plot obtained with the *PlotSpectrogram* function and the *freqRange* parameter.

5.2 Reading several file formats

RHRV provides a lot of functionality for importing data files containing heart beat positions. Supported formats include ASCII (*LoadBeatAscii* function), EDF (*LoadBeatEDFPlus*), Polar (*LoadBeatPolar*), Suunto (*LoadBeatSuunto*) and WFDB data files (*LoadBeatWFDB*) [26]. We have already dealt with the *LoadBeatASCII* function. In this section, we will discuss the remaining functions for reading heart beat data.

5.2.1 Reading *RR* files

There exists another *RHRV* function that reads ASCII files: the *LoadBeatRR* function. This function reads ASCII files storing the *RR* intervals (and not the heart beat times). The parameters of the *LoadBeatRR* function are exactly the same as those of the *LoadBeatAscii* function, that is:

```
> LoadBeatRR(HRVData, RecordName, RecordPath=". ", scale = 1,  
+           datetime = "1/1/1900 0:0:0")
```

5.2.2 Reading files in WFDB format

PhysioNet [14] is a free web resource that provides large collections of recorded physiologic signals (PhysioBank) and related open-source software (PhysioToolkit). In most cases, a record from PhysioBank consists of at least three files, which are named using the record name followed by different extensions that indicate their content. Almost all records include a binary .dat file, containing digitized samples

5.2. READING SEVERAL FILE FORMATS

of one or more signals. The .hea (header) file is a short text file that describes the signals. Most records include one or more binary annotation files. For example, .qrs files contain an annotation for each QRS complex (heart beat) in the recording; .apn files contain apnea annotations; etc. For the sake of simplicity, we call “WFDB file” (WaveForm DataBase) to such a collection of files containing data on the same recording [26]. Further details about PhysioBank can be found in the PhysioNet website.

The RHRV package provides the *LoadBeatWFDB* function for reading WFDB files. This function takes as input parameters the name of the WFDB file to be used without any extension (*RecordName* argument), the relative path of the file (*RecordPath* argument) and the extension of the file with the heart beats annotations (*annotator*, its default value is “qrs”, so in most cases the user won’t have to specify it). As an example, we are going to create a data structure that will read the “a03” register from the PhysioBank’s Apnea-ECG database [28].

```
> hrv.wfdb = CreateHRVData()  
  
> hrv.wfdb = SetVerbose(hrv.wfdb, TRUE)  
  
> hrv.wfdb = LoadBeatWFDB(hrv.wfdb, "a03", RecordPath = ". ",  
+ annotator = "qrs")  
  
** Loading beats positions for record: a03 **  
  
Path: .  
  
Opening header file: a03.hea  
  
No time information in header: 00:00:00
```

```
No date information in header: 01/01/1900  
Date: 01/01/1900  
Time: 00:00:00  
Number of beats: 34254
```

5.2.3 Other formats

Since the remaining functions have a similar behaviour than those we have already discussed, we will just discuss their prototypes. The *LoadBeatEDFPlus* function allows the user to read EDF+ (European Data Format) data [19]. Its format is similar to the *LoadBeatWFDB* function:

```
> LoadBeatEDFPlus(HRVData, RecordName, RecordPath = ". ",  
+ annotationType ="QRS")
```

Finally, RHRV provides functionality for reading Polar and Suunto files with the *LoadBeatPolar* and *LoadBeatSuunto* functions. These functions only receive as arguments the record name and the record path:

```
> LoadBeatPolar(HRVData, RecordName, RecordPath=". ")  
> LoadBeatSuunto(HRVData, RecordName, RecordPath=". ")
```

5.2.4 A general function

The *LoadBeat* function provides a common interface to access all the functions responsible for loading heart beats. Thus, the prototype of the *LoadBeat* function contains all the *parameters* needed for the loading functions. It also offers the *fileType*

fileType	function called
“WFDB”	<i>LoadBeatWFDB</i>
“Ascii”	<i>LoadBeatAscii</i>
“RR”	<i>LoadBeatRR</i>
“Polar”	<i>LoadBeatPolar</i>
“Suunto”	<i>LoadBeatSuunto</i>
“EDFPlus”	<i>LoadBeatEDFPlus</i>

Table 5.1: *LoadBeat* operation depending on the *fileType* parameter.

parameter so the user can specify which type of file is going to be read. Depending on the *fileType* value, the *LoadBeat* function delegates on one of the previous loading functions. The possible values of the *fileType* parameter and the function that is called are summarized in table 5.1.

Thus, if we want to read a WFDB file, we could use either the *LoadBeatWFDB* function or the *LoadBeat* function:

```
> hrv.wfdb = CreateHRVData()  
  
> hrv.wfdb = SetVerbose(hrv.wfdb, TRUE)  
  
> hrv.wfdb = LoadBeat(hrv.wfdb, fileType = "WFDB", "a03", RecordPath = ".",  
+ annotator = "qrs")  
  
** Loading beats positions for record: a03 **  
  
Path: .  
  
Opening header file: a03.hea  
  
No time information in header: 00:00:00  
  
No date information in header: 01/01/1900  
  
Date: 01/01/1900
```

5.3. PERFORMING ANALYSIS IN DIFFERENT INTERVALS OF A RECORDING

Time: 00:00:00

Number of beats: 34254

5.3 Performing analysis in different intervals of a recording

Intervals of the HR time series with pathophysiological interest may be annotated in the so-called episode files. For example, it may be interesting to compare the heart rate series before, during and after an apnea episode (apneas are cessations of a patient’s respiratory airflow during the nocturnal rest) [21]. Such a study could be useful for searching for significant differences in the HRV caused by the apneas. The RHRV package provides functions for loading episode information. The supported formats for this information are either ASCII (*LoadEpisodesAscii*) or WFDB (*LoadApneaWFDB*). Episodes may also be added programmatically to the time series using the *AddEpisodes* function. All episodes are stored under the *Episodes* field of the *HRVData* structure (see Figure 5.1). The plotting functions allow the user to include episodic information in the graphics. We will discuss all these points in more detail below.

5.3.1 AddEpisodes

The simplest way of adding episodic information is using the *AddEpisodes* function. *AddEpisodes* adds information of episodes by specifying the initial times of each

5.3. PERFORMING ANALYSIS IN DIFFERENT INTERVALS OF A RECORDING

episode (*InitTimes* argument, in seconds), the names of the episodes (*Tags*), the duration of each episode (*Durations*, in seconds) and a numerical identifier for each episode(*Values*). The *Values* field is useful for those episodes that store some numerical values. For example, an apnea episode could store information about the Oxygen saturation level in the *Values* field. Note that all the parameters specified by the user will be stored in the *HRVData* structure in its corresponding fields as shown in Figure 5.1.

Let us read our example file “example.beats” and add three episodes to it: a first type “A” episode in the [700,1600] s interval; a second episode of the same type as the first one (“A”) in the [5000,5600] s interval; and a third episode in the [2000,4500] s interval of type “B”:

```
> hrv.data = CreateHRVData( )

> hrv.data = LoadBeatAscii(hrv.data, "example.beats", "beatsFolder")

> hrv.data = SetVerbose(hrv.data, TRUE)

> hrv.data = AddEpisodes(hrv.data, InitTimes = c(700,5000,2000),
+ Tags = c("A","A","B"), Durations = c(900,600,2500), Values = c(0,0,0))

** Adding new episodes **

Added 3 episodes from file

Number of episodes: 3
```

5.3. PERFORMING ANALYSIS IN DIFFERENT INTERVALS OF A RECORDING

5.3.2 Plotting episodic information

The *plotHR* and *PlotNIHR* functions allow the user to include episodic information in the plot. The user can specify a list of tags to specify which episodes are included in the plot (*Tag* parameter). *Tag*=“all” plots all episodes present in the data. Thus we could execute, after the code of the previous paragraph:

```
> hrv.data = BuildNIHR(hrv.data)

> # plot all tags

> PlotNIHR(hrv.data, Tag="all")

> hrv.data = InterpolateNIHR(hrv.data, freqhr = 4)

> # Plot only the "A" episodic information

> PlotHR(hrv.data , Tag=c("A"))
```

The plots obtained with *PlotNIHR* and *PlotHR* are shown in Figures 5.6 and 5.7, respectively.

RHRV is also capable of including episodic information when representing the spectrograms obtained with the *CalculatePowerBand*. For this purpose, the *PlotPowerBand* includes the *Tag* input parameter. Thus, if we want to perform a frequency analysis and plot the power bands with the episodic information that we have added in the previous paragraphs, we could execute:

```
> hrv.data = CreateFreqAnalysis(hrv.data)

> # perform frequency analysis
```

5.3. PERFORMING ANALYSIS IN DIFFERENT INTERVALS OF A RECORDING

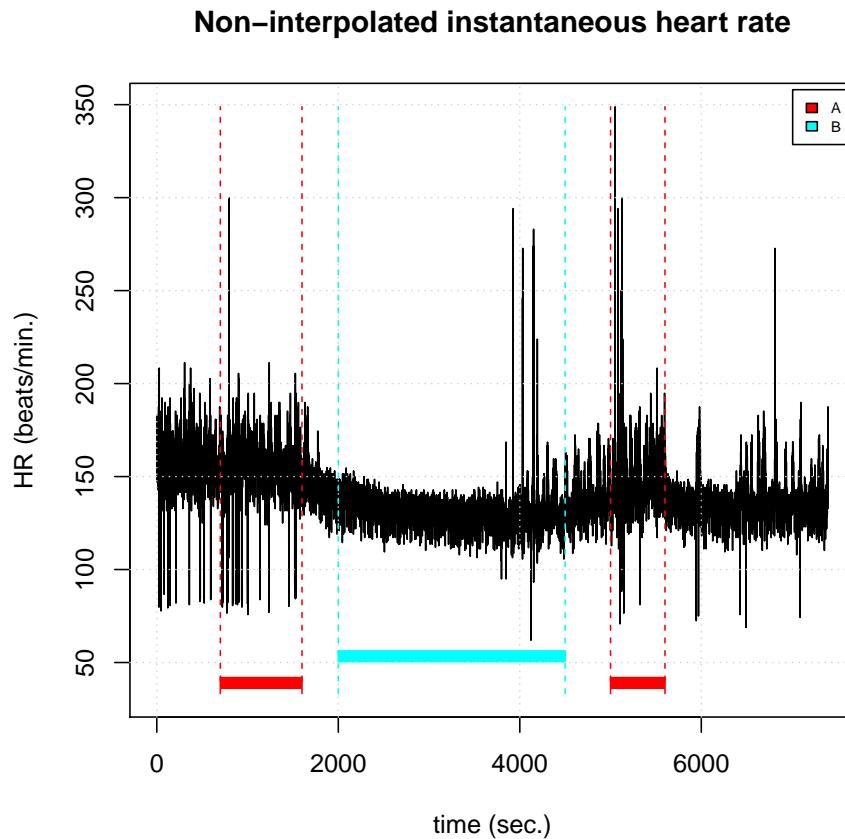


Figure 5.6: Episodic information in the Non interpolated Heart Rate time series.

```
> hrv.data = CalculatePowerBand( hrv.data , indexFreqAnalysis= 1,  
+   type = "wavelet", wavelet = "la8", bandtolerance = 0.01, relative = FALSE)  
> # plot episodic information  
> PlotPowerBand(hrv.data, indexFreqAnalysis = 1, ymax = 5000, ymaxratio = 50,  
+                           Tag = "all")
```

The resulting plot is shown in Figure 5.8.

5.3. PERFORMING ANALYSIS IN DIFFERENT INTERVALS OF A RECORDING

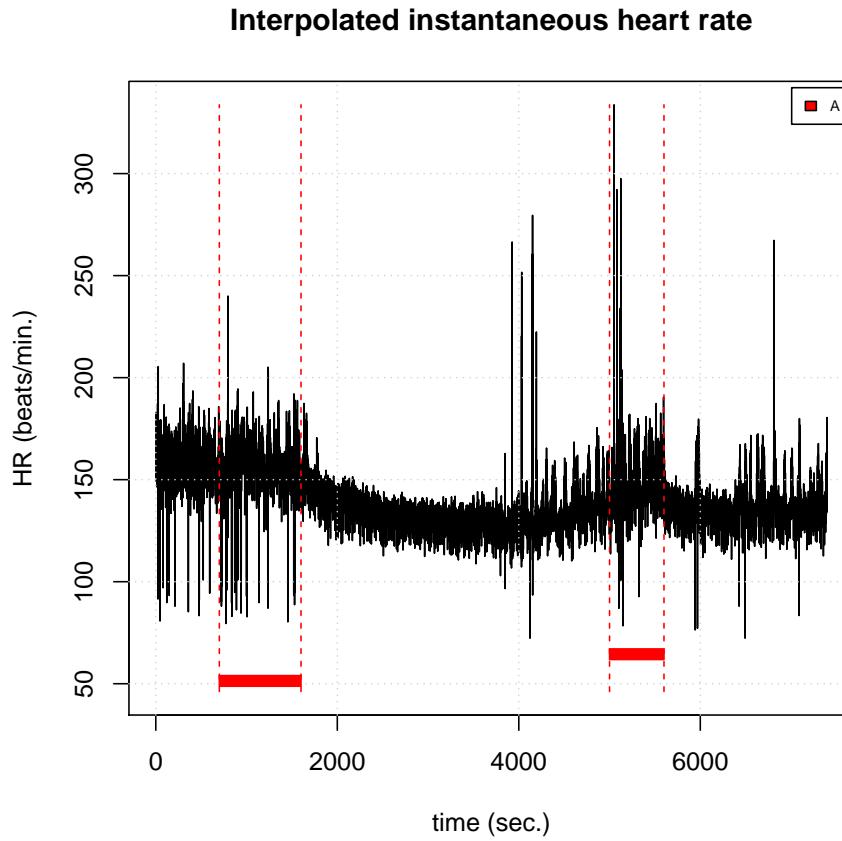


Figure 5.7: Episodic information in the interpolated Heart Rate time series.

5.3.3 LoadEpisodesAscii

The *LoadEpisodesAscii* function allows the user to read episodic information stored in an ASCII file saving it into the *HRVData* data structure. The expected format of each line is:

InitTime	Tag	Duration	Value
HH:MM:SS	“Tag name”	double	integer

5.3. PERFORMING ANALYSIS IN DIFFERENT INTERVALS OF A RECORDING

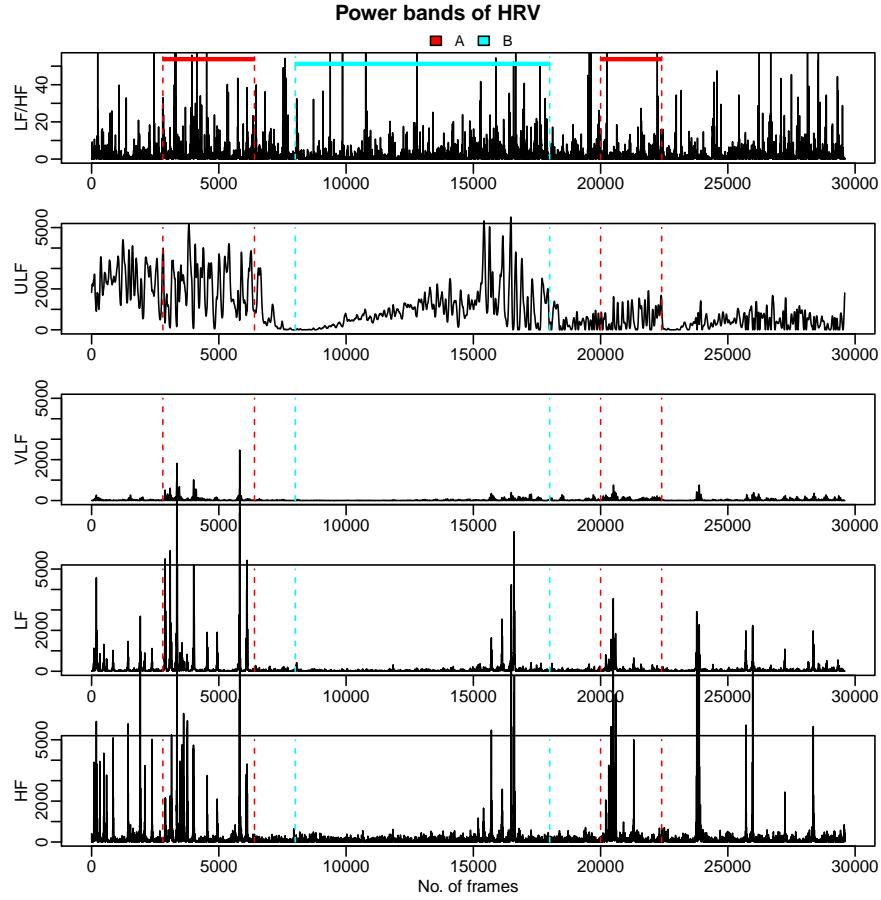


Figure 5.8: *Episodic information in all the power bands.*

The first column is the start time of the episode with the “HH:MM:SS” format. The second column serves the same purpose as the *Tag* parameter in *AddEpisodes*. The third column specifies the duration of the episode in seconds. Finally, the fourth column assigns a numerical value for each episode (the same function performed by the *Value* parameter in the *AddEpisodes* function).

5.3. PERFORMING ANALYSIS IN DIFFERENT INTERVALS OF A RECORDING

It must be taken into account that the *LoadEpisodesAscii* function skips the first line of the ASCII file because it assumes that the first line will contain a header (the format of this header does not really matter). If there is no header line in the file, the user can specify it to the function using the *header* parameter (by default, it is setted to TRUE).

As an example, we are going to create an ASCII file containing the same information as the episodes “A” and “B” that we programmatically introduced in the *AddEpisodes*(5.3.1) paragraph. We will also add information about a “C” episode that the *LoadEpisodesAscii* function will skip. We shall call this file “annotationsFile.txt”. This file will have the following content:

	InitTime	Type	Duration	Value
	00:11:40	”A”	900	1
	01:23:20	“A”	600	2
	00:33:20	”B”	2500	3
	01:00:00	”C”	100	4

The *LoadEpisodesAscii* function takes as input parameters: the absolute path to the episodes file to be read (*FileName*), the types of episodes that should be read (*Tag*) and the time (“HH:MM:SS”) at which the recording began (*InitTime*). This last parameter enables reading those files in which the initial time of episodes was specified in absolute time, and not relative to the start of the recording. Since we wrote relative times in the ASCII file, we should use *InitTime*=“0:0:0”. Thus, in order to read just the episodes tagged as “A” from our file, we could write:

5.3. PERFORMING ANALYSIS IN DIFFERENT INTERVALS OF A RECORDING

```
> hrv.data = CreateHRVData( )

> hrv.data = LoadBeatAscii(hrv.data, "example.beats", "beatsFolder")

> hrv.data = SetVerbose(hrv.data, TRUE)

> hrv.data = LoadEpisodesAscii(hrv.data, Tag=c("A"), InitTime="0:0:0",
+ FileName="beatsFolder/annotationsFile.txt")

** Loading episodes file: beatsFolder/annotationsFile.txt **

Path: .

Tag: A

Initial time: 00:00:0.000

Data includes values associated to episodes

Loaded 2 episodes from file

Number of episodes: 2
```

5.3.4 LoadApneaWFDB

The *LoadApneaWFDB* function allows the user to load apnea annotations from a WFDB file (the user must ensure that there is a file with the .apn extension between the WFDB recording's files). The function takes as input parameters the name of the WFDB file (*RecordName*), the path to the WFDB file (*RecordPath*) and a name for the apnea episodes (*Tag*, its default value is “APNEA”). The use of this function requires the installation of the WFDB tools (see chapter 3).

As an illustrative example, we are going to read the apnea episodes for the “a03” file of the ApneaECG database from PhysioBank (see the *Reading files in WFDB format* paragraph). The plot of the non interpolated HR series is shown in Figure 5.9.

5.3. PERFORMING ANALYSIS IN DIFFERENT INTERVALS OF A RECORDING

```
> hrv.wfdb = CreateHRVData()

> hrv.wfdb = SetVerbose(hrv.wfdb, TRUE)

> hrv.wfdb = LoadBeat(hrv.wfdb, fileType = "WFDB", "a03", RecordPath = ".",
+                      annotator = "qrs")

** Loading beats positions for record: a03 **

Path: .

Opening header file: a03.hea

No time information in header: 00:00:00

No date information in header: 01/01/1900

Date: 01/01/1900

Time: 00:00:00

Number of beats: 34254

> hrv.wfdb = LoadApneaWFDB(hrv.wfdb, RecordName="a03", Tag="Apnea",
+ RecordPath=". ")

** Loading apnea episodes for record: a03 **

Path: .

Header info already present for: a03

Command: rdann -r a03 -a apn

Number of labels: 518

** Adding new episodes **

Added 11 episodes from file

Number of episodes: 11
```

5.3. PERFORMING ANALYSIS IN DIFFERENT INTERVALS OF A RECORDING

```
> hrv.wfdb = BuildNIHR(hrv.wfdb)

** Calculating non-interpolated heart rate **

Number of beats: 34254

> PlotNIHR(hrv.wfdb, Tag="all")

** Plotting non-interpolated instantaneous heart rate **

Number of points: 34254

Episodes in plot: Apnea

No of episodes: 11

No of classes of episodes: 1
```

5.3.5 Analyzing HRV inside and outside the episodes

RHRV provides basic functionality for comparing data inside and outside each episode. The simplest function that the user can use for this purpose is the *SplitHRbyEpisodes* function, that splits the interpolated heart rate series into two vectors containing samples inside (the “InEpisodes” vector) and outside the episode (the “OutEpisodes” vector) specified in the *Tag* argument. For example, if we want to compare the HR series inside and outside the apnea episodes from the “a03” file, we could use:

```
> # remember to interpolate the Heart Rate series!!
> hrv.wfdb = InterpolateNIHR (hrv.wfdb, freqhr = 4)
```

5.3. PERFORMING ANALYSIS IN DIFFERENT INTERVALS OF A RECORDING

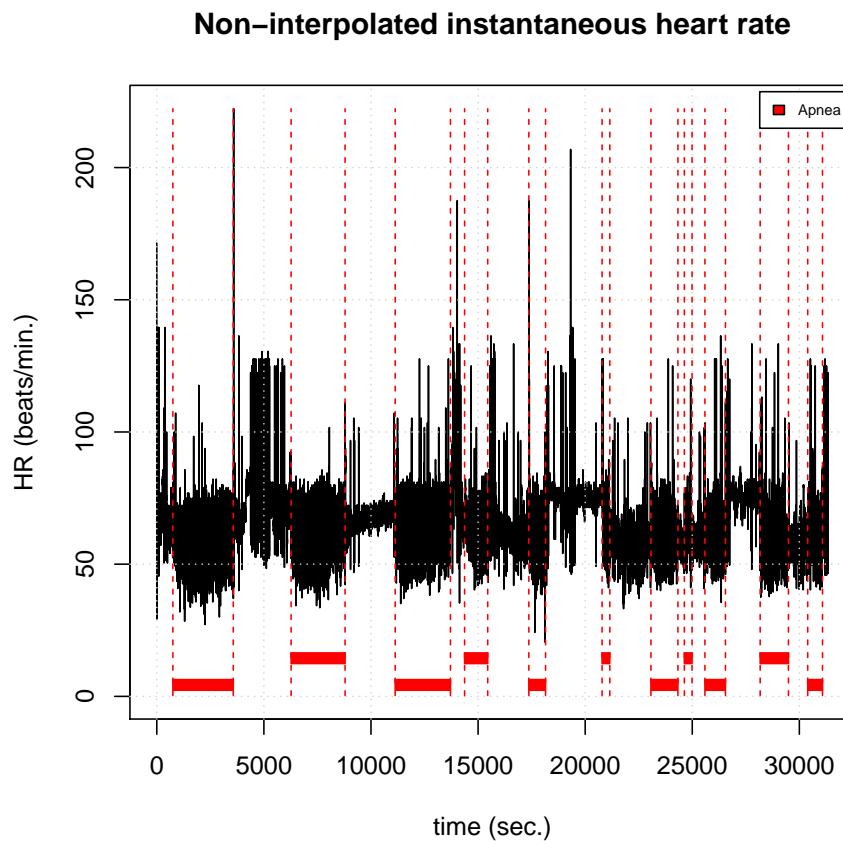


Figure 5.9: Loading Apnea episodes using the *LoadApneaWFDB* function.

** Interpolating instantaneous heart rate **

Frequency: 4Hz

Number of beats: 34254

Number of points: 125383

```
> splitting.data = SplitHRbyEpisodes(hrv.wfdb, Tag = c("Apnea"))
```

** Splitting heart rate signal using episodes **

Using episodes with tag: Apnea

5.3. PERFORMING ANALYSIS IN DIFFERENT INTERVALS OF A RECORDING

```
Number of episodes: 11  
Inside episodes: 58919 points  
Outside episodes: 66464 points
```

It is straightforward to use a statistical function for comparing both vectors. For example:

```
> cat("comparing the mean inside and outside Apnea episodes...\\n")
```

```
comparing the mean inside and outside Apnea episodes...
```

```
> cat("Apnea mean: ",mean(splitting.data$InEpisodes), "\\n")
```

```
Apnea mean: 61.94101
```

```
> cat("Normal mean: ",mean(splitting.data$OutEpisodes), "\\n")
```

```
Normal mean: 69.11033
```

Although the previous example illustrates how to access the data inside and outside a certain type of episode, we could have used the *AnalyzeHRbyEpisodes* for comparing its means. This function analyzes the heart rate series evaluating the desired function (*func* parameter) inside and outside the episodes of interest (*Tag* parameter). Thus, we could have executed:

```
> cat("comparing the mean inside and outside the Apnea episodes...\\n")
```

```
comparing the mean inside and outside the Apnea episodes...
```

5.3. PERFORMING ANALYSIS IN DIFFERENT INTERVALS OF A RECORDING

```
> result = AnalyzeHRbyEpisodes(hrv.wfdb, Tag ="Apnea", "mean")  
  
** Applying function to heart rate signal using episodic information **  
Function: "mean"()  
Using episodes with tag: Apnea  
  
** Splitting heart rate signal using episodes **  
Using episodes with tag: Apnea  
Number of episodes: 11  
Inside episodes: 58919 points  
Outside episodes: 66464 points  
  
> cat("Apnea mean:: ",result$resultIn,"\n")  
  
Apnea mean:: 61.94101  
  
> cat("Normal mean: ",result$resultOut,"\n")  
  
Normal mean: 69.11033
```

There also exists a splitting function that separates the spectral power per band in two lists using a specific episode type: the *SplitPowerBandByEpisodes* (however there is no analogue function to the *AnalyzeHRbyEpisodes* function). In addition to the *Tag* parameter, the *SplitPowerBandByEpisodes* receives as input parameters, the *HRVData* structure (*HRVData*) and the frequency analysis index to which apply the splitting function (*indexFreqAnalysis*). The function returns a list with two lists: “InEpisodes” and “OutEpisodes”, both lists include the ULF, VLF, LF and HF bands:

5.3. PERFORMING ANALYSIS IN DIFFERENT INTERVALS OF A RECORDING

```
> # ...

> hrv.wfdb = CreateFreqAnalysis(hrv.wfdb)

** Creating frequency analysis

Data has now 1 frequency analysis

> hrv.wfdb = CalculatePowerBand( hrv.wfdb , indexFreqAnalysis= 1,
+   type = "wavelet", wavelet = "la8", bandtolerance = 0.01, relative = FALSE)

** Calculating power per band **

** Using Wavelet analysis **

Power per band calculated

> splitting.data = SplitPowerBandByEpisodes(hrv.wfdb,
+ indexFreqAnalysis = 1, Tag = c("Apnea"))

** Splitting power bands using episodes**

Using episodes with tag: Apnea

Number of episodes: 11

No. of frames: 125383

No. of frames in episodes: 58931

No. of frames outside episodes: 66452

> cat("comparing the mean power in the LF band
+   inside and outside A episodes...\n")

comparing the mean power in the LF band

inside and outside A episodes...
```

```
> cat("LF power in Apnea episodes: ",  
+ mean(splitting.data$InEpisodes$LF), "\n")  
  
LF power in Apnea episodes: 3194.988  
  
> cat("LF power in Normal episodes: ",  
+ mean(splitting.data$OutEpisodes$LF), "\n")  
  
LF power in Normal episodes: 811.5181
```

5.4 Storing and reading HRVData

In order to save interesting results RHRV provides functionality for storing and reading *HRVData* structures. For example, if the user wants to store the *hrv.wfdb* structure from the previous section, he just has to write:

```
> WriteToFile(hrv.wfdb, name="HRVstructure")  
  
** Writing file: HRVstructure.hrv  
  
File HRVstructure.hrv already exists  
18468167 bytes written
```

The *WriteToFile* function will store the *HRVData* structure in a file called “HRVstructure.hrv”. Note that the “.hrv” suffix has been added. Additionally, the user may specify the behaviour of the function if the file already exists with the *overwrite* parameter. The default value overwrites existing files. If the user wants to prevent losing previous data stored in the “HRVstructure.hrv” file, he can write:

5.4. STORING AND READING HRVDATA

```
> WriteToFile(hrv.wfdb, name="HRVstructure", overwrite = FALSE)

Error in WriteToFile(hrv.data, name = "HRVstructure", overwrite = FALSE) :
  --- File exists... Not overwriting it!! ---
  --- Quitting now!! ---
```

Note that the function informs about the existence of a previous file named “HRVstructure.hrv”.

In order to read *HRVData* structures that had been previously stored, the *ReadFromFile* function is provided:

```
> data = ReadFromFile(name = "HRVstructure", verbose = TRUE)

** Reading file: HRVstructure.hrv

18468167 bytes read
```

Note that the “.hrv” prefix is not included in the file’s name, although the *HRVData* structure was stored as “HRVstructure.hrv”. The user can control the verbosity level using the *verbose* argument.

Chapter 6

HRV nonlinear analysis

6.1 An introduction to nonlinear analysis techniques

RHRV implements all the nonlinear statistics presented in Section 2.2.3. As we shall see in the next sections, most of the nonlinear statistics share a common name convention for their functions: the *CalculateX*, the *EstimateX* and the *PlotX* functions, being X the name of the statistic (for example *corrDim*). This naming convention reflects that these statistics share a similar estimation process: first, the *CalculateX* function performs some heavy computations that are required for obtaining the value of the statistic. The statistic is then obtained as the slope of a regression involving these computations. The regression is performed with the *EstimateX* function. The *PlotX* function draws the variables involved in the re-

6.1. AN INTRODUCTION TO NONLINEAR ANALYSIS TECHNIQUES

gression, which may be useful to check if the statistics follow the expected behaviour.

In order to load the example file that we shall use to illustrate the nonlinear functions included in RHRV, we may execute:

```
> library(RHRV)  
  
> hrv.data = CreateHRVData()  
  
> hrv.data = LoadBeatAscii(hrv.data, RecordName="nonlinearHB.beats",  
+                               RecordPath="beatsFolder")  
  
> hrv.data = BuildNIHR(hrv.data)  
  
> hrv.data = SetVerbose(hrv.data, TRUE)
```

The beats from this file were generated synthetically from a nonlinear process in order to show the expected behaviour of a truly nonlinear RR time series and avoid some issues when dealing with a noisy process. Thus, this file has no physiological meaning. An example of nonlinear HRV analysis with a real RR series is given in Section 6.2.

Before performing the nonlinear analysis we must create the data analysis structure that will store all the results from the analysis: the *NonLinearAnalysis* list, under the *HRVData* structure. Just as with the others analysis structures, each nonlinear analysis structure is identified by a unique number. To create an analysis structure, the *CreateNonLinearAnalysis* function is used.

```
> hrv.data = CreateNonLinearAnalysis(hrv.data)
```

```
** Creating non linear analysis
```

```
  Data has now 1 nonlinear analysis
```

In the next few sections we shall present the basics for performing the most widely-used nonlinear analysis algorithms (Although some of the algorithms will not be presented here, but in Section 6.2). Due to the heterogeneity of the nonlinear algorithms, we will also explain how to access the computations for each algorithm as they are presented.

6.1.1 Nonlinearity Test

Before applying nonlinear analysis to the RR time series we should ensure that the HR shows, indeed, some degree of nonlinearity. If we don't , there is a risk of obtaining unreliable results. There exist two functions in *RHRV* that allow running nonlinearity tests: *NonlinearityTests* and *SurrogateTest*.

The *NonlinearityTests* permits running a wide variety of nonlinearity tests including: two tests for neglected nonlinearity that are based on neural networks, the Keenan test for nonlinearity, the McLeod and Li test for nonlinearity and the Tsay's test.

Surrogate data testing tests the null hypothesis that the data was generated from a stationary linear stochastic process with Gaussian inputs. Surrogate data testing consists on generating a surrogate data set showing the same linear properties of the RR time series. Then, a statistic is calculated for the RR time series and all the surrogate set. If the value of the statistic is significantly different for the RR series

6.1. AN INTRODUCTION TO NONLINEAR ANALYSIS TECHNIQUES

and all the surrogate set, the null hypothesis is rejected and nonlinearity assumed. In practice, the null hypothesis is rejected when the statistic calculated for the RR series is smaller or larger than those calculated for the surrogate set.

Both functions take as input parameters the *HRVData* structure and the *indexNonLinearAnalysis*. The *indexNonLinearAnalysis* is an integer denoting the *NonLinearAnalysis* structure that will contain the results of the nonlinear analysis. Additionally, the *SurrogateTest* allows the user to specify the significance of the test (*significance* parameter) and specify the function that will compute the discriminating statistic (*useFunction*) as well as its parameters. Finally, it's also possible to obtain a graphical representation of the statistic values of both surrogate data and the RR time series using the *doPlot* parameter.

The next lines show how to use both functions for nonlinearity testing:

```
> # Testing  
> hrv.data = NonlinearityTests(hrv.data)  
  
--- Performing nonlinearity tests ---  
  
      ** Teraesvirta's neural network test **  
  
Null hypothesis: Linearity in "mean"  
X-squared =  903.5991  df =  2  p-value =  0  
  
      ** White neural network test **  
  
Null hypothesis: Linearity in "mean"
```

6.1. AN INTRODUCTION TO NONLINEAR ANALYSIS TECHNIQUES

```
X-squared = 877.33 df = 2 p-value = 0

** Keenan's one-degree test for nonlinearity **

Null hypothesis: The time series follows some AR process
F-stat = 516.1898 p-value = 8.409783e-109

** McLeod-Li test **

Null hypothesis: The time series follows some ARIMA process
Maximum p-value = 0

** Tsay's Test for nonlinearity **

Null hypothesis: The time series follows some AR process
F-stat = 29.45 p-value = 0

** Likelihood ratio test for threshold nonlinearity **

Null hypothesis: The time series follows some AR process
Alternative hypothesis: The time series follows some TAR process
X-squared = 2600.526 p-value = 0

> hrv.data = SurrogateTest(hrv.data, significance = 0.05,
+   useFunction = timeReversibility, tau=4, doPlot = TRUE)

Computing statistics

Null Hypothesis: Original data comes from a linear stochastic process
Reject Null hypothesis: Original data's statistic is the greatest one
```

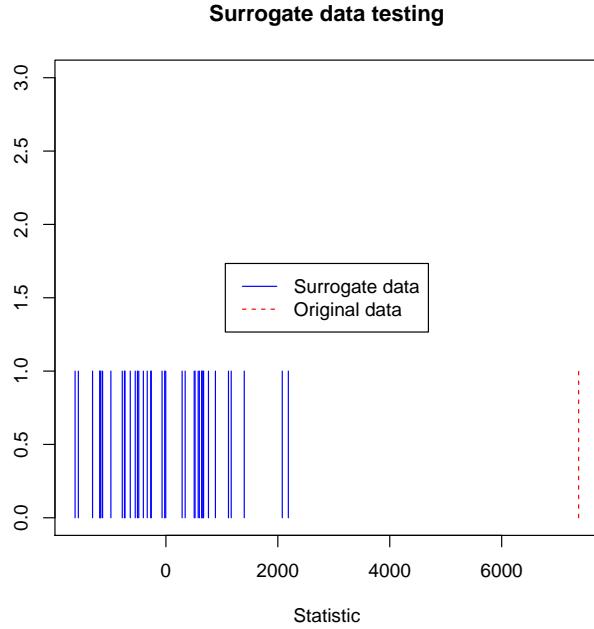


Figure 6.1: *Surrogate data testing.*

The *SurrogateTest* uses the *timeReversibility* function from the *nonlinearTseries* package. This function implements the third-order statistic

$$\psi(\tau) = \frac{1}{N - \tau} \sum_{i=\tau+1}^N (s_n - s_{n-\tau})^3,$$

that is useful for measuring the asymmetry of a series under time reversal. Since linear stochastic series are symmetric under time reversal, this statistic may be used for testing if the data was generated from a stationary linear stochastic process (the null hypothesis). The *tau* parameter is a parameter of the *timeReversibility* function that sets $\tau = \text{tau}$. It must be noted that the discrimination power of the time

simmetry test is low. We have use it here for illustrative purposes. Figure 6.1 shows the results of the surrogate data testing.

6.1.2 Phase space reconstruction

As exposed in Section 2.2.3.1 the Takens embedding theorem provides a method for phase space reconstruction. However, it does not provide information on how to select τ and m parameters. Fortunately, the *RHRV* package provides functionality for estimating both parameters.

6.1.2.1 Time lag estimation

In practical applications, the τ parameter is firstly selected. Then, the embedding dimension is estimated for a fixed value of the τ parameter.

The τ parameter can be estimated in *RHRV* by using the *CalculateTimeLag* function. This function selects the time lag based on the following reasoning: if the time lag used to build the Takens' vectors is too small, the coordinates will be too highly temporally correlated and the embedding will tend to cluster around the diagonal in the phase space. If the time lag is chosen too large, the resulting coordinates may be almost uncorrelated and the resulting embedding will be very complicated. There is a wide variety of methods for estimating an appropriate time lag based on the study of the autocorrelation function of a given time series:

- Select the time lag where the autocorrelation function decays to 0 (*first.zero* method).

- Select the time lag where the autocorrelation function decays to $1/e$ (*first.e.decay* method).
- Select the time lag where the autocorrelation function reaches its first minimum (*first.minimum* method).
- Select the time lag where the autocorrelation function decays to the value specified by the user (*first.value* method and *value* parameter).

The *CalculateTimeLag* function takes as input parameters the *HRVData* and the *method*, *value*, *lagMax* and *doPlot* parameters. The *method* parameter indicates the method that we shall use to estimate the time lag . Its value must be “*first.zero*”, “*first.e.decay*”, “*first.minimum*” or “*first.value*”. The *value* parameter denotes the value that the autocorrelation function must cross in order to select the time lag if the “*first.value*” method is used. The *lagMax* parameter specifies the maximum lag at which to calculate the autocorrelation function. By default, the length of the *timeSeries* is used. Finally, the logical *doPlot* parameter indicates if a plot of the autocorrelation function is shown or not.

In the following example we estimate a proper time lag by using the different methods explained above:

```
> # method "first.zero"  
> CalculateTimeLag(hrv.data,method="first.zero",lagMax=100)
```

```
--- Calculating optimum time lag ---  
--- Time Lag = 1 ---  
[1] 1  
  
> # method "first.minimum"  
  
> CalculateTimeLag(hrv.data,method="first.minimum",lagMax=100,  
+                      doPlot=FALSE)  
  
--- Calculating optimum time lag ---  
--- Time Lag = 1 ---  
[1] 1  
  
> # method "first.value"  
  
> CalculateTimeLag(hrv.data,method="first.value",value=0.1,  
+                      lagMax=100, doPlot=FALSE)  
  
--- Calculating optimum time lag ---  
--- Time Lag = 1 ---  
[1] 1  
  
> # method "first.e.decay" (default)  
  
> kTimeLag = CalculateTimeLag(hrv.data,lagMax=100, doPlot=FALSE)  
  
--- Calculating optimum time lag ---  
--- Time Lag = 1 ---  
  
> print(kTimeLag)
```

6.1. AN INTRODUCTION TO NONLINEAR ANALYSIS TECHNIQUES

```
[1] 1
```

Most of the nonlinear algorithms in *RHRV* accept as input both time lag and embedding dimension parameters. If they are not specified, *RHRV* estimates them by using the functions explained in this section and the next one. However we strongly recommend to estimate them just once because of computational efficiency. Thus, we store the *kTimeLag* parameter to use it in all the examples of this section.

Figure 6.2 shows the autocorrelation function used in the time lag estimation.

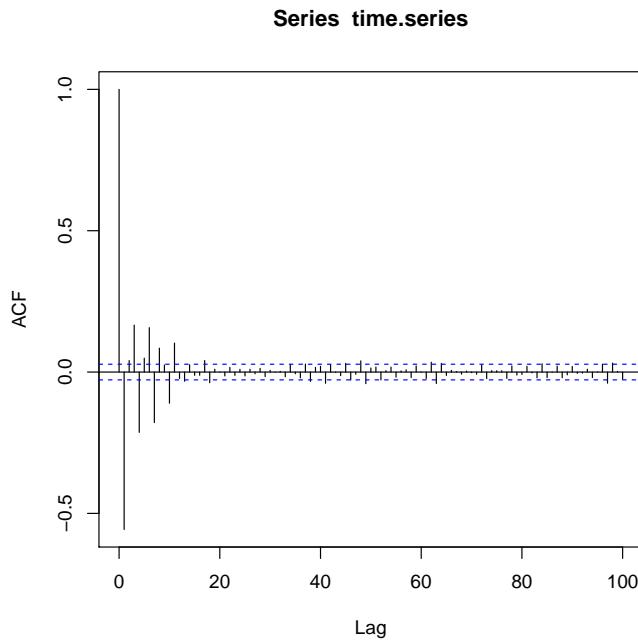


Figure 6.2: Autocorrelation function of the niHR time series used for calculate the optimum TimeLag.

6.1.2.2 Embedding dimension estimation

Once the time lag has been estimated, it is possible to obtain a proper embedding dimension by using the algorithm explained in [8]. The Cao's algorithm uses 2 functions in order to estimate the embedding dimension from a time series: the $E1(d)$ and the $E2(d)$ functions, where d denotes the dimension.

$E1(d)$ stops changing when d is greater than or equal to the embedding dimension m , staying close to 1. On the other hand, $E2(d)$ is used to distinguish deterministic signals from stochastic signals. For deterministic signals, there exists some d fulfilling $E2(d) \neq 1$. For stochastic signals, $E2(d)$ is approximately 1 for all the values. The *CalculateEmbeddingDim* function implements this algorithm in the *RHRV* package.

The *CalculateEmbeddingDim* takes as input parameters the *HRVData* structure, the number of points of the time series to be used for estimating the embedding dimension (*numberPoints*); the time lag calculated in the previous Section (*timeLag*); the maximum possible embedding dimension (*maxEmbeddingDim*); the threshold that $E1(d)$ must cross for considering that it is close to the limit value 1 (*threshold*; the default value is 0.95); the maximum relative change in $E1(d)$ with respect to $E1(d - 1)$ in order to consider that the $E1$ function has been stabilized and that it will stop changing (*maximumRelativeValue*; the default value is 0.01) and the *doPlot* parameter (If TRUE a plot of $E1(d)$ and $E2(d)$ is shown).

6.1. AN INTRODUCTION TO NONLINEAR ANALYSIS TECHNIQUES

In the following example we estimate a proper embedding dimension using the time lag calculated in the previous section. Figure 6.3 shows the resulting plot.

```
> kEmbeddingDim = CalculateEmbeddingDim(hrv.data,
+                                         numberPoints = 10000,
+                                         timeLag = kTimeLag,
+                                         maxEmbeddingDim = 15)

--- Calculating optimum embedding dimension ---
```

>

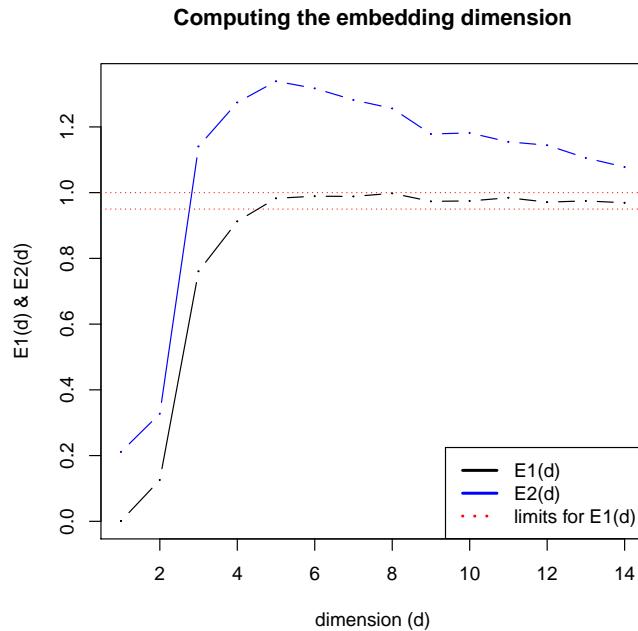


Figure 6.3: Estimation of the embedding dimension using the Cao's algorithm.

We store the result in the *kEmbeddingDim* variable in order to use it in the remainder of this section.

6.1.3 Computing nonlinear statistics

6.1.3.1 The classic correlation dimension

In order to calculate the correlation dimension of our RR time series we can use the *CalculateCorrDim* function. This function takes as input the HRVData structure to be analyzed and the index of the *NonlinearAnalysis* structure that shall store the results. The embedding dimensions in which the *CalculateCorrDim* will be computed are specified with the *minEmbeddingDim* and the *maxEmbeddingDim* parameters (remember that we should compute the correlation dimension for several embedding dimensions). The time lag needed for reconstructing the phase space is provided by the *timeLag* parameter. The *minRadius* and *maxRadius* parameters specify the radius in which the correlation sum is going to be calculated. The number of points used to compute this radius range can be specified with the *pointsRadius* parameter (i.e. we shall use *pointsRadius* points in order to cover the $[minRadius, maxRadius]$ interval). The Theiler window is specified using the *theilerWindow* parameter. Finally, a log-log plot of the correlation sum Vs. the radius can be obtained by setting *doPlot* = *TRUE* (default).

For analyzing our *hrv.data* structure we shall use the embedding dimension and time lag obtained in Section 6.1.2. Thus, we set *timeLag* = *kTimeLag* and we se-

lect a sequence of dimensions surrounding $kEmbeddingDim$. For example, we may use $minEmbeddingDim = kEmbeddingDim - 1$ and $maxEmbeddingDim = kEmbeddingDim + 2$. In order to select the radius we have to take into account that we are analyzing RR intervals in milliseconds. Thus, the selection $minRadius = 1$ and $maxRadius = 100$ will probably cover the most important range of the correlation sum $C(r)$. By selecting $pointsRadius = 100$, the correlation sum $C(r)$ will be computed in $r = 1, 2, \dots, 100$. In order to select a proper Theiler window, the autocorrelation function may be used. Figure 6.2 shows that the RR series is practically uncorrelated after the time lag 10. In order to be sure that we got rid of the temporal correlations we set $theilerWindow = 20$. Finally, we set $doPlot=FALSE$ (later, we will plot the correlation sum with the *PlotCorrDim* function).

```
> my.index = 1
> hrv.data = CalculateCorrDim(hrv.data,
+                               indexNonLinearAnalysis = my.index,
+                               minEmbeddingDim = kEmbeddingDim - 1,
+                               maxEmbeddingDim = kEmbeddingDim + 2,
+                               timeLag = kTimeLag, minRadius=1,
+                               maxRadius=100, pointsRadius = 100,
+                               theilerWindow = 20, doPlot = FALSE)
--- Computing the Correlation sum ---
```

As usual, we may use the `$` operator from the R language to access the correlation sum calculated with the *CalculateCorrDim* function and stored under the *correlation\$computations* field of the *NonLinearAnalysis* structure:

```
> corr.struct = hrv.data$NonLinearAnalysis[[my.index]]$correlation
> corrSum = corr.struct$computations
```

Now, in order to access to the *corrSum* attributes we may write:

```
> # Let's print the four first rows and columns of the correlation matrix
> print(corrSum$corr.matrix[1:4,1:4])
```

	100	95.4548456661834	91.1162756115489	86.9749002617783
4	0.9988547	0.9917903	0.9854796	0.9778878
5	0.9984679	0.9896422	0.9818793	0.9728309
6	0.9980811	0.9874943	0.9782899	0.9678070
7	0.9972923	0.9849646	0.9743409	0.9624882

```
> # Access the radius and embedding dimensions used for computations
> radius = corrSum$radius
> embeddingDims = corrSum$embedding.dims
```

Note the correlation matrix stores all the correlation sums that have been computed. Each row stores the correlation sum for a concrete embedding dimension whereas each column stores the correlation sum for a specific radius. Note that the names of the matrix dimensions denote the embedding dimension and the radius (in descending order).

A graphical representation of the correlation sum can be obtained by using the *Plot-CorrDim* function. This function shows a log-log plot of the correlation sum Vs the radius and the local slopes of $\log_{10}(C(r))$ Vs $\log_{10}(C(r))$.

6.1. AN INTRODUCTION TO NONLINEAR ANALYSIS TECHNIQUES

```
> PlotCorrDim(hrv.data, indexNonLinearAnalysis=my.index)
```

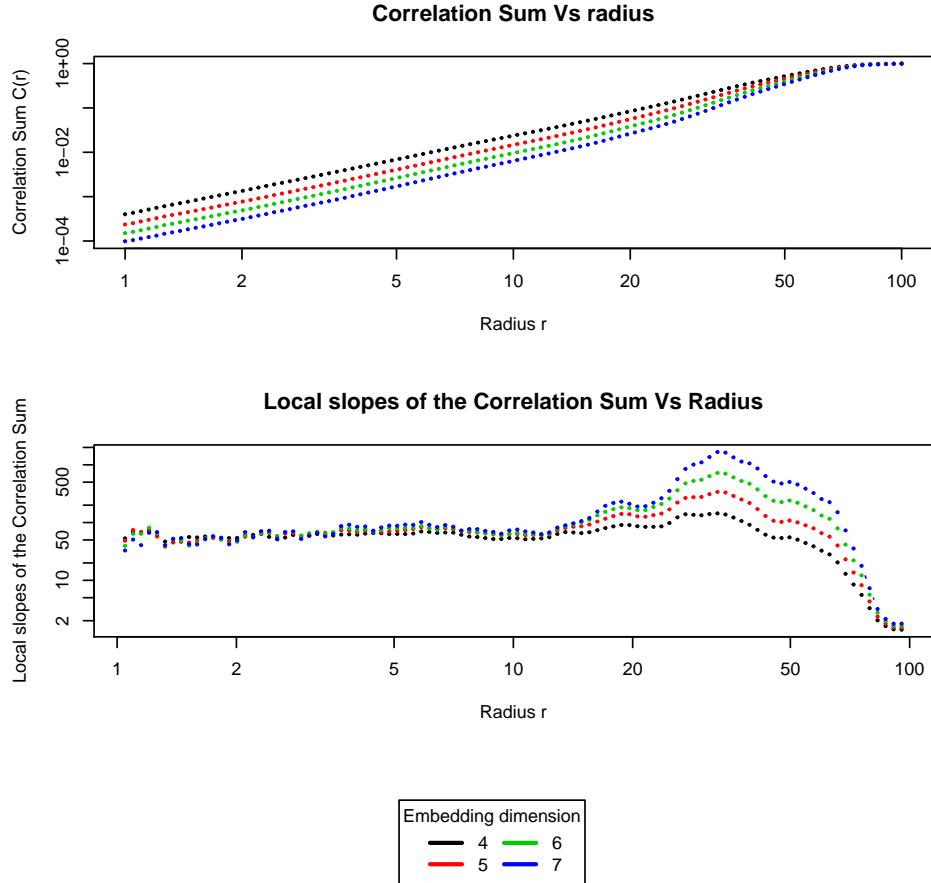


Figure 6.4: Correlation sums calculation.

In order to get an estimation of the correlation dimension, the *EstimateCorrDim* function can be used. This function estimates the correlation dimension of the RR time series by averaging the slopes of the embedding dimensions specified in the *useEmbeddings* parameter. The slopes are determined by performing a linear regression over the radius' range specified in *regressionRange*. The user must select

the dimensions and radius in which the slopes of the correlation sum are approximately equal. If such a region does not exist, the estimation should be discarded. In our case, we may select all the embedding dimensions and set the *regressionRange* to [1.5, 10] (see Figure 6.4). If *doPlot=TRUE* a graphic of the regression over the data is shown. The results are returned into a new field called *statistic* under the *hrv.data\$NonLinearAnalysis[[indexNonLinearAnalysis]]\$correlation* list:

```
> hrv.data = EstimateCorrDim(
+                               hrv.data,
+                               indexNonLinearAnalysis=my.index,
+                               regressionRange=c(1.5,10),
+                               useEmbeddings=(kEmbeddingDim-1):(kEmbeddingDim+2),
+                               doPlot=TRUE)

--- Estimating the Correlation dimension ---
--- Correlation dimension = 1.828973 ---

> # We may add a legend to the plot
> legend(x=10,y=-3, lty = rep(1,3), col = 1:3,
+         title="Embedding dimension",
+         legend = (kEmbeddingDim-1):(kEmbeddingDim + 2))
> cat("The correlation dimension is ",
+      hrv.data$NonLinearAnalysis[[my.index]]$correlation$statistic, "\n")
```

The correlation dimension is 1.828973

The plot obtained in the estimation process is shown in Figure 6.5. It must be noted

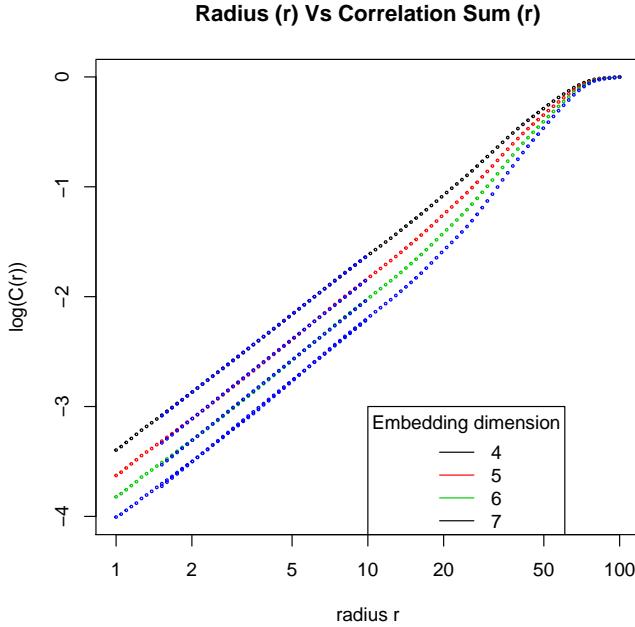


Figure 6.5

that all the nonlinear statistics following the the *CalculateX-EstimateX-PlotX* procedure store their results in the two fields presented in this section: the *computations* field shall store the calculations whereas that the *statistic* field shall store the final estimation of the statistic being computed.

Since the correlation dimension is the most important of the nonlinear dimensions, the generalized dimensions (including the information dimension) shall be presented in Section 6.2.

6.1.3.2 Sample entropy

In Section 2.2.3.5 we saw that in order to compute the sample entropy of order q (H_q) of a time series we first have to calculate:

$$h_q(m, r) = \log \left(\frac{C_q(m, r)}{C_q(m + 1, r)} \right),$$

and then, we may obtain H_q by taking the limit:

$$H_q = \lim_{\substack{r \rightarrow 0 \\ m \rightarrow \infty}} h_q(m, r).$$

Thus, in *RHRV*, in order to use the *CalculateSampleEntropy* function a “corrDim” object with several embedding dimensions must exist under the *NonLinearAnalysis* structure. Also, the embedding dimensions must be high enough (in order to “simulate” the limit $m \rightarrow \infty$).

Let’s compute the sample entropy of our example file. First of all, we shall compute several correlations sums using high embedding dimensions. In this example, we shall use at least four times the optimum embedding dimension *kEmbeddingDim*. Then, we use the *CalculateSampleEntropy* function to compute the $h_q(m, r)$ function. This function is quite simple since only takes as input parameters the *HRVData* being analyzed, the *indexNonLinearAnalysis* and the boolean argument *doPlot* (allowing to obtain a plot of $h_q(m, r)$).

```
> hrv.data = CreateNonLinearAnalysis(hrv.data)
```

```

** Creating non linear analysis

Data has now 2 nonlinear analysis

> my.index = 2

> hrv.data = CalculateCorrDim(hrv.data,
+                               indexNonLinearAnalysis = my.index,
+                               minEmbeddingDim = 4*kEmbeddingDim,
+                               maxEmbeddingDim = 4*kEmbeddingDim+5,
+                               timeLag = kTimeLag, minRadius = 1,
+                               maxRadius = 100, pointsRadius = 100,
+                               theilerWindow = 20, doPlot = FALSE)

--- Computing the Correlation sum ---

> hrv.data = CalculateSampleEntropy(hrv.data,
+                                     indexNonLinearAnalysis= my.index,
+                                     doPlot = FALSE)

--- Computing the sample entropy of order 2 ---

```

We can also plot the $h_q(m, r)$ function with the *PlotSampleEntropy* function. The resulting figure is shown in Figure 6.6.

```
> PlotSampleEntropy(hrv.data, indexNonLinearAnalysis=my.index)
```

Since we require H_q not to depend on m nor r , we shall select a radius range in which the $h_q(m, r)$ is approximately flat. Also, the height of this flat region should be approximately the same for some embedding dimensions (this height will be the value

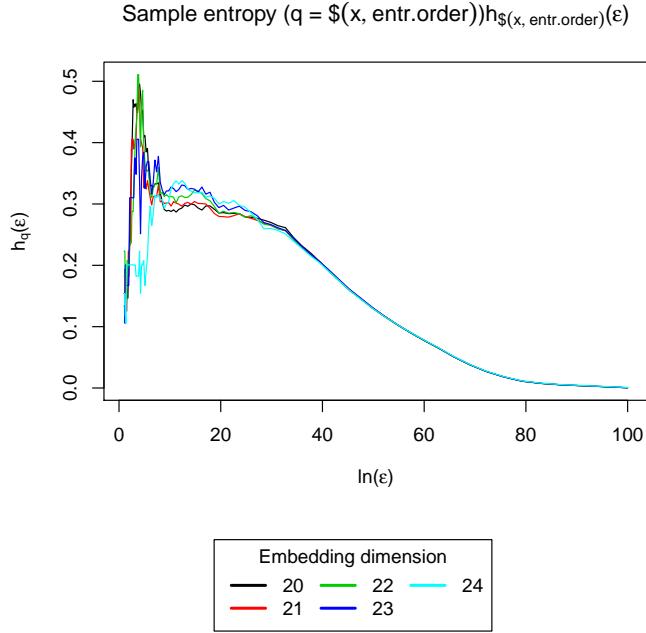


Figure 6.6: Sample entropy computations.

of H_q). The *EstimateSampleEntropy* function takes as input the radius range (*regressionRange*) and the embedding dimensions selected (*useEmbeddings*) to compute an estimation of the sample entropy H_q .

```
> hrv.data = EstimateSampleEntropy(hrv.data,
+                                     indexNonLinearAnalysis=my.index,
+                                     regressionRange=c(10,20),
+                                     useEmbeddings = 20:22,
+                                     doPlot = TRUE)

--- Computing the sample entropy---
--- Sample entropy with its order: ---
```

	20	21	22
0.2930925	0.2975794	0.3095050	

Figure 6.7 shows the plot generated by the *EstimateSampleEntropy* function.

The results of the sample entropy estimations are stored under the *sampleEntropy* field of the *NonLinearAnalysis* structure as expected: The *calculations* field stores $h_q(m, r)$ whereas that the *statistic* field stores the estimate of the sample entropy. The attributes of the *calculations* field can be accessed as usual using `$. calculations$sample.entropy` returns a matrix containing the $h_q(m, r)$ function (the rows store the computations for an specific embedding whereas that the columns depend on the radius. `calculations$radius` and `calculations$embedding.dims` return the radius and the embedding dimensions used for calculations:

```

> se = hrv.data$NonLinearAnalysis[[my.index]]$  

+ sampleEntropy$calculations  

> # obtaining the  $h_q(m,r)$  function, the radius and the embedding dimensions  

> sample.function = se$sample.entropy  

> radius = se$radius  

> dims = se$embedding.dims  

> # obtaining the sample entropy estimation  

> estimation =  

+ hrv.data$NonLinearAnalysis[[my.index]]$sampleEntropy$statistic  

> cat("The sample entropy is ")

```

The sample entropy is

```
> print(estimation)
```

```
20      21      22
0.2930925 0.2975794 0.3095050
```

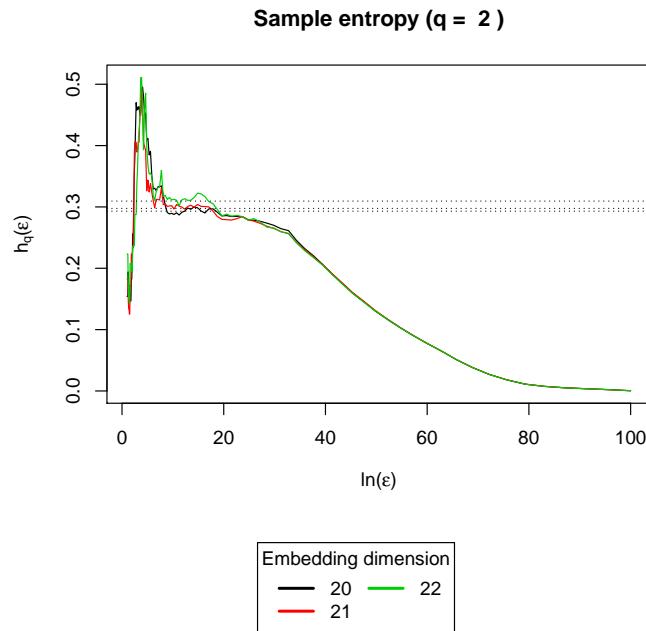


Figure 6.7: Sample entropy estimation.

6.1.3.3 Maximum Lyapunov exponent

As presented in Section 2.2.3.6, the first step in order to calculate the maximum Lyapunov exponent should be computing the $S(t)$ function. In *RHRV*, the $S(t)$ shall be estimated with the *CalculateMaxLyapunov* function. The *CalculateMaxLyapunov*

function will average the divergence of several reference points in a m-dimensional space. The *minEmbeddingDim*, *maxEmbeddingDim* and *timeLag* specify the parameters for the phase space reconstruction. The number of reference points that the routine will try to use in order to compute the divergence is specified with the *minRefPoints* parameter. A point in the phase space is considered to be a reference point if there exist a minimum number of close neighbours. 500 points are usually enough (default). The number of close neighbours needed to be considered a reference point can be specified with the *minNeighs* parameter (default value is 5). The *radius* specifies the maximum distance in which the routine will look for close neighbours. Since we are working with the RR time series, a radius in the radius' range [1 – 10] ms seems a proper choice for selecting very close phase space points.

In order to compute the maximal Lyapunov exponent for our example file we could write:

```
> my.index = 1  
> hrv.data = CalculateMaxLyapunov(  
+                               hrv.data,  
+                               indexNonLinearAnalysis = my.index,  
+                               minEmbeddingDim= kEmbeddingDim,  
+                               maxEmbeddingDim= kEmbeddingDim+2,  
+                               timeLag = kTimeLag,
```

```

+
          radius = 3, theilerWindow = 20,
+
          doPlot = FALSE)

--- Computing the divergence of the time series ---

```

As usual, the user should plot $S(t)$ Vs t when looking for the maximal Lyapunov exponent. This can be done using the *doPlot* parameter of the *CalculateMaxLyapunov* routine or using the *PlotMaxLyapunov* function.

```
> PlotMaxLyapunov(hrv.data, indexNonLinearAnalysis = 1)
```

The resulting plot is shown in Figure 6.8. After plotting the $S(t)$ function we should check if $S(t)$ shows a linear behaviour for some temporal range. If that's the case, its slope is an estimate of the maximal Lyapunov exponent per unit of time. The *EstimateMaxLyapunov* routine allows the user to get always an estimate of the maximal Lyapunov exponent, but the user must check that there is a linear region in the $S(t)$ Vs t . If such a region does not exist, the estimation should be discarded.

In our example we found a strong linear region in the [1, 6] interval. The regression performed by the *EstimateMaxLyapunov* is shown in Figure 6.9.

```

> hrv.data = EstimateMaxLyapunov(
+
          hrv.data,
+
          indexNonLinearAnalysis = my.index,
+
          regressionRange = c(1, 6),
+
          useEmbeddings = (kEmbeddingDim):(kEmbeddingDim+2),
+
          doPlot = TRUE)

```

6.1. AN INTRODUCTION TO NONLINEAR ANALYSIS TECHNIQUES

--- Estimating the Maximum Lyapunov exponent ---

--- Maximum Lyapunov exponent = 0.4863784 ---

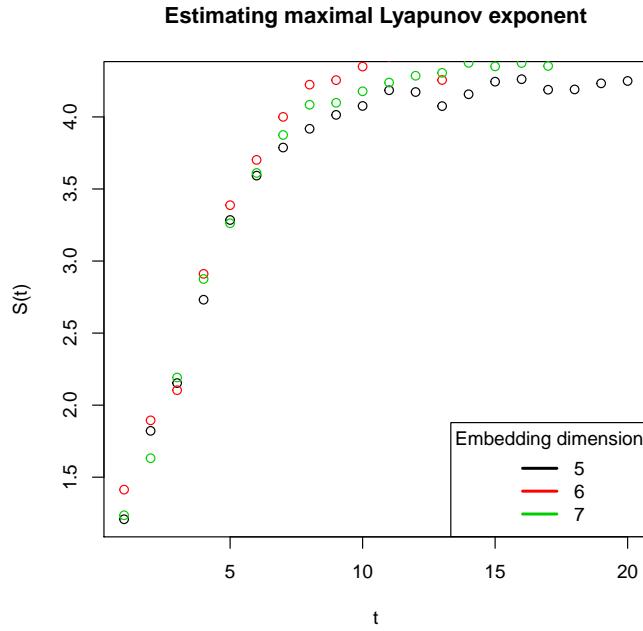


Figure 6.8: Divergence computations for the maximum Lyapunov exponent.

As always, we can access the computations under the *NonLinearAnalysis* list. The divergence computations are stored under the *lyapunov\$computations* list of the *NonLinearAnalysis* structure. The maximum Lyapunov exponent is stored under the *lyapunov\$statistic* field.

```
> lyapunov.results = hrv.data$NonLinearAnalysis[[my.index]]$lyapunov
> divergence.structure = lyapunov.results$computations
> max.exponent = lyapunov.results$statistic
> # get the S(t) function
```

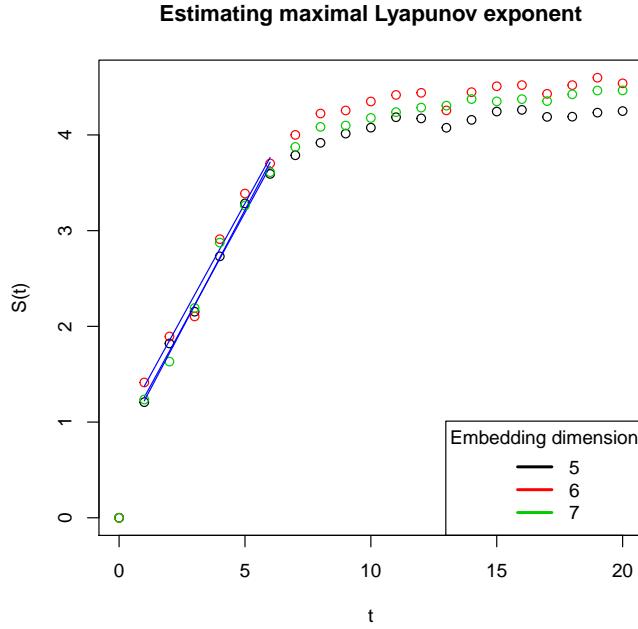


Figure 6.9: Maximum Lyapunov exponent estimation.

```
> div = divergence.structure$s.function
> tim = divergence.structure$time
> cat("The Max Lyapunov exponent is ",max.exponent,"\\n")
```

The Max Lyapunov exponent is 0.4863784

6.1.3.4 Detrended Fluctuation Analysis

Before introducing the *RHRV* functionality for performing DFA, we have to review the DFA procedure presented in Section 2.2.3.7.

1. Integrate the time series to be analyzed. The time series resulting from the integration will be referred to as the *profile*.

6.1. AN INTRODUCTION TO NONLINEAR ANALYSIS TECHNIQUES

2. Divide the profile into N non-overlapping segments.
3. Calculate the local trend for each of the segments using least-square regression.
Compute the total error for each of the segments.
4. Compute the average of the total error over all segments and take its root square. By repeating the previous steps for several segment sizes (let's denote it by t : number of beats), we obtain the so-called *fluctuation function* $F(t)$.
5. If the data presents long-range power law correlations: $F(t) \propto t^\alpha$, we can estimate the exponent using regression.
6. Usually, when plotting $\log(F(t))$ Vs $\log(t)$ we may distinguish two linear regions. By performing a regression on each of them separately, we obtain two scaling exponents, α_1 (the exponent for small values of t , characterizing short-term fluctuations) and α_2 (the exponent for large values of t , characterizing long-term fluctuations).

Steps 1-4 are performed in *RHRV* using the *CalculateDFA* function. In order to obtain a estimate of some scaling exponent, the user must use the *EstimateDFA* function specifying the regression range (window sizes used to detrend the series). α_1 is usually obtained by performing the regression in the $4 \leq t \leq 16$ range whereas that α_2 is obtained in the $16 \leq t \leq 64$ range (However the $F(t)$ function must be linear in these ranges to obtain reliable results).

Besides the *HRVData*, *indexNonLinearAnalysis* and *doPlot* parameters, the *CalculateDFA* function accepts as inputs the range of values for the window sizes that will

6.1. AN INTRODUCTION TO NONLINEAR ANALYSIS TECHNIQUES

be used to estimate the fluctuation function (*windowSizeRange*, the default value is *c(10,300)*) and the number of different window sizes in that range that will be used to estimate the Fluctuation function (*npoints*). The *EstimateDFA* and *PlotDFA* functions work as usual.

```
> hrv.data = CalculateDFA(hrv.data,
+                           indexNonLinearAnalysis = 1,
+                           windowSizeRange = c(6, 300),
+                           npoints = 25,
+                           doPlot = FALSE)

--- Performing Detrended Fluctuation Analysis---

> hrv.data = EstimateDFA(hrv.data,
+                           indexNonLinearAnalysis = 1,
+                           regressionRange = c(20,100), doPlot = TRUE)

--- Estimating Scaling exponent ---
--- Scaling Exponent number 1 = 0.1442544 ---
```

Figure 6.10 show the regression performed over the Fluctuation function.

The *dfa* computations are stored in the *dfa\$computations* list under the *NonLinearAnalysis* structure. The fluctuation function and the windows used for the computations can be obtained using *\$fluctuation.function* and *\$window.sizes* on the *dfa\$computations* list, respectively. All the estimated exponents are stored under

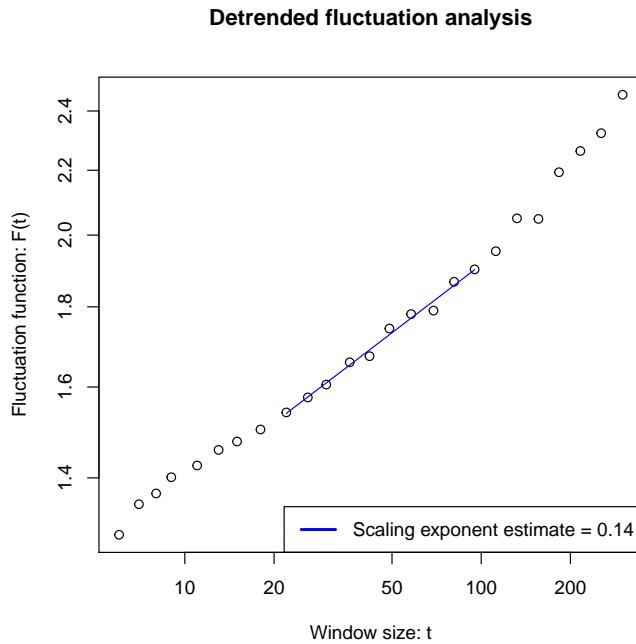


Figure 6.10: DFA.

the `NonLinearAnalysis[[index]]dfastatistic` list. It must be noted that, in this case, the `statistic` field may contain more than just one value for the statistic (depending on the regression range used).

```
> #obtaining the object pf class "dfa"
> dfa.object = hrv.data$NonLinearAnalysis[[1]]$dfa$computations
> windows = dfa.object>window.sizes
> fluctuation.f = dfa.object$fluctuation.function
> # get the exponent. Note the index 1 in the statistic field!!
> scaling.exp = hrv.data$NonLinearAnalysis[[1]]$dfa$statistic[[1]]
> print(scaling.exp)
```

```
$range  
[1] 20 100
```

```
$estimate  
[1] 0.1442544
```

6.1.3.5 Recurrence Quantification Analysis

In order to perform Recurrence Quantification analysis of the RR time series *RHRV* provides the *RQA* routine. The *RQA* function accepts as parameters the *HRVData*, *indexNonLinearAnalysis*, *embeddingDim*, *timeLag* and *doPlot* with their usual meaning. The user may also specify the *radius* (maximum distance between two phase-space points to be considered a recurrence) and *numberPoints* (number of points to be used in the RQA computation). Since this method requires heavy computations, this last parameter is specially useful.

Let's apply the *RQA* function to our example:

```
> hrv.data = RQA(hrv.data, indexNonLinearAnalysis = my.index,  
+                   embeddingDim=kEmbeddingDim, timeLag = kTimeLag,  
+                   radius = 2, doPlot=TRUE)  
  
--- Plotting recurrence plot ---  
  
> # let's see which statistics have been computed...  
> names(hrv.data$NonLinearAnalysis[[1]]$rqa)
```

6.1. AN INTRODUCTION TO NONLINEAR ANALYSIS TECHNIQUES

```
[1] "REC"           "RATIO"        "DET"  
[4] "DIV"          "Lmax"         "Lmean"  
[7] "LmeanWithoutMain"  "ENTR"        "TREND"  
[10] "LAM"          "Vmax"         "Vmean"  
[13] "diagonalHistogram" "recurrenceRate"  
  
> #... and access one of them  
> cat("Entropy of the diagonal lines: ",  
+     hrv.data$NonLinearAnalysis[[1]]$rqa$ENTR,  
+     "\n")
```

Entropy of the diagonal lines: 1.805314

Figure 6.11 shows the resulting recurrence plot. The only issue of the previous example is the selection of the radius... Why did we set *radius*=2? Since we are analyzing RR time series in milliseconds it seems that a reasonable choice for analyzing close phase space points could be our selection: RR vectors whose maximum difference is less than 2 ms. However, other similar values for the radius can be used. An useful rule of thumb for selecting the *radius* parameter is choosing the radius so that the recurrence matrix is sparse although it must contain certain vertical or diagonal lines (deterministic structures). It must also be noted that high values for the radius may result in very heavy computations.

As seen in the previous example, in addition to the usual statistics (presented in table 2.1), the *RQA* function also returns the histogram of the length of diagonal

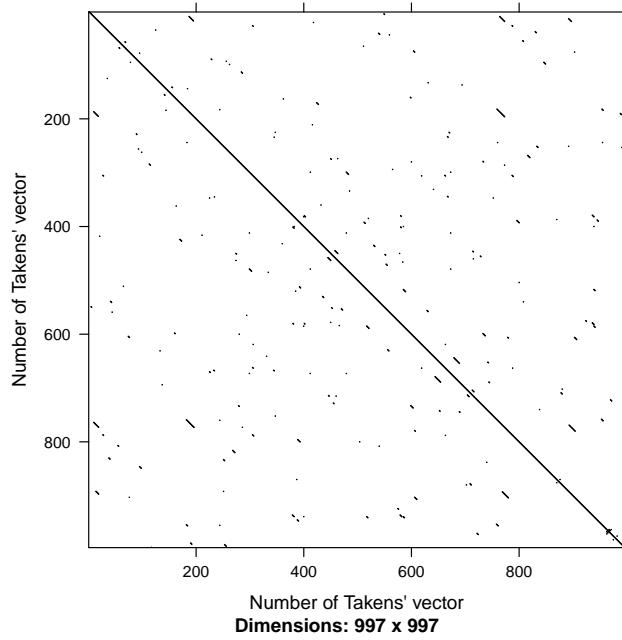


Figure 6.11: Recurrence plot.

lines (*diagonalHistogram*) and the number of recurrence points among all possible points depending on the distance to the main diagonal (*recurrenceRate*). Let's get a plot of the recurrence rate (see Figure 6.12):

```
> recurrence.rate = hrv.data$NonLinearAnalysis[[1]]$rqa$recurrenceRate
> plot(1:length(recurrence.rate),recurrence.rate,type="l",
+       xlab="Distance to main diagonal",ylab="Recurrence Rate",
+       main="Recurrence Rate")
```

Note the border effects in the plot. Since near the borders there are very few points, the recurrence rate is usually biased. Of course we can ignore the end of the plot, but we could have avoided the border effects by using the *distanceToBorder* parameter

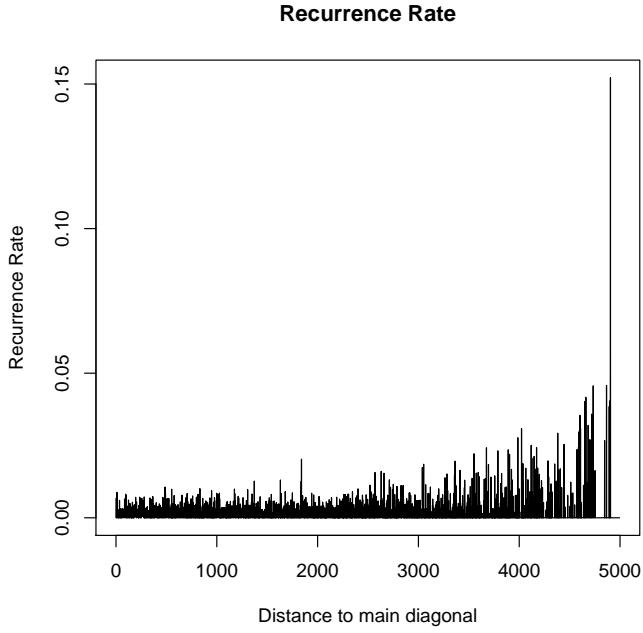


Figure 6.12: Recurrence plot.

(The points that are closer than *distanceToBorder* to the border of the recurrence matrix are ignored when computing the RQA parameters).

The histogram of the length of diagonal lines is not widely used because short diagonals prevail. The *lmin* parameter can be used in order to set a minimum length of a diagonal line to be considered in the RQA (default value is *lmin* = 2). Similarly, the *vmin* parameter can be used in order to set the minimum length of a vertical line to be considered in the RQA (default value is *vmin* = 2).

6.2 Advanced nonlinear analysis techniques

All the nonlinear functions presented in 6.1 used as example file an RR time series generated synthetically from a nonlinear process. In this section we will deal with a real RR time series and we will present all the remaining nonlinear functions.

The example file “example2.beats” that we shall use in the next sections can be downloaded from the project’s website <http://rhrv.r-forge.r-project.org/>. The file was obtained from an ECG of one of the authors of the *RHRV* package and thus we expect the file to be from a healthy subject!

```
> hrv.data = CreateHRVData( )  
  
> hrv.data = LoadBeatAscii(hrv.data,  
+                               RecordName="example2.beats",  
+                               RecordPath="beatsFolder")  
  
> hrv.data = BuildNIHR(hrv.data)  
  
> hrv.data = FilterNIHR(hrv.data)  
  
> hrv.data = InterpolateNIHR (hrv.data, freqhr = 4)  
  
> hrv.data = CreateNonLinearAnalysis(hrv.data)  
  
> hrv.data = SetVerbose(hrv.data, TRUE)
```

First of all, we should run the nonlinearity tests in order to make sure that the time series shows some degree of nonlinearity. We leave the test to the reader. Now, we shall estimate both time lag and embedding dimension parameters of our new RR time series. As always, we first estimate the time lag. The autocorrelation function

6.2. ADVANCED NONLINEAR ANALYSIS TECHNIQUES

used for the selection of the time lag is shown in Figure 6.13. This figure also suggest that a suitable choice for the Theiler window could be ≈ 200 .

```
> kTimeLag = CalculateTimeLag(hrv.data, method = "first.minimum",
+                               lagMax = 300)

--- Calculating optimum time lag ---

--- Time Lag =  2  ---
```

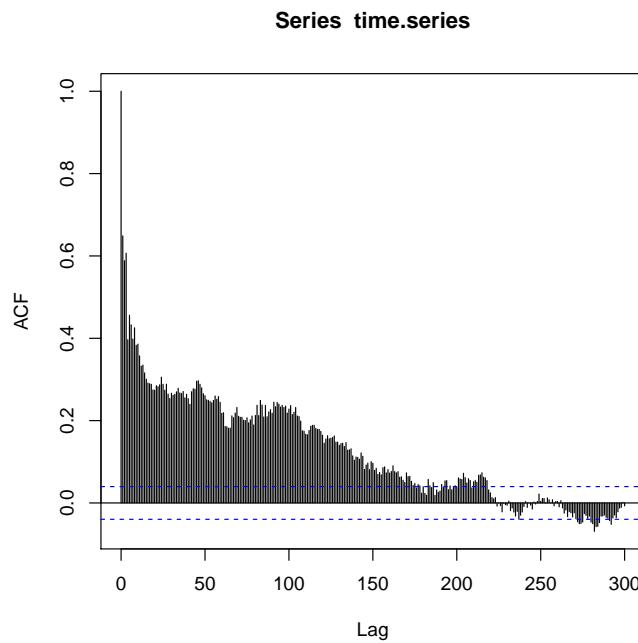


Figure 6.13: Autocorrelation function for the RR time series.

We can appreciate that there is high autocorrelation in the time series even for large time lag values. Although our estimation is reasonable it may happen that some RR series have their “optimum” time lag at 100 or even more. However, in most HRV

publications, the time lag value is not usually selected above 10 or 15.

Now, we may try to obtain an estimation of the embedding dimension (see Figure 6.14). The optimum embedding dimension is 12.

```
> kEmbeddingDim = CalculateEmbeddingDim(hrv.data,
+                                         numberPoints = 10000,
+                                         timeLag = kTimeLag,
+                                         maxEmbeddingDim = 18)

--- Calculating optimum embedding dimension ---
```

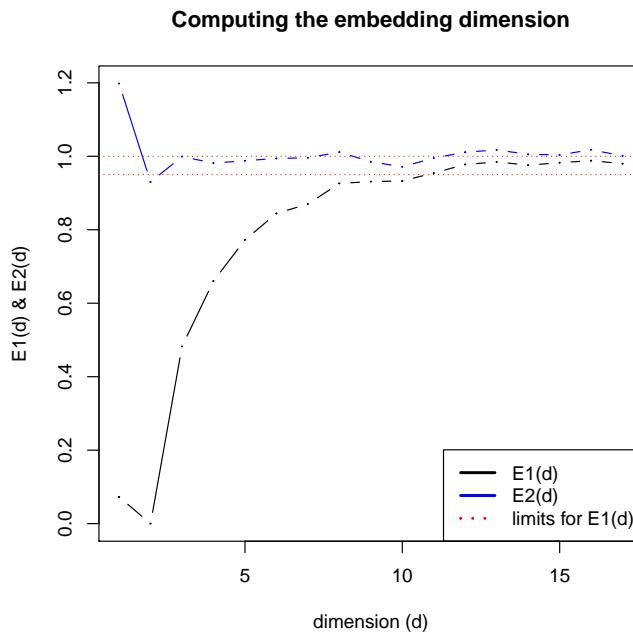


Figure 6.14: Automatic estimation of the embedding dimension.

6.2.1 Nonlinear noise reduction

Since the RR intervals are derived from an ECG, the RR time series may suffer from discretization problems. If the ECG was discretized very coarsely, many RR intervals shall have the same value. Thus, there will be several identical points in phase space, which may bias the nonlinear algorithms. *RHRV* provides functionality that addresses those situations in which the ECG was discretized very coarsely. In order to deal with the discretization problem of the RR time series, *RHRV* provides the *NonLinearNoiseReduction* function. This function adds uniform white noise of a magnitude equal to the resolution of the RR intervals. If the ECG from which the RR intervals were derived was sampled at f_s Hz, uniform $[-0.5, 0.5]/f_s$ white noise is added. Then, a nonlinear noise reduction algorithm is applied. The noise reduction algorithm performs noise reduction by averaging each Takens' vector in an m-dimensional space with his neighbours (time lag=1). Each neighbourhood is specified with balls of a given radius (max norm is used). Although this procedure will certainly not solve the discretization problem, it will alleviate it:

```
> hrv.data = NonLinearNoiseReduction(hrv.data,
+                                         embeddingDim = kEmbeddingDim)

** Denoising RR time series using nonlinear techniques **

** Calculating non-interpolated heart rate **

Number of beats: 2434
```

Additional parameters for the *NonLinearNoiseReduction* function are *ECGsamplingFreq*: the ECG sampling frequency (optional, but it can improve the performance of the algorithm if provided) and *radius*: the radius used to looking for neighbours in the phase space.

6.2.2 Generalized correlation dimensions

Section 6.1.3.1 presented the *CalculateCorrDim*, *EstimateCorrDim* and *PlotCorrDim* functions for computing the correlation dimension of an RR time series. These functions can also be used for computing the generalized correlation dimension. The user only has to specify the order of the correlation dimension with the *corrOrder* parameter in the *CalculateCorrDim* (the default value is 2, the correlation dimension). The order must fulfill $corrOrder > 1$. The generalized correlation dimension of order $q = 1$ is known as the information dimension, that is computed with the *CalculateInfDim* function (see Section 6.2.3).

As a quick example, we will try to calculate the correlation dimension of order 4 for our new *hrv.data*. Figures 6.15 and 6.16 shows the results of the correlation sum computation and the estimation, respectively.

```
> hrv.data = CreateNonLinearAnalysis(hrv.data)

** Creating non linear analysis

Data has now 2 nonlinear analysis
```

```
> my.index = 1

> hrv.data = CalculateCorrDim(hrv.data,
+                               indexNonLinearAnalysis = my.index,
+                               minEmbeddingDim = kEmbeddingDim - 2,
+                               maxEmbeddingDim = kEmbeddingDim + 2,
+                               timeLag = kTimeLag, minRadius = 1,
+                               maxRadius = 200, pointsRadius = 100,
+                               theilerWindow = 200, doPlot = TRUE,
+                               corrOrder = 4)

--- Computing the generalized Correlation sum of order 4 ---


> hrv.data = EstimateCorrDim(hrv.data, indexNonLinearAnalysis = 1,
+                               regressionRange=c(3*10^5,9*10^5),
+                               useEmbeddings = 13:14)

--- Estimating the generalized Correlation dimension of order 4 ---
--- Generalized Correlation dimension of order 4 = 8.585802 ---

>
```

6.2.3 Information dimension

As seen in Section 2.2.3.4 the information dimension is a particular case of the generalized correlation dimension when setting the order $q = 1$. In the *RHRV* package,

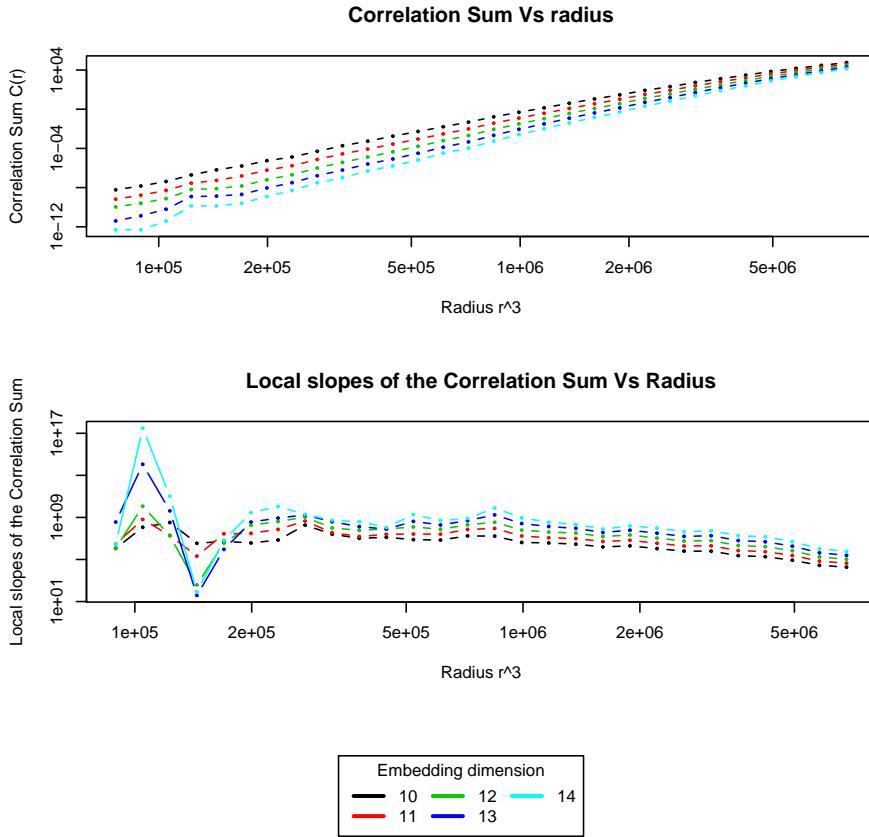


Figure 6.15: Generalized correlation dimension computations.

the information dimension is computed using the *CalculateInfDim*, *EstimateInfDim* and *PlotInfDim*.

Since the *EstimateInfDim* and *PlotInfDim* are analogous to the *EstimateCorrDim* and *PlotCorrDim* functions, we shall only review the *CalculateInfDim* function.

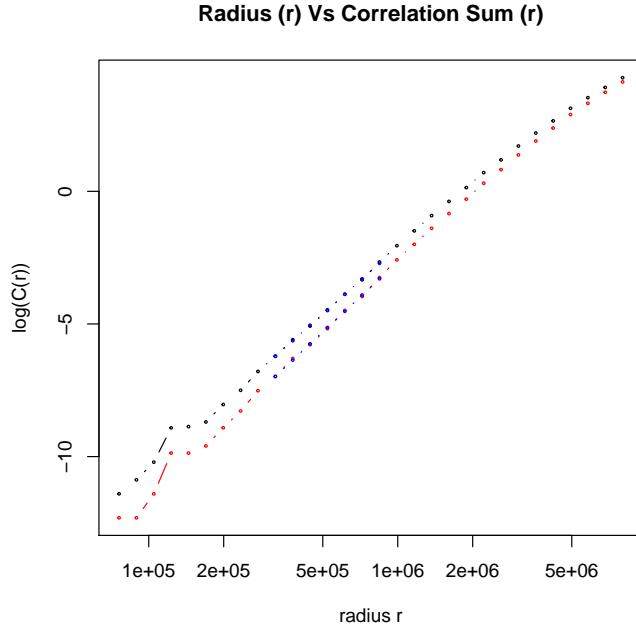


Figure 6.16: Generalized correlation dimension estimation.

In Section 2.2.3.4 we introduced the method for estimating the information dimension in practical applications. This algorithm looks for the scaling behaviour of the average radius that contains a given portion (a "fixed-mass") of the total points in the phase space. By performing a linear regression of $\log(p)$ Vs. $\log(< r >)$ (being p the fixed-mass of the total points), an estimate of the information dimension (D_1) is obtained.

It must be noted that the calculations for the information dimension are heavier than those needed for the correlation dimension. However, the user should run the method for different embedding dimensions to check if D_1 saturates.

```
> my.index = 1

> hrv.data = CalculateInfDim(hrv.data,
+                               indexNonLinearAnalysis=my.index,
+                               minEmbeddingDim=kEmbeddingDim-1,
+                               maxEmbeddingDim=kEmbeddingDim+1,
+                               timeLag=kTimeLag,
+                               minFixedMass=2*10^-3,maxFixedMass=0.25,
+                               numberFixedMassPoints=10,
+                               radius=0.8,increasingRadiusFactor=1.05,
+                               numberPoints=500, theilerWindow=200,
+                               doPlot=TRUE)

--- Computing the Information dimension ---
```

Figure 6.17 shows $\langle \log p(r) \rangle$ Vs. $\log(r)$ as obtained from the *CalculateInfDim* function. As usual, we may use the *Estimate* function (*EstimateInfDim* function) in order to obtain the estimation and then accessing it using \$ under the *statistic* field.

```
> hrv.data = EstimateInfDim(hrv.data,
+                               indexNonLinearAnalysis=my.index,
+                               regressionRange=c(0.0025,0.100),
+                               useEmbeddings = (kEmbeddingDim-1):(kEmbeddingDim+1),
+                               doPlot=TRUE)
```

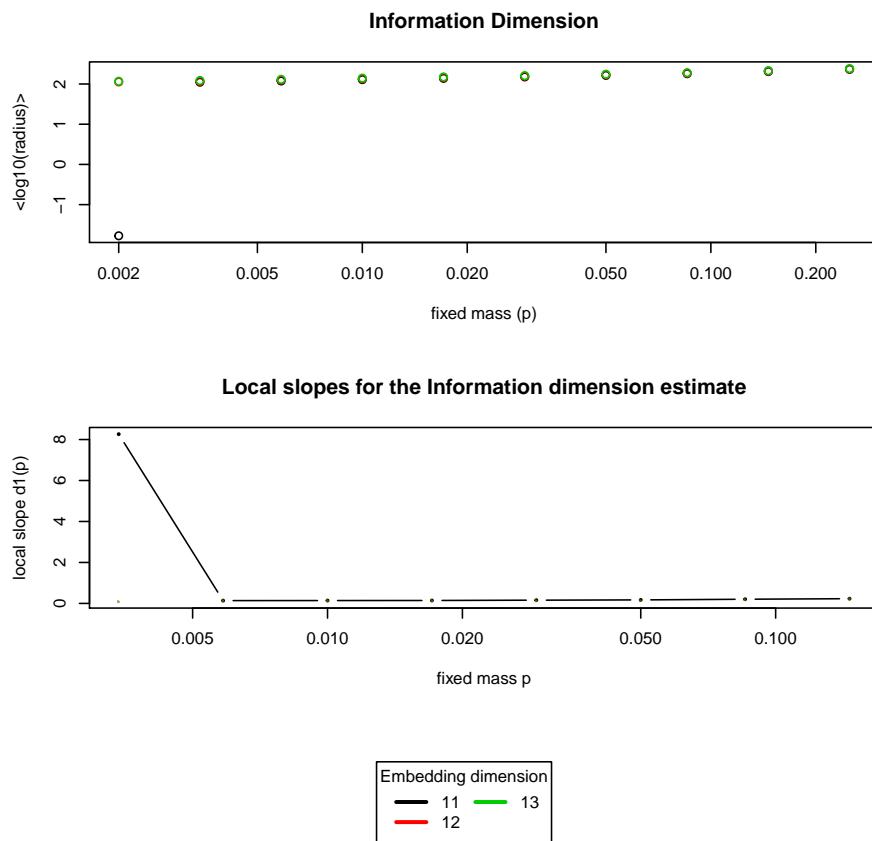


Figure 6.17: Information dimension computations in RHRV.

--- Estimating the information dimension ---

--- Information dimension = 6.945107 ---

```
> # Let's print again the value of the information dimension!
> cat("The information dimension is ",
+     hrv.data$NonLinearAnalysis[[my.index]]$infDim$statistic, "\n")
```

The information dimension is 6.945107

Figure 6.18 shows the regression performed for obtaining the estimate of the information dimension.

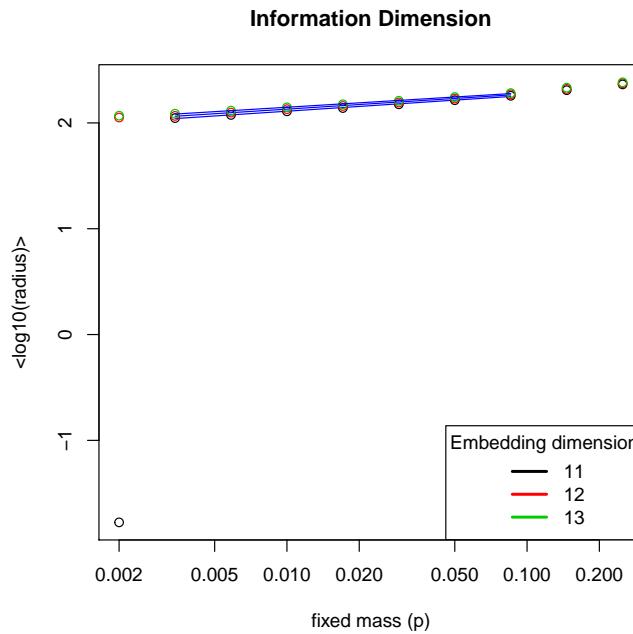


Figure 6.18: Information dimension estimation.

6.2.4 Poincaré Plot

The *PoincarePlot* function implements all the functionality for performing Poincaré plot analysis, whether you employ the “classic” calculation methods or those based on confidence regions (see Section 2.2.3.9).

6.2. ADVANCED NONLINEAR ANALYSIS TECHNIQUES

In order to compute the “classic” SD_1 and SD_2 parameters in the Poincaré plot, the most important arguments of the *RHRV* routine are the *timeLag* (τ , that must be setted to 1) and *doPlot* (if TRUE, the Poincaré plot is shown). The following lines illustrate the use of this function. The resulting Poincaré plot is shown in Figure 6.19.

```
> # Set timeLag = 1 to obtain the "classic" Poincare parameters  
> hrv.data = PoincarePlot(hrv.data,  
+                           indexNonLinearAnalysis=1,  
+                           timeLag=1, doPlot=TRUE)  
  
--- Calculating SD1 and SD2 parameters ---  
--- Creating Poincare Plot with time lag = 1 ---  
--- SD1 = 3.284373 ---  
--- SD2 = 7.931162 ---
```

Figure 6.19 shows the resulting Poincaré plot. *RHRV* also provides functionality for fitting the ellipse (and computing both SD_1 and SD_2 parameters) using the theory of the confidence regions. In order to enable the confidence region estimation the user can set *confidenceEstimation* = TRUE. The confidence level can be selected with the *confidence* parameter (default value is 0.95). It must be noted that when *timeLag* > 1, the confidence region approach is always used.

Figure 6.20 shows the result of applying the following piece of code:

```
> hrv.data = CreateNonLinearAnalysis(hrv.data)
```

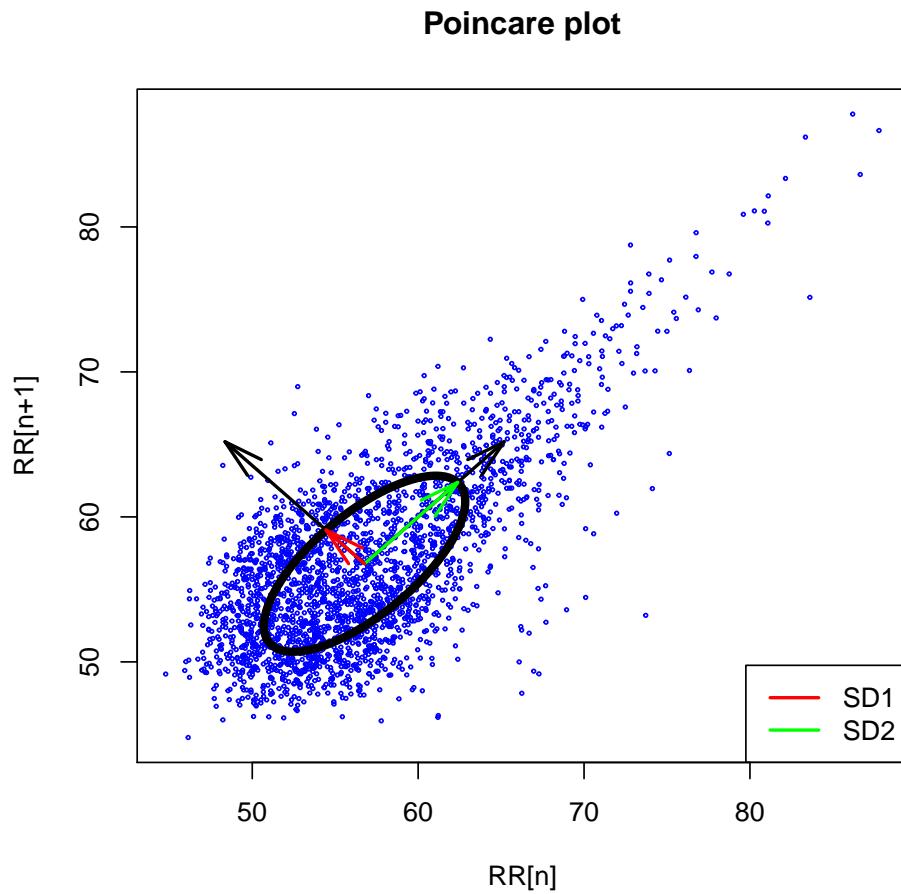


Figure 6.19: “Classic” Poincaré Plot.

** Creating non linear analysis

Data has now 3 nonlinear analysis

```
> hrv.data = PoincarePlot(hrv.data,
+                           indexNonLinearAnalysis=2,
+                           timeLag=1, confidenceEstimation = TRUE,
```

```

+           confidence = 0.9,
+
+           doPlot=TRUE)

--- Calculating SD1 and SD2 parameters ---
--- Creating Poincare Plot with time lag = 1 ---
--- SD1 = 7.048153 ---
--- SD2 = 17.01799 ---

```

The SD_1 and SD_2 parameters are stored under the *PoincarePlot* list of the *NonLinearAnalysis*. Of course, they can be accessed as usual:

```

> # results of the first fit ...
> print(hrv.data$NonLinearAnalysis[[1]]$PoincarePlot)

$SD1
[1] 3.284373

$SD2
[1] 7.931162

> # ... vs results of the second one
> print(hrv.data$NonLinearAnalysis[[2]]$PoincarePlot)

$SD1
[1] 7.048153

$SD2
[1] 17.01799

```

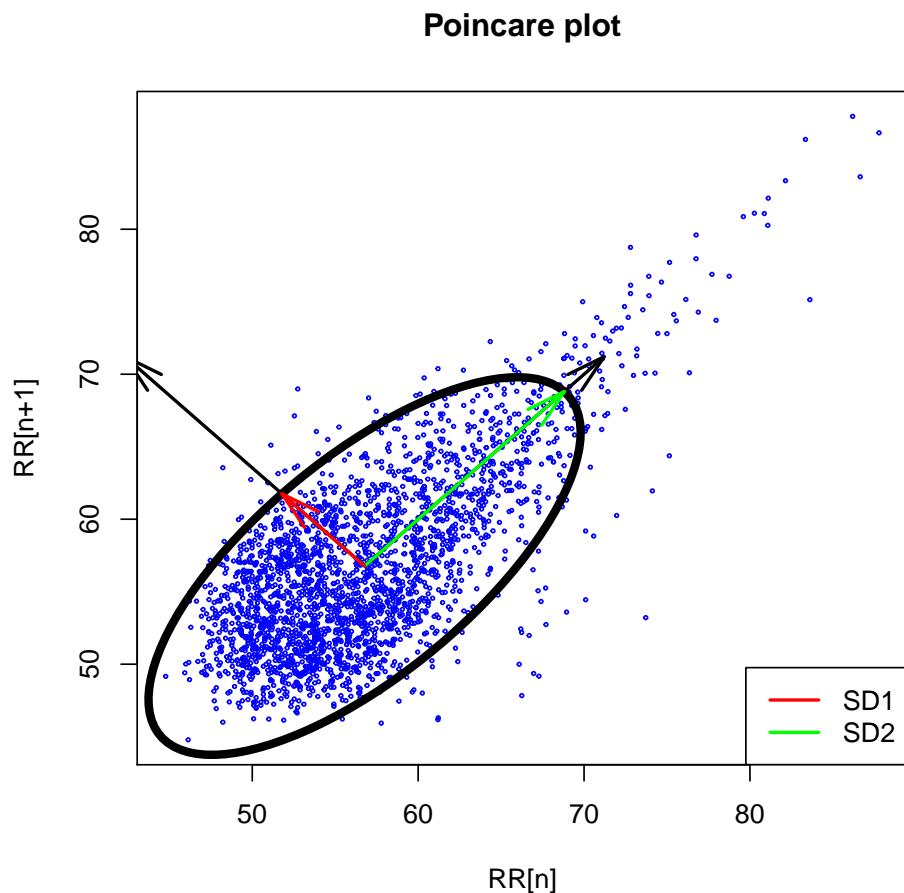


Figure 6.20: Poincaré Plot using the confidence region estimation.

Bibliography

- [1] Autonomic nervous system,
http://www.scholarpedia.org/article/Autonomic_nervous_system.
- [2] The comprehensive R archive network. <http://cran.r-project.org/>.
- [3] The R project for statistical computing. <http://www.r-project.org/>.
- [4] S. Akselrod, D. Gordon, F.A. Ubel, D.C. Shannon, A.C. Berger, and R.J. Cohen. Power spectrum analysis of heart rate fluctuation: a quantitative probe of beat-to-beat cardiovascular control. *Science*, 213(4504):220, 1981.
- [5] M.L. Appel, R.D. Berger, J.P. Saul, J.M. Smith, and R.J. Cohen. Beat to beat variability in cardiovascular variables: noise or music? *Journal of the American College of Cardiology*, 14(5):1139–1148, 1989.
- [6] R.D. Berger, S. Akselrod, D. Gordon, and R.J. Cohen. An efficient algorithm for spectral analysis of heart rate variability. *Biomedical Engineering, IEEE Transactions on*, (9):900–904, 1986.
- [7] G.G. Berntson. Heart rate variability: Origins, methods and interpretive caveats. *Psychophysiology*, 34:623–648, 1997.
- [8] Liangyue Cao. Practical method for determining the minimum embedding dimension of a scalar time series. *Physica D: Nonlinear Phenomena*, 110(1):43–50, 1997.
- [9] T. Force. Heart rate variability: standards of measurement, physiological interpretation and clinical use. task force of the european society of cardiology and the north american society of pacing and electrophysiology. *Circulation*, 93(5):1043–65, 1996.

BIBLIOGRAPHY

- [10] R. Furlan, S. Guzzetti, W. Crivellaro, S. Dassi, M. Tinelli, G. Baselli, S. Cerutti, F. Lombardi, M. Pagani, and A. Malliani. Continuous 24-hour assessment of the neural regulation of systemic arterial pressure and rr variabilities in ambulant subjects. *Circulation*, 81(2):537–547, 1990.
- [11] D. Gabor. Theory of communication. *Electrical Engineers-Part III: Radio and Communication Engineering, Journal of the Institution of*, 93(26):429–441, 1946.
- [12] C.A. García, A. Otero, X.A. Vila, and M.J. Lado. An open source tool for heart rate variability wavelet-based spectral analysis. In *International Joint Conference on Biomedical Engineering Systems and Technologies, BIOSIGNALS 2012*, 2012.
- [13] Constantino A. García, Abraham Otero, Xosé Vila, and David G. Márquez. A new algorithm for wavelet-based heart rate variability analysis. *Biomedical Signal Processing and Control*, 8(6):542–550, 2013.
- [14] A. L. Goldberger, L. A. N. Amaral, L. Glass, J. M. Hausdorff, P. Ch. Ivanov, R. G. Mark, J. E. Mietus, G. B. Moody, C.-K. Peng, and H. E. Stanley. PhysioBank, PhysioToolkit, and PhysioNet: Components of a new research resource for complex physiologic signals. *Circulation*, 101(23):e215–e220, 2000 (June 13). Circulation Electronic Pages: <http://circ.ahajournals.org/cgi/content/full/101/23/e215> PMID:1085218; doi: 10.1161/01.CIR.101.23.e215.
- [15] B.W. Hyndman and R.K. Mohn. A pulse modulator model for pacemaker activity. In *Digest of the 10th Int. Conf. Med. & Biol. Eng*, page 223, 1973.
- [16] M.V. Kamath, E.L. Fallen, et al. Power spectral analysis of heart rate variability: a noninvasive signature of cardiac autonomic function. *Critical reviews in biomedical engineering*, 21(3):245, 1993.
- [17] J. Kautzner and A. John Camm. Clinical relevance of heart rate variability. *Clinical cardiology*, 20(2):162–168, 1997.
- [18] J. Kautzner and A. John Camm. Clinical relevance of heart rate variability. *Clinical cardiology*, 20(2):162–168, 1997.
- [19] B. Kemp and J. Olivan. European data format plus (EDF+), an EDF alike standard format for the exchange of physiological data. *Clinical Neurophysiology*, 114(9):1755–1761, 2003.

BIBLIOGRAPHY

- [20] B-U Kohler, Carsten Hennig, and Reinhold Orlmeister. The principles of software qrs detection. *Engineering in Medicine and Biology Magazine, IEEE*, 21(1):42–57, 2002.
- [21] M.J. Lado, A.J. Méndez, L. Rodríguez-Liñares, A. Otero, and X.A. Vila. Nocturnal evolution of heart rate variability indices in sleep apnea. *Computers in Biology and Medicine*, 2012.
- [22] M. Malik and A.J. Camm. *Heart rate variability*. Futura Pub. Co., Armonk, NY, 1995.
- [23] A. Malliani, M. Pagani, F. Lombardi, and S. Cerutti. Cardiovascular neural regulation explored in the frequency domain. *Circulation*, 84(2):482–492, 1991.
- [24] Matlab. <http://www.mathworks.com/index.html>, 2011.
- [25] G.B. Moody. Wfdb applications guide. *Harvard-MIT Division of Health Sciences and Technology*, 10, 2003.
- [26] G.B. Moody and R.G. Mark. The MIT-BIH arrhythmia database on cd-rom and software for use with it. In *Computers in Cardiology*, pages 185–188, 1990.
- [27] Nevrokard-aHRV. <http://www.nevrokard.eu/index.html>, 2011.
- [28] T. Penzel, GB Moody, RG Mark, AL Goldberger, and JH Peter. The apnea-ecg database. In *Computers in Cardiology 2000*, pages 255–258. IEEE, 2000.
- [29] D.B. Percival and A.T. Walden. *Wavelet methods for time series analysis*, volume 4. Cambridge Univ Pr, 2006.
- [30] L. Rodríguez-Liñares, A.J. Méndez, M.J. Lado, D.N. Olivieri, X.A. Vila, and I. Gómez-Conde. An open source tool for heart rate variability spectral analysis. *Computer Methods and Programs in Biomedicine*, 2010.
- [31] J.P. Saul, P. Albrecht, R.D. Berger, and R.J. Cohen. Analysis of long term heart rate variability: methods, 1/f scaling and implications. *Computers in cardiology*, 14:419, 1988.
- [32] M.P. Tarvainen and J.P. Niskanen. Kubios HRV version 2.0 user's guide. *Department of Physics, University of Kuopio, Kuopio, Finland*, 2008.

BIBLIOGRAPHY

- [33] J. Vila, F. Palacios, J. Presedo, M. Fernandez-Delgado, P. Felix, and S. Barro. Time-frequency analysis of heart-rate variability. *Engineering in Medicine and Biology Magazine, IEEE*, 16(5):119–126, 1997.
- [34] XA Vila, MJ Lado, AJ Mendez, DN Olivieri, and L.R. Linares. An R package for heart rate variability analysis. In *Intelligent Signal Processing, 2009. WISP 2009. IEEE International Symposium on*, pages 217–222. IEEE, 2009.